




IBM Training

# IBM Operations Analytics Log Analysis 1.3 Administration

## Student Exercises

Course code TN630G ERC 1.0

April 2016



All files and material for this course are IBM copyright property covered by the following copyright notice.

© Copyright IBM Corp. 2016. All Rights Reserved.

US Government Users Restricted Rights: Use, duplication or disclosure restricted by GSA ADP Schedule Contract with IBM Corp.

IBM, the IBM logo, and [ibm.com](http://ibm.com) are trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the web at "Copyright and trademark information" at [www.ibm.com/legal/copytrade.shtml](http://www.ibm.com/legal/copytrade.shtml).

Trademarks or trademarks of Adobe Systems Incorporated in the United States, and/or other countries.

IT Infrastructure Library is a registered trademark of the Central Computer and Telecommunications Agency which is now part of the Office of Government Commerce.

Intel, Intel logo, Intel Inside, Intel Inside logo, Intel Centrino, Intel Centrino logo, Celeron, Intel Xeon, Intel SpeedStep, Itanium, and Pentium are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries.

Linux is a registered trademark of Linus Torvalds in the United States, other countries, or both.

Microsoft, Windows, Windows NT, and the Windows logo are trademarks of Microsoft Corporation in the United States, other countries, or both.

ITIL is a registered trademark, and a registered community trademark of the Office of Government Commerce, and is registered in the U.S. Patent and Trademark Office.

UNIX is a registered trademark of The Open Group in the United States and other countries.

Java and all Java-based trademarks and logos are trademarks or registered trademarks of Oracle and/or its affiliates.

Cell Broadband Engine is a trademark of Sony Computer Entertainment, Inc. in the United States, other countries, or both and is used under license therefrom.

Linear Tape-Open, LTO, the LTO Logo, Ultrium, and the Ultrium logo are trademarks of HP, IBM Corp. and Quantum in the U.S. and other countries.

The information contained in this publication is provided for informational purposes only. While efforts were made to verify the completeness and accuracy of the information contained in this publication, it is provided AS IS without warranty of any kind, express or implied. In addition, this information is based on IBM's current product plans and strategy, which are subject to change by IBM without notice. IBM shall not be responsible for any damages arising out of the use of, or otherwise related to, this publication or any other materials. Nothing contained in this publication is intended to, nor shall have the effect of, creating any warranties or representations from IBM or its suppliers or licensors, or altering the terms and conditions of the applicable license agreement governing the use of IBM software.

References in this publication to IBM products, programs, or services do not imply that they will be available in all countries in which IBM operates. Product release dates and/or capabilities referenced in this presentation may change at any time at IBM's sole discretion based on market opportunities or other factors, and are not intended to be a commitment to future product or feature availability in any way. Nothing contained in these materials is intended to, nor shall have the effect of, stating or implying that any activities undertaken by you will result in any specific sales, revenue growth, savings or other results.



# Contents

<b>About these exercises</b> .....	<b>v</b>
host image details (host2.tivoli.edu) .....	v
host image details (bivm.ibm.com) .....	v
<b>1 Overview and basic administration exercises</b> .....	<b>1</b>
Exercise 1 Managing the application and users .....	1
Exercise 2 Adding data sources .....	6
Adding a WebSphere SystemOut data source .....	7
Adding a WebSphere SystemErr Data source .....	11
Adding a web access data source .....	14
Testing the data sources .....	17
Configuring search dashboards .....	22
Saving searches and creating dashboards .....	27
Exercise 3 Deleting historical data .....	38
<b>2 Common configuration tasks exercises</b> .....	<b>44</b>
Exercise 1 Using the Generic Annotation Insight Pack .....	44
Exercise 2 Using the DSV toolkit .....	57
Creating the Insight Pack .....	57
Testing the Insight Pack .....	62
<b>3 Troubleshooting exercises</b> .....	<b>69</b>
There are no student exercises for this unit.	
<b>4 Alerts exercises</b> .....	<b>70</b>
Exercise 1 Creating alert actions .....	70
Index alert action .....	70
Log alert action .....	70
Email alert action .....	71
Exercise 2 Creating base conditions .....	73
WebSphere SystemOut log base condition .....	73
WebSphere SystemErr log base condition .....	74
Web access log base condition .....	75
Exercise 3 Testing base conditions .....	76
Exercise 4 Creating composite conditions .....	79
WebSphere multi-condition-window composite condition .....	79
Web access single-condition-count composite condition .....	80
Exercise 5 Testing composite conditions .....	81

<b>5 Hadoop Distributed File System (HDFS) integration exercises . . . . .</b>	<b>84</b>
Exercise 1 Configuring passwordless SSH . . . . .	84
Exercise 2 Configuring BigInsights and Hadoop . . . . .	89
Exercise 3 Configuring Log Analysis . . . . .	94
Exercise 4 Verifying the integration . . . . .	96
Exercise 5 Disabling the HDFS integration . . . . .	102
<b>6 Performance tuning exercises . . . . .</b>	<b>104</b>
Exercise 1 Tuning the EIF Receiver . . . . .	104
Exercise 2 Solr administration . . . . .	107
Optional steps: Verifying the Solr configuration change . . . . .	114
<b>7 Backing up and restoring IBM Operations Analytics Log Analysis exercises . . . . .</b>	<b>117</b>
There are no student exercises for this unit.	



# About these exercises

These exercises use two virtual machine images for the lab environment:

**host2:** This virtual machine runs the IBM® Operations Analytics Log Analysis application. The host name is **host2.tivoli.edu**. The IP address is **192.168.100.161**.

**bivm:** This virtual machine runs the IBM InfoSphere® BigInsights® software. The host name is **bivm.ibm.com**. The IP address is **192.168.100.166**.

## host image details (host2.tivoli.edu)

The following table describes the user names and passwords that are used with the **host2.tivoli.edu** virtual machine (host2).

User name	Password	Description
netcool	object00	Operating system user
root	object00	Operating system super user
unityadmin	unityadmin	Operations Analytics Log Analysis super user
unityuser	unityuser	Operations Analytics Log Analysis user

Python version 2.6.6 is installed on this host. Use the following command to verify the Python version:

```
python -V
```

## host image details (bivm.ibm.com)

The following table describes the user names and passwords that are used with the **bivm.ibm.com** virtual machine.

User name	Password	Description
biadmin	object00	Operating system user
root	object00	Operating system super user
biadmin	object00	InfoSphere BigInsights user






# 1 Overview and basic administration exercises

In these exercises, you perform basic administration tasks, including application management, user management, and data storage housekeeping.

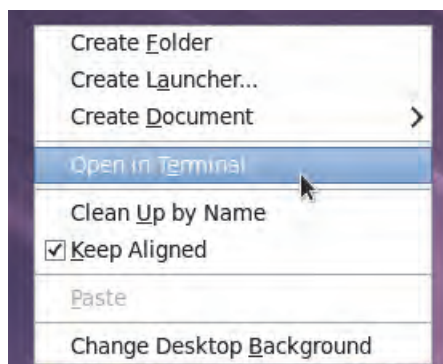
## Exercise 1 Managing the application and users

In this exercise, you start and stop the application and add users.

1. Log in to the host2 virtual machine (host2.tivoli.edu) with the user name **netcool** and the password **object00**.



2. Right-click the desktop and select **Open In Terminal**.



3. Use the following command to verify that you are logged in to the correct host:

```
hostname  
host2.tivoli.edu
```

4. Run the following command to verify that IBM Operations Analytics Log Analysis is running.

```
/opt/IBM/LogAnalysis/utilities/unity.sh -status
```

5. Stop all IBM Operations Analytics Log Analysis components. Verify that they are stopped. Use the following commands.

```
/opt/IBM/LogAnalysis/utilities/unity.sh -stop
```

```
/opt/IBM/LogAnalysis/utilities/unity.sh -status
```

Mon Apr 6 15:41:08 UTC 2016

IBM Operations Analytics - Log Analysis v1.3.0.0 Application Services Status:

```
-----  
No.  Service                      Status  Process ID  
-----  
1    Derby Network Server          DOWN  
2    ZooKeeper                      DOWN  
3    Websphere Liberty Profile      DOWN  
4    EIF Receiver                   DOWN  
5    Log File Agent instance        DOWN  
-----
```

Getting status of Solr on host2.tivoli.edu

Status of Solr Nodes:

```
-----  
No.  Instance Name                Host                Status  State  
-----  
1    SOLR_NODE_LOCAL                host2.tivoli.edu    DOWN    ACTIVE  
-----
```

All Application Services are in Stopped State



6. Start the IBM Operations Analytics Log Analysis components. Verify that they are running. Use the following commands.

```
/opt/IBM/LogAnalysis/utilities/unity.sh -start
```

```
/opt/IBM/LogAnalysis/utilities/unity.sh -status
```

```
Mon Apr 6 15:46:13 UTC 2015
```

```
IBM Operations Analytics - Log Analysis v1.3.0.0 Application Services Status:
```

```
-----  
No.   Service                               Status   Process ID  
-----  
1     Derby Network Server                   UP       8432  
2     ZooKeeper                               UP       8476  
3     Websphere Liberty Profile               UP       8638  
4     EIF Receiver                             UP       8810  
5     Log File Agent instance                 UP       9057  
-----
```

```
Getting status of Solr on host2.tivoli.edu
```

```
Status of Solr Nodes:
```

```
-----  
No. Instance Name                        Host                Status   State  
-----  
1   SOLR_NODE_LOCAL                      host2.tivoli.edu    UP       ACTIVE  
-----
```

```
All Application Services are in Running State
```

```
Checking server initialization status: Server has initialized!
```

7. Use the `unity.sh` utility again to show product version information. Run the following command.

```
/opt/IBM/LogAnalysis/utilities/unity.sh -version
```

```
Product: IBM Operations Analytics - Log Analysis
```

```
Version: 1
```


```
Release: 3
```

```
Modification: 0
```

```
Fix pack: 0
```

```
Build Id: 201503090631
```

8. Verify that you can log in to the IBM Operations Analytics Log Analysis user interface. Use the user name **unityadmin** and the password **unityadmin**.
  - a. Double-click the **Firefox** icon on the desktop.



- b. Browse to the following address.  
`https://host2.tivoli.edu:9987/Unity`
  - c. Log in to the user interface with the user name **unityadmin** and the password **unityadmin**. This action verifies that the user interface is running. Log out when you finish.



You add users to IBM Operations Analytics Log Analysis by editing the basic user registry file. This file is named `unityUserRegistry.xml`. Passwords in this file are encoded.

9. Edit the basic user registry file and add the following two users. Make the password for both users **object00**.

- **admin**: Add this user to the UnityAdmins and the UnityUsers groups.
- **loguser**: Add this user to the UnityUsers group.

- a. Use the following commands to encode the password **object00**. The output of the command is the encoded version of the password. Record the encoded password.

```
cd /opt/IBM/LogAnalysis/wlp/bin
./securityUtility encode object00
{xor}MD01Ojwrb28=
```

- b. Open the `unityUserRegistry.xml` file in a text editor. This example uses `vi`.

```
cd /opt/IBM/LogAnalysis/wlp/usr/servers/Unity
vi unityUserRegistry.xml
```

- c. Add the following five lines to the file. Make the unityUserRegistry.xml file look like the following example.

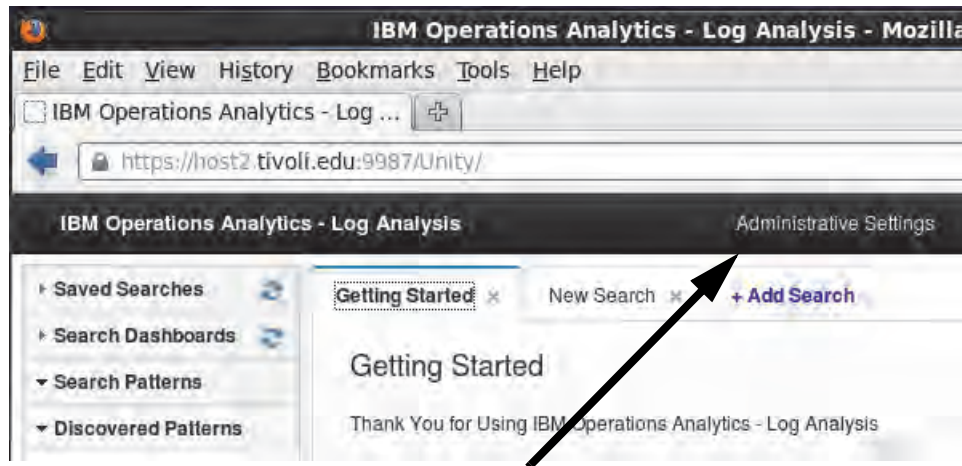
```
<server>
  <basicRegistry id="basic" realm="UnityRealm">
    <user name="unityuser" password="{xor}KjE2KyYqLDot" />
    <user name="unityadmin" password="{xor}KjE2KyY+OzI2MQ==" />
    <user name="admin" password="{xor}MD010jwrb28=" />
    <user name="loguser" password="{xor}MD010jwrb28=" />
    <group name="UnityUsers">
      <member name="unityuser" />
      <member name="unityadmin" />
      <member name="admin" />
      <member name="loguser" />
    </group>
    <group name="UnityAdmins">
      <member name="unityadmin" />
      <member name="admin" />
    </group>
  </basicRegistry>
</server>
```

```
<server>
  <basicRegistry id="basic" realm="UnityRealm">
    <user name="unityuser" password="{xor}KjE2KyYqLDot" />
    <user name="unityadmin" password="{xor}KjE2KyY+OzI2MQ==" />
    <user name="admin" password="{xor}MD010jwrb28=" />
    <user name="loguser" password="{xor}MD010jwrb28=" />
    <group name="UnityUsers">
      <member name="unityuser" />
      <member name="unityadmin" />
      <member name="admin" />
      <member name="loguser" />
    </group>
    <group name="UnityAdmins">
      <member name="unityadmin" />
      <member name="admin" />
    </group>
  </basicRegistry>
</server>
```

- d. Save and exit the file after you finish editing it.

## 10. Test the new users.

- a. Open a Firefox browser and browse to <https://host2.tivoli.edu:9987/Unity>. Log in to the user interface with the user name **admin** and the password **object00**.
- b. Verify that there is an **Administrative Settings** link in the user interface, which means that the **admin** user has administrative privileges. Log out of the user interface.



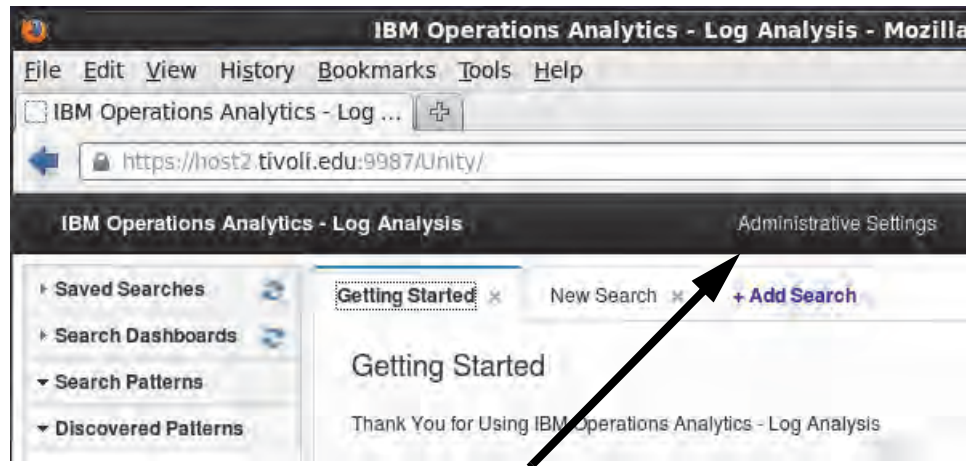
- c. Log back in to the user interface with the user name **loguser** and the password **object00**. Notice that **loguser** does not have a link to the Administrative Settings page.
- d. Log out of the user interface.

## Exercise 2 Adding data sources

In this exercise, you add some data sources and view log data in the user interface.

1. Open the administration user interface.
  - a. Double-click the **Firefox** icon on the desktop.
  - b. Browse to the following address.  
<https://host2.tivoli.edu:9987/Unity>
  - c. Log in to the user interface with the user name **unityadmin** and the password **unityadmin**.

- d. Click **Administrative Settings**. The administration user interface opens in a new Firefox tab.



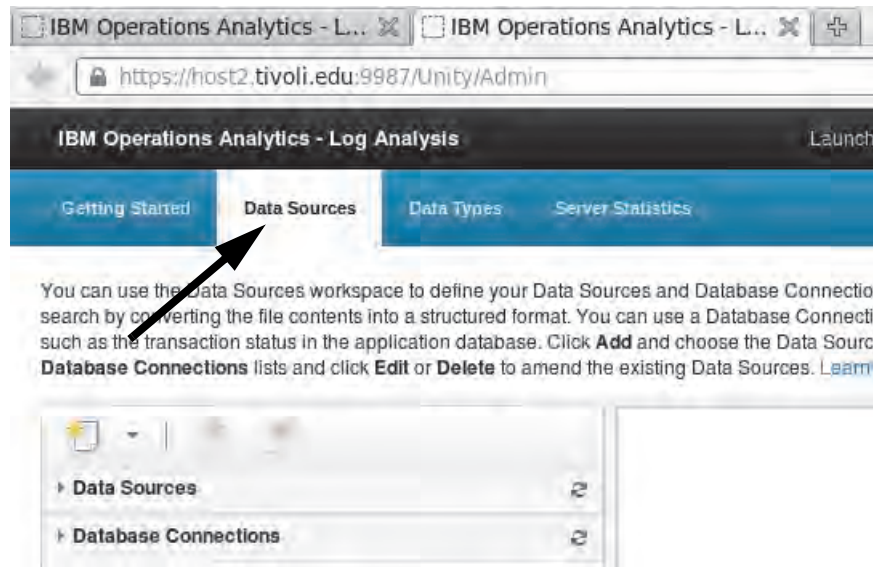
## Adding a WebSphere SystemOut data source

A **data source** is a reference to a log file.

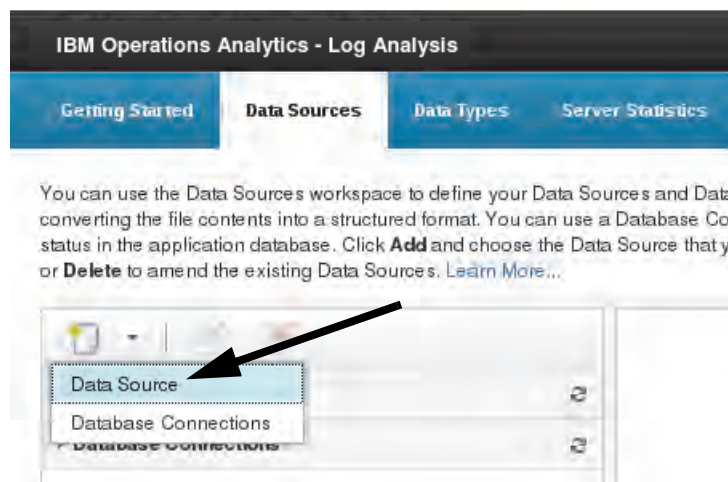
2. Create a data source named **WAS\_SystemOut**. Use the values in the following table to complete the data source wizard.

Field	Value
Location	Select Local file
Host name	host2.tivoli.edu
File Path	/software/log_samples/WAS_logs/SystemOut.log
Type	WASSystemOut
Collection	Leave this field blank
Name	WAS_SystemOut
Description	This source uses WAS SystemOut logs
Group	Leave this field blank

- a. Click the **Data Sources** tab in the administration user interface. The administration user interface is in the second Firefox tab.



- b. Click **Add > Data Source**.





- c. Select **Local file**. Click **Next**.

#### Add Data Source

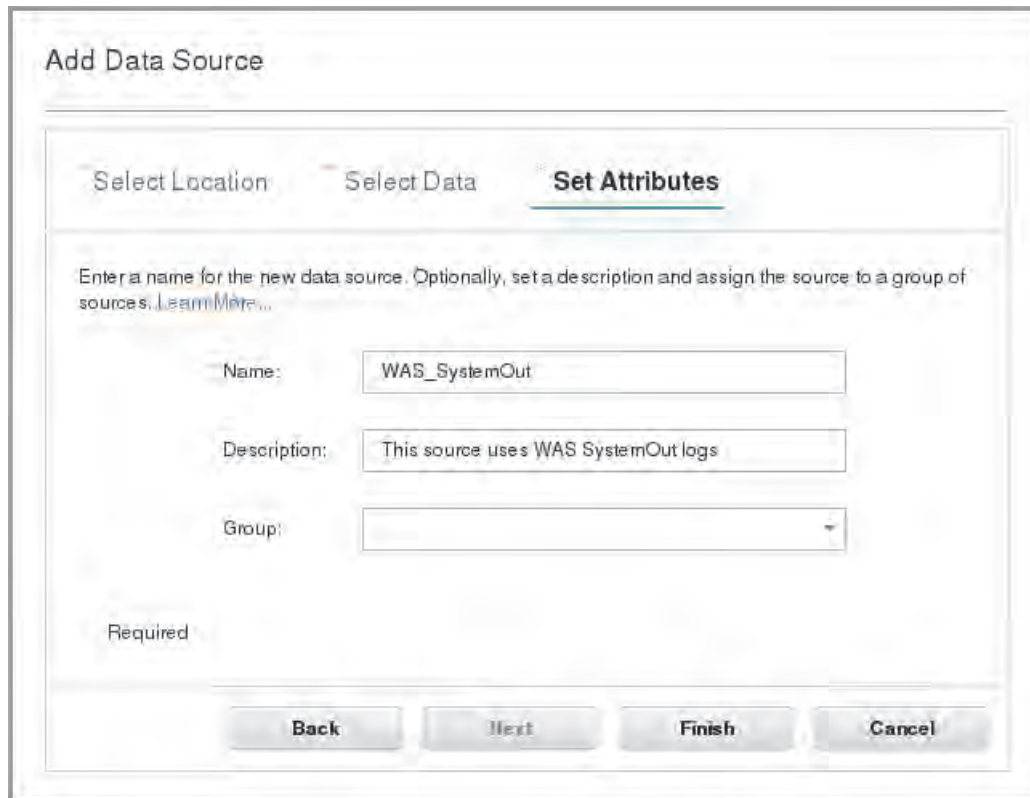
The screenshot shows the 'Add Data Source' wizard with three tabs: 'Select Location', 'Select Data', and 'Set Attributes'. The 'Select Location' tab is active. It contains a text box with instructions: 'If you want to ingest data into the Log Analytics server, use the wizard to configure a data source. Select Local or Remote file to monitor changes to a file. Select Custom when data is sent to the Log Analytics server from external sources such as a remote log file agent, Logstash, or the data collector client. [Learn More](#).' Below this are three radio buttons: 'Local file' (selected), 'Remote file', and 'Custom'. A 'Host name:' label is followed by a text box containing 'host2.livoli.edu'. A 'Required' label is at the bottom left. At the bottom are four buttons: 'Back', 'Next', 'Finish', and 'Cancel'.

- d. Enter `/software/log_samples/WAS_logs/SystemOut.log` as the file path. Select **WASSystemOut** as the type. Click **Next**.

#### Add Data Source

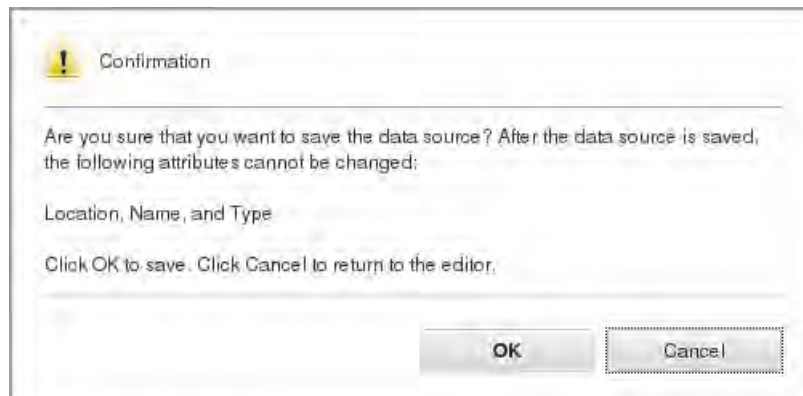
The screenshot shows the 'Add Data Source' wizard with three tabs: 'Select Location', 'Select Data', and 'Set Attributes'. The 'Select Data' tab is active. It contains a text box with instructions: 'Enter the location and type of data for this data source. The file path is not validated when you select the custom option. [Learn More](#).' Below this are two fields: 'File path:' with a text box containing '/software/log\_samples/WAS\_logs/SystemOut.log' and a 'Browse' button, and 'Type:' with a dropdown menu showing 'WASSystemOut'. There is a checkbox labeled 'This source of data is a rolling file' which is unchecked. A 'Required' label is at the bottom left. At the bottom are four buttons: 'Back', 'Next', 'Finish', and 'Cancel'.

- e. Enter **WAS\_SystemOut** as the name of the data source. Use the description from the preceding table. Click **Finish**.



The image shows the 'Add Data Source' dialog box with the 'Set Attributes' tab selected. The dialog has three tabs: 'Select Location', 'Select Data', and 'Set Attributes'. Below the tabs, there is a text area with the instruction: 'Enter a name for the new data source. Optionally, set a description and assign the source to a group of sources. [Learn More...](#)'. There are three input fields: 'Name' with the value 'WAS\_SystemOut', 'Description' with the value 'This source uses WAS SystemOut logs', and 'Group' which is an empty dropdown menu. At the bottom left, there is a 'Required' label. At the bottom right, there are four buttons: 'Back', 'Next', 'Finish', and 'Cancel'.

- f. Click **OK** in the confirmation windows.



The image shows a 'Confirmation' dialog box with a yellow warning icon. The text inside reads: 'Are you sure that you want to save the data source? After the data source is saved, the following attributes cannot be changed: Location, Name, and Type. Click OK to save. Click Cancel to return to the editor.' At the bottom right, there are two buttons: 'OK' and 'Cancel'.



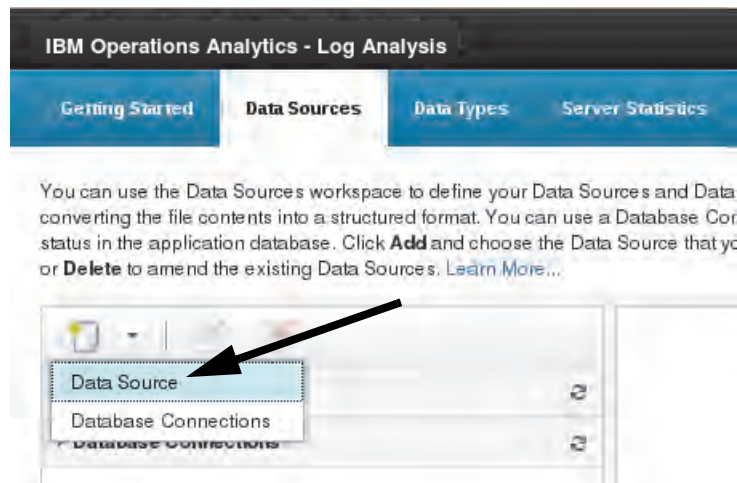


## Adding a WebSphere SystemErr Data source

3. Create a data source named **WAS\_SystemErr**. Use the values in the following table to complete the data source wizard.

Field	Value
Location	Select Local file
Host name	host2.tivoli.edu
File Path	/software/log_samples/WAS_logs/SystemErr.log
Type	WASSystemErr
Collection	Leave this field blank
Name	WAS_SystemErr
Description	This source uses WAS SystemErr logs
Group	Leave this field blank

- a. Click the **Data Sources** tab in the administration user interface. The administration user interface is in the second Firefox tab.
- b. Click **Add > Data Source**.



- c. Select **Local file**. Click **Next**.

#### Add Data Source

The screenshot shows the 'Add Data Source' wizard. The 'Select Location' step is active. The 'Local file' radio button is selected. The 'Host name' field contains 'host2.fivoli.edu'. The 'Next' button is highlighted.

- d. Enter `/software/log_samples/WAS_logs/SystemErr.log` as the file path. Select **WASSystemErr** as the Type. Click **Next**.

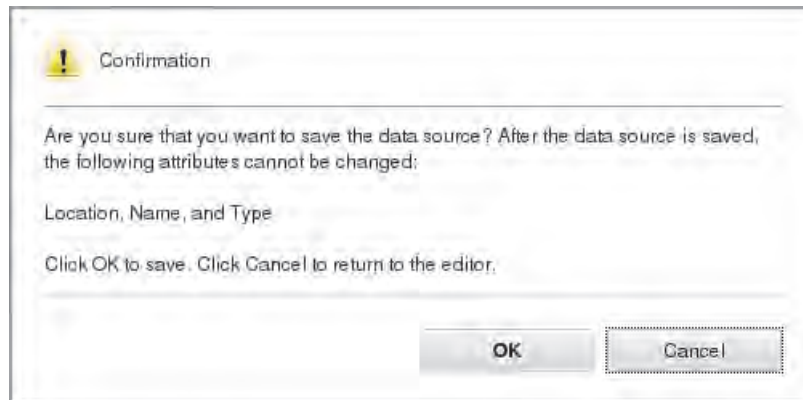
#### Add Data Source

The screenshot shows the 'Add Data Source' dialog with the 'Select Data' tab selected. The dialog has three tabs: 'Select Location', 'Select Data', and 'Set Attributes'. Below the tabs, there is a text box for 'File path' containing '/software/log\_samples/WAS\_logs/SystemErr.log' and a 'Browse' button. Below that is a dropdown menu for 'Type' with 'WASSystemErr' selected. There is a checkbox labeled 'This source of data is a rolling file' which is unchecked. At the bottom, there are four buttons: 'Back', 'Next', 'Finish', and 'Cancel'.

- e. Enter **WAS\_SystemErr** as the name of the data source. Use the description from the preceding table. Click **Finish**.

The screenshot shows the 'Add Data Source' dialog with the 'Set Attributes' tab selected. The dialog has three tabs: 'Select Location', 'Select Data', and 'Set Attributes'. Below the tabs, there is a text box for 'Name' containing 'WAS\_SystemErr'. Below that is a text box for 'Description' containing 'This source uses WAS SystemErr logs'. Below that is a dropdown menu for 'Group' which is empty. At the bottom, there are four buttons: 'Back', 'Next', 'Finish', and 'Cancel'.

- f. Click **OK** in the confirmation windows.

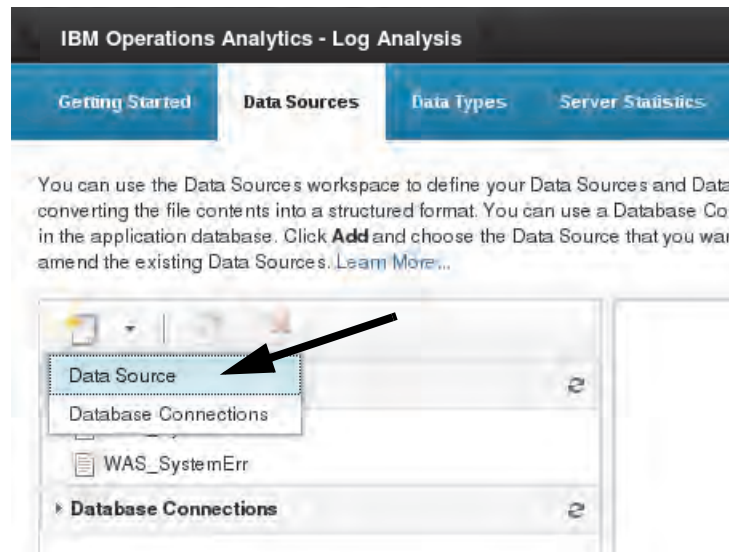


## Adding a web access data source

4. Create a data source named **Web\_Server**. Use the values in the following table to complete the data source wizard.

Field	Value
Location	Select Remote file
Host name	host2.tivoli.edu
User name	<b>netcool</b>
Password	<b>object00</b>
File Path	/software/log_samples/IHS_logs/IHS-access.log
Type	WebAccessLog
Collection	Leave this field blank
Name	Web_Server
Description	This source uses an IBM HTTP Server log
Group	Leave this field blank

- a. Click the **Data Sources** tab in the administration user interface. The administration user interface is in the second Firefox tab.
- b. Click **Add > Data Source**.



- c. Select **Remote file**. Enter `host2.tivoli.edu` as the host name. Enter **netcool** and **object00** as the user name and the password. Click **Next**.

#### Add Data Source

The screenshot shows the 'Add Data Source' wizard. The 'Select Location' tab is selected. The 'Remote file' radio button is selected. The 'Host name' field contains 'host2.tivoli.edu', the 'User name' field contains 'netcool', and the 'Password' field contains 'object00'. The 'Next' button is highlighted.

- d. Enter `/software/log_samples/IHS_logs/IHS-access.log` as the file path. Select **WebAccessLog** as the Type. Click **Next**.

## Add Data Source

The screenshot shows the 'Add Data Source' wizard with the 'Select Data' tab selected. The 'File path' field contains '/software/log\_samples/IHS\_logs/IHS-access.log' and the 'Type' dropdown is set to 'WebAccessLog'. A 'Browse' button is next to the file path field. Below these fields is a checkbox labeled 'This source of data is a rolling file' which is unchecked. At the bottom are buttons for 'Back', 'Next', 'Finish', and 'Cancel'.

Select Location **Select Data** Set Attributes

Enter the location and type of data for this data source. The file path is not validated when you select the custom option. [Learn More...](#)

File path:

Type:

☐ This source of data is a rolling file

Required

- e. Enter **Web\_Server** as the name of the data source. Use the description from the preceding table. Click **Finish**.

## Add Data Source

The screenshot shows the 'Add Data Source' wizard with the 'Set Attributes' tab selected. The 'Name' field contains 'Web\_Server', the 'Description' field contains 'This source uses an IBM HTTP Server log', and the 'Group' dropdown is empty. At the bottom are buttons for 'Back', 'Next', 'Finish', and 'Cancel'.

Select Location Select Data **Set Attributes**

Enter a name for the new data source. Optionally, set a description and assign the source to a group of sources. [Learn More...](#)

Name:

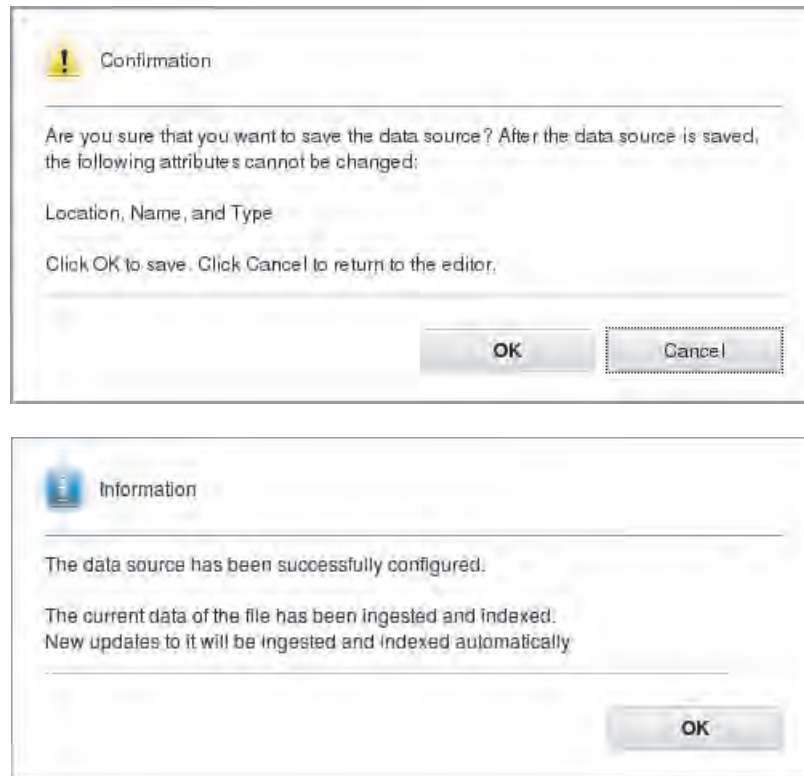
Description:

Group:

Required



- f. Click **OK** in the confirmation windows.



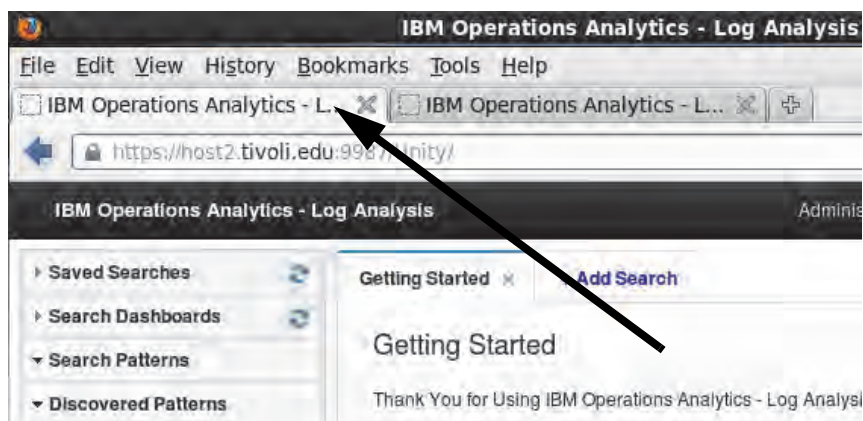
## Testing the data sources

The WebSphere® log files, `SystemOut.log` and `SystemErr.log`, for this lab are generated by scripts.

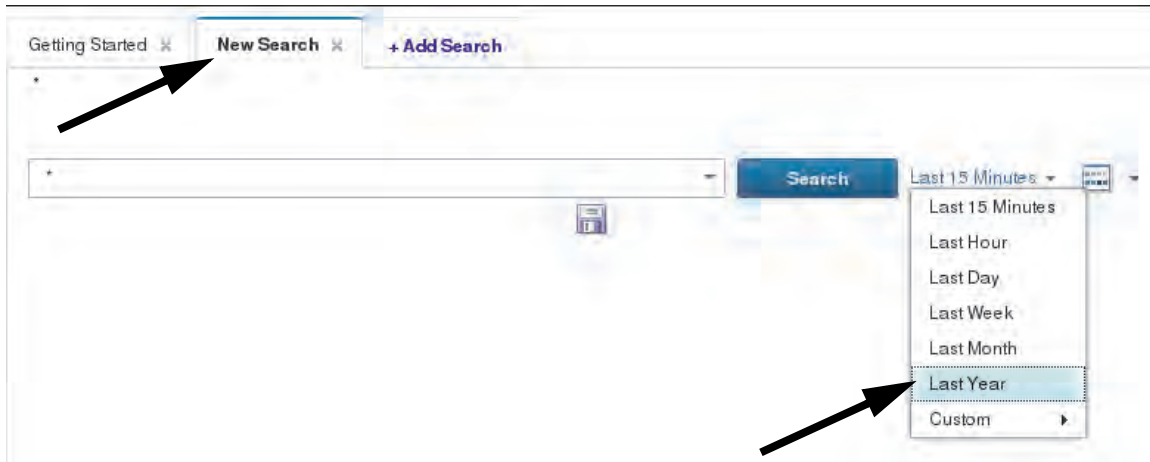
5. Return to a terminal window. Run the following command to generate events in the WebSphere log files.

```
/software/log_samples/scripts/WAS_Logs.sh
```

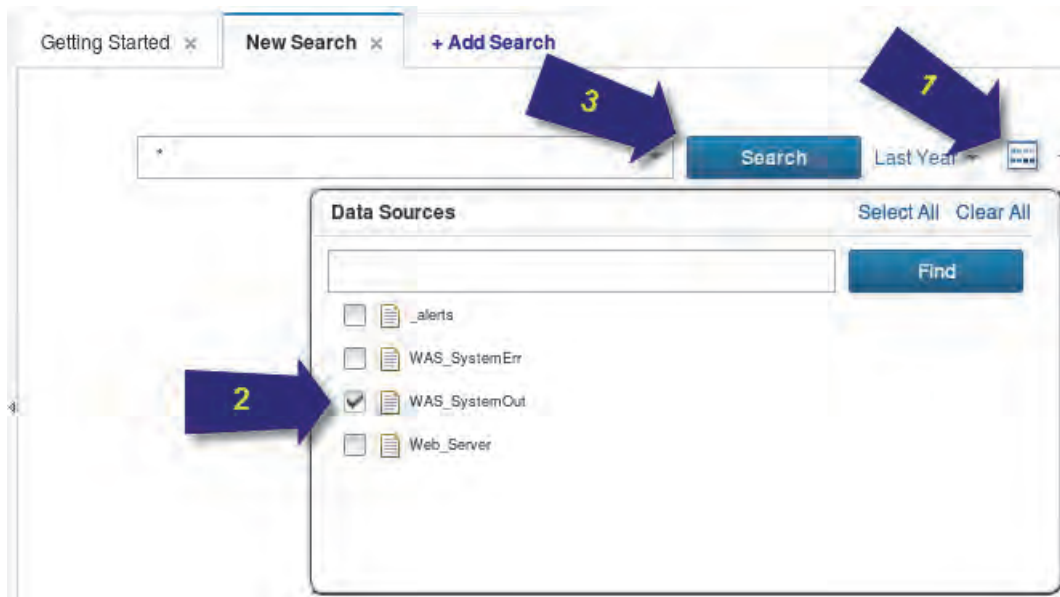
6. Return to the user interface by clicking the first Firefox tab.



7. Search through the WebSphere logs.
  - a. Click the **New Search** tab. Select **Last Year** as the time filter.



- b. Select only **WAS\_SystemOut** as the log source. Click **Search**.

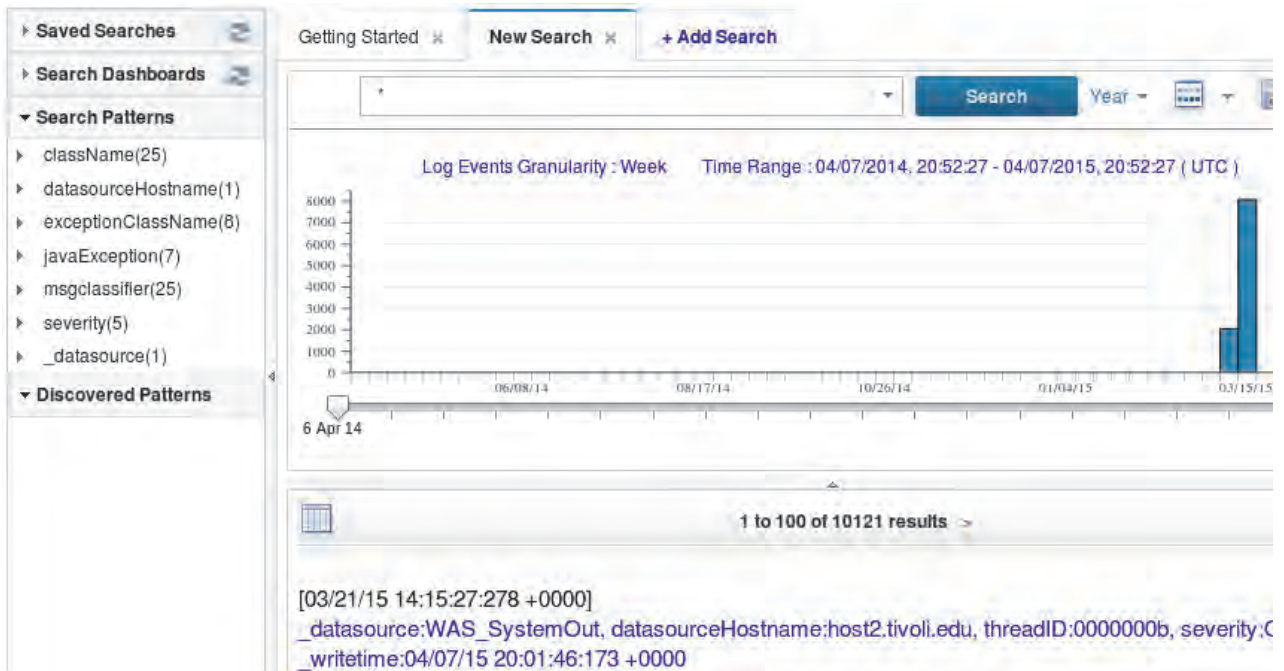


- c. Log events load in to the user interface. There are many events in this sample. Notice the summary on the left of the window. This summary shows the patterns that are found in the log file and the number of events that follow these patterns.

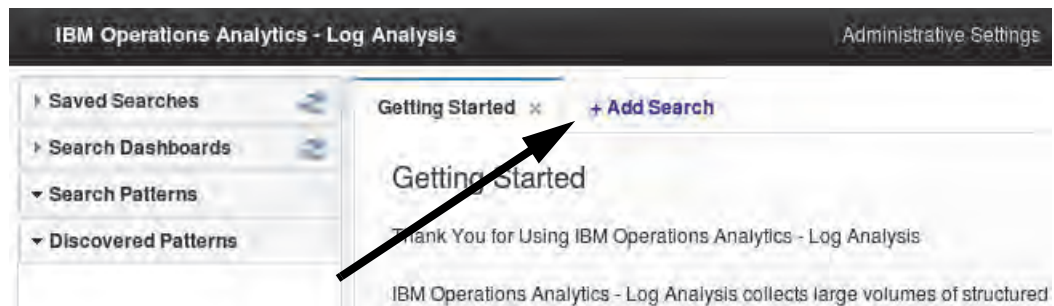


**Note:** New messages in the log file are processed only after a time-out setting expires. If you do not see any events in the search results, wait 90 seconds and try the search again.

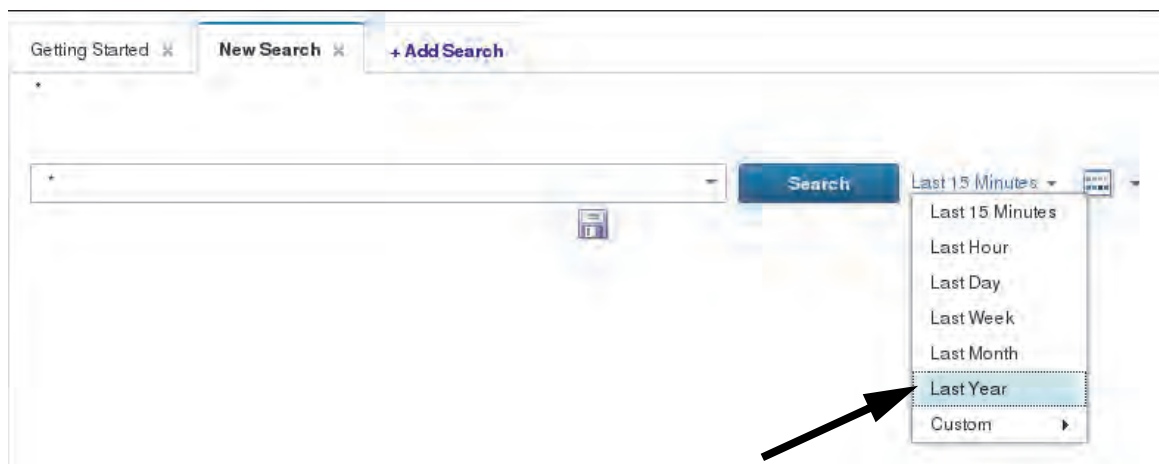




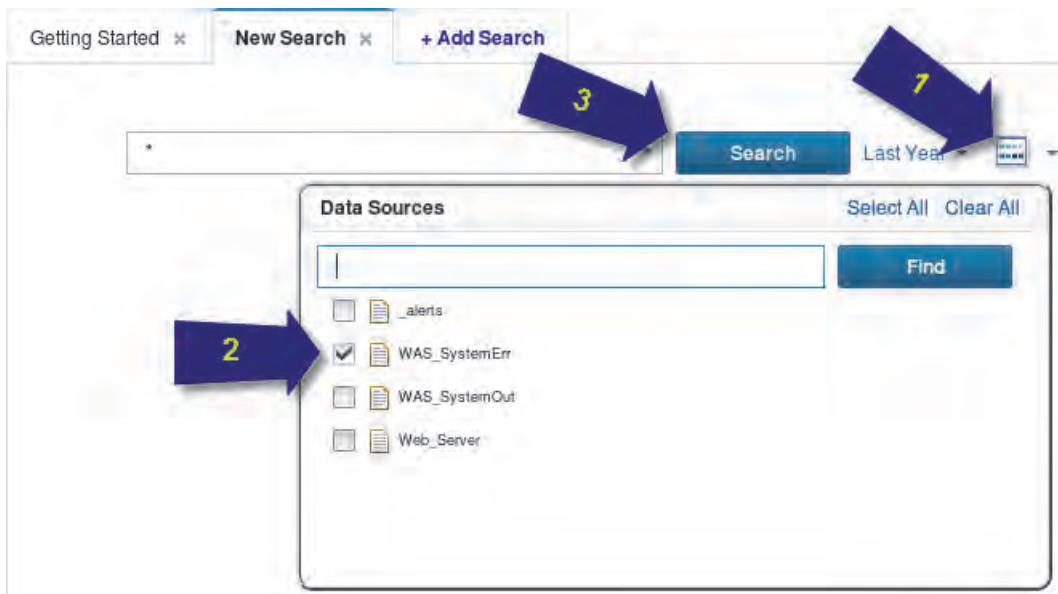
- d. Close the **New Search** tab. Click **OK** to confirm.
- e. Click **Add Search**. A new search tab opens.



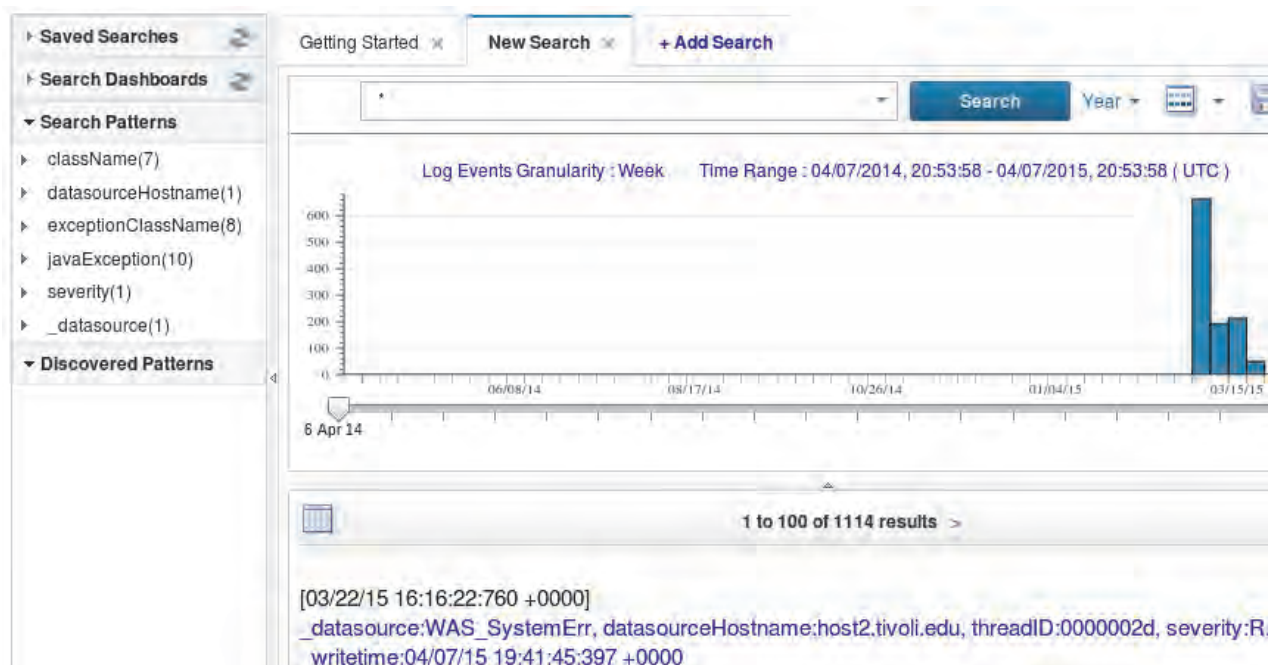
- f. Select **Last Year** as the time filter.



- g. Select only **WAS\_SystemErr** as the log source. Click **Search**.

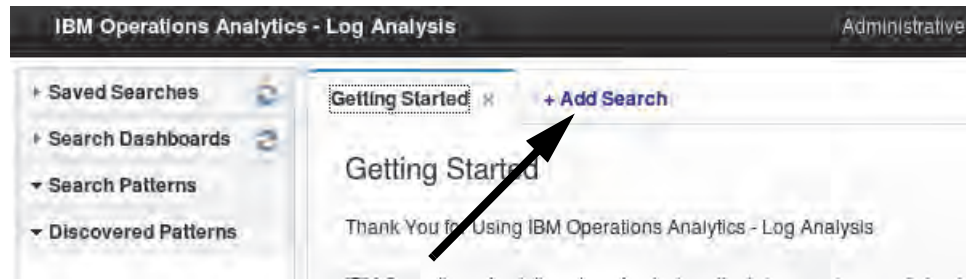


- h. Log events load in to the user interface. There are many events in this sample. Close the **New Search** tab when you finish.

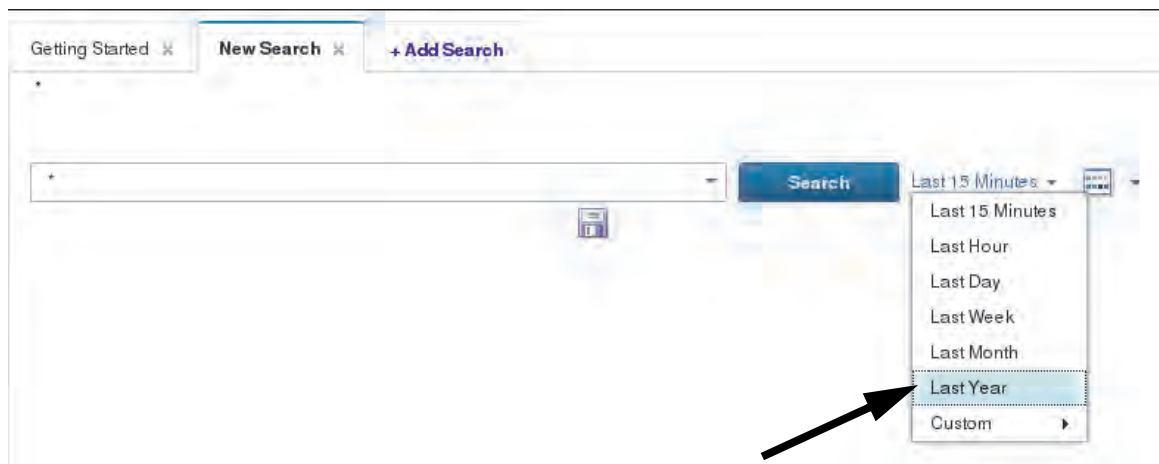


8. Run the following command to generate events in the web access log.  
`/software/log_samples/scripts/Web_Logs.sh`

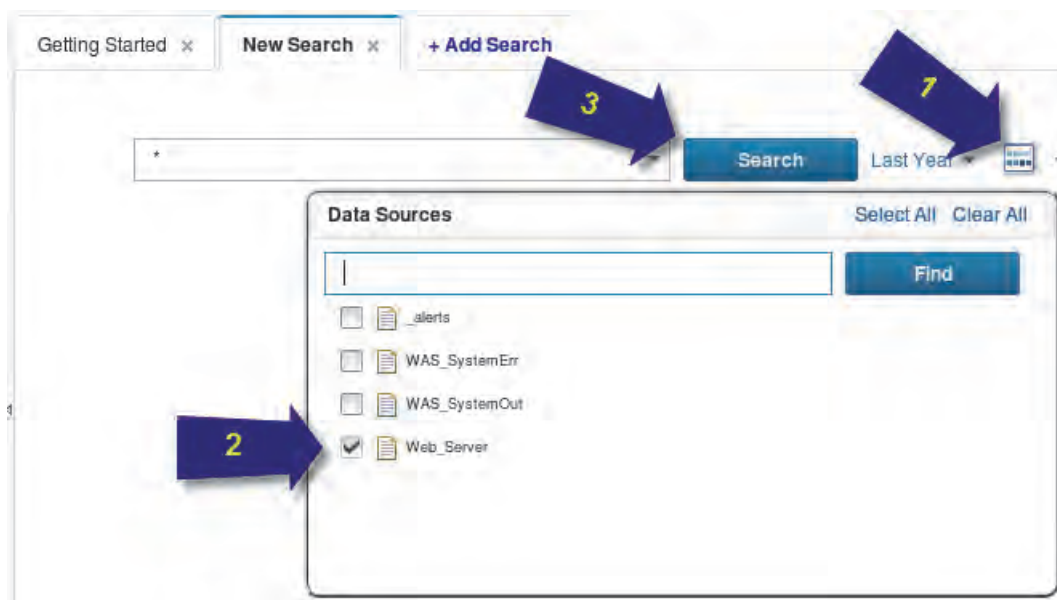
9. Search through the web access log.
- a. Click **Add Search**. A **New Search** tab opens.



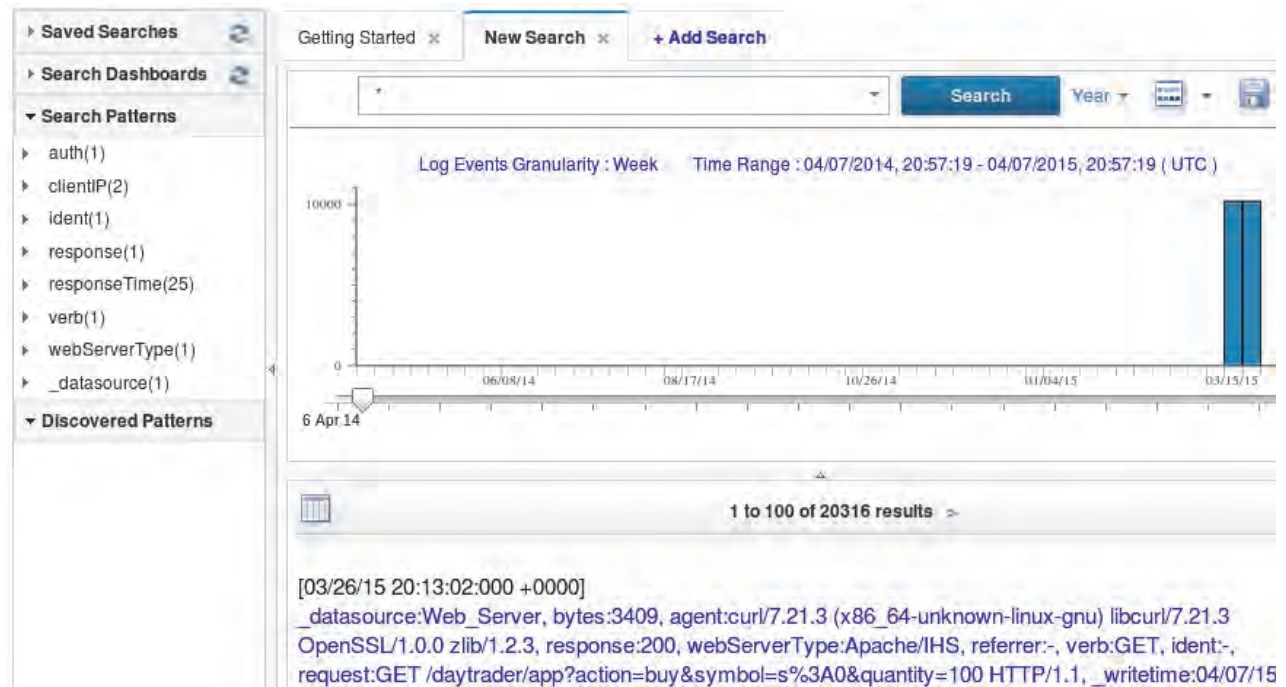
- b. Select **Last Year** as the time filter.



- c. Select only **Web\_Server** as the log source. Click **Search**.



- d. Log events load in to the user interface. Close the **New Search** tab when you finish.



## Configuring search dashboards

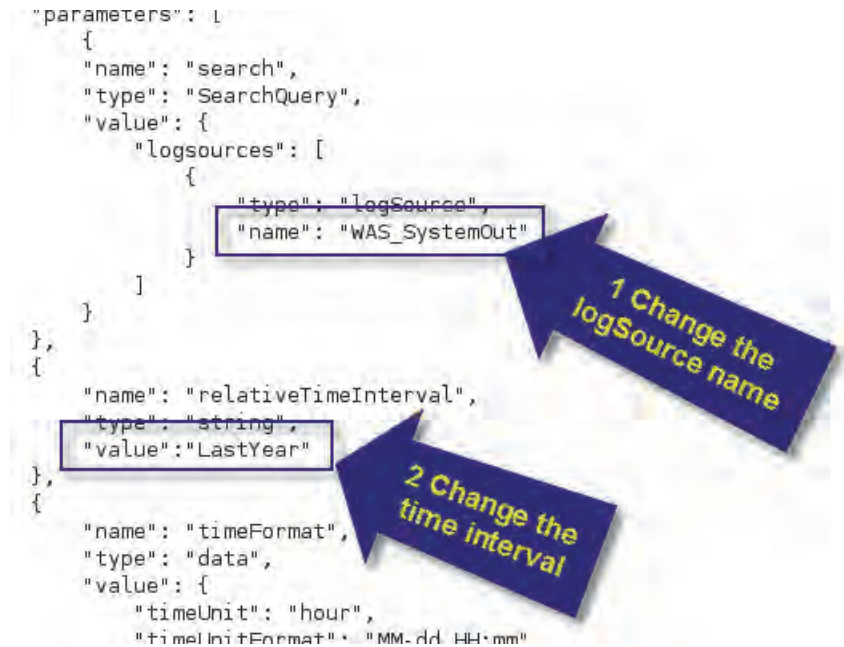
Search dashboards show event data in charts and graphs. You must configure the dashboards that are included with Insight Packs before you can use them.

10. Edit the WebSphere search dashboard to match your environment.
  - a. Return to the terminal window. Run the following commands to back up the WebSphere dashboard.
 

```
cd /opt/IBM/LogAnalysis/AppFramework/Apps/WASAppInsightPack_v1.1.0.3
cp WAS_Troubleshooting.app WAS_Troubleshooting.app.orig
```
  - b. Open the WAS\_Troubleshooting.app file in a text editor. This example uses vi.
 

```
vi WAS_Troubleshooting.app
```

- c. Change the following two lines in the file. Make the `WAS_Troubleshooting.app` file look like the following example.
- ◆ Change the name of the **logSource** to `WAS_SystemOut`.
  - ◆ Change the **relativeTimeInterval** to **LastYear**.



- d. Save and exit the file after you finish editing it.
11. Edit the Web Health Check dashboard to match your environment.
- a. Run the following commands to back up the Web Health Check dashboard.
 

```

cd /opt/IBM/LogAnalysis/AppFramework/Apps/WebAccessLogInsightPack_v1.1.0.2/
cp Web\ Health\ Check.appExmpl Web\ Health\ Check.app

```
  - b. Open the `Web\ Health\ Check.app` file in a text editor. This example uses `vi`.
 

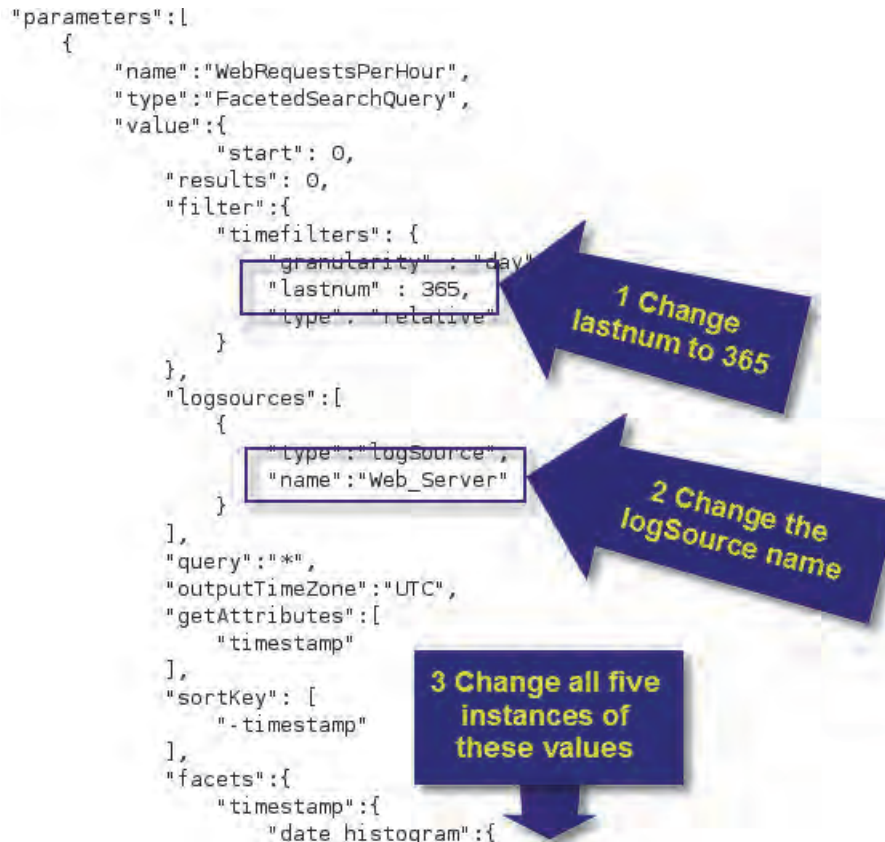
```

vi Web\ Health\ Check.app

```



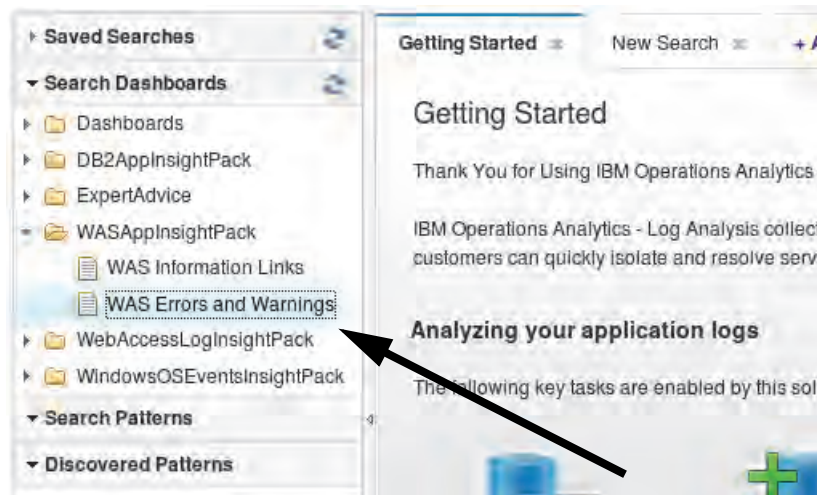
- c. Change the following 10 lines in the file. Make the `Web\ Health\ Check.app` file look like the following example.
- ◆ Change **lastnum** to **365**. There are five instances of this line in the file. Change all instances.
  - ◆ Change the name of **logSource** to **Web\_Server**. There are five instances of this line in the file. Change all instances.



- d. Save and exit the file after you finish editing it.

12. Return to the user interface. Test the WebSphere dashboard.

- a. Double-click **Search Dashboards > WASAppInsightPack > WAS Errors and Warnings** on the left of the window.



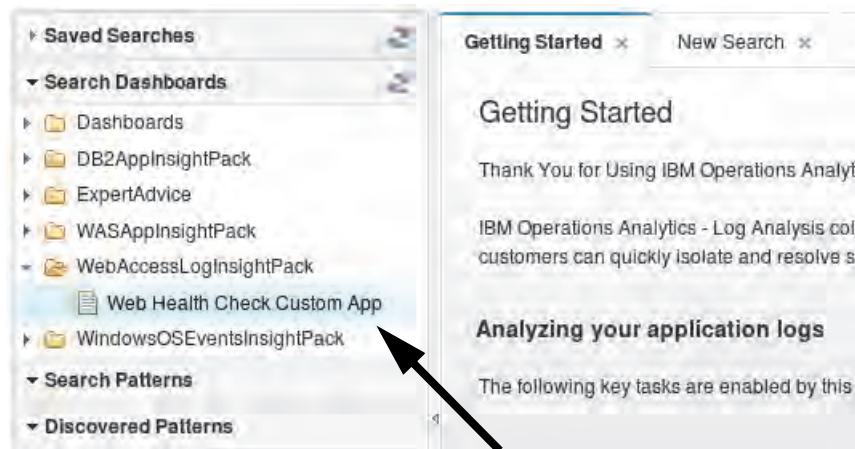
- b. Scroll down and look at the WAS Errors and Warnings page. This page shows useful troubleshooting graphics, such as number of errors and warnings or top five message counts.



- c. Close the dashboard page.

13. Test the Web Health Check dashboard.

- a. Double-click **Search Dashboards > WebAccessLogInsightPack > Web Health Check Custom App** on the left of the window.



- b. Scroll down and look at the Web Health Check dashboard page.

This page shows useful troubleshooting graphics, such as response time or number of requests. The charts on this page are customizable. The data points in these charts are interactive. You can drill down to specific log messages from this dashboard.



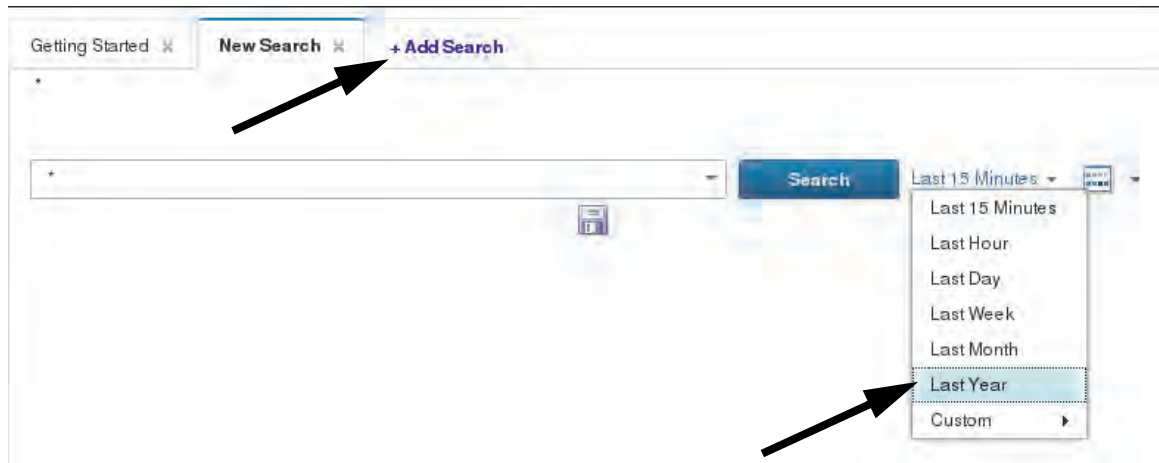
- c. Close the dashboard page.



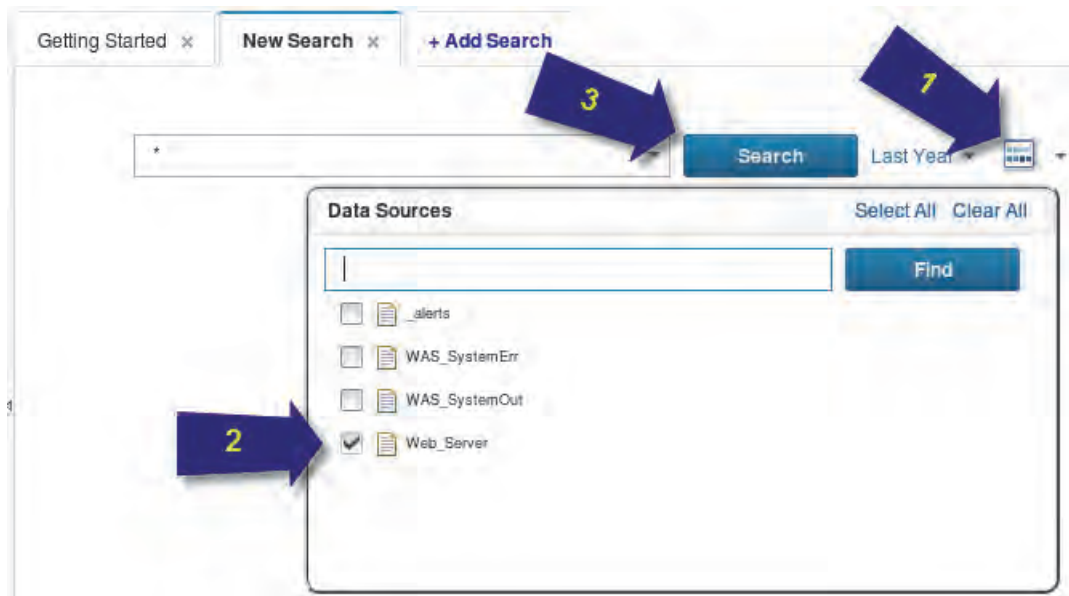
## Saving searches and creating dashboards

14. Search through the web server log:

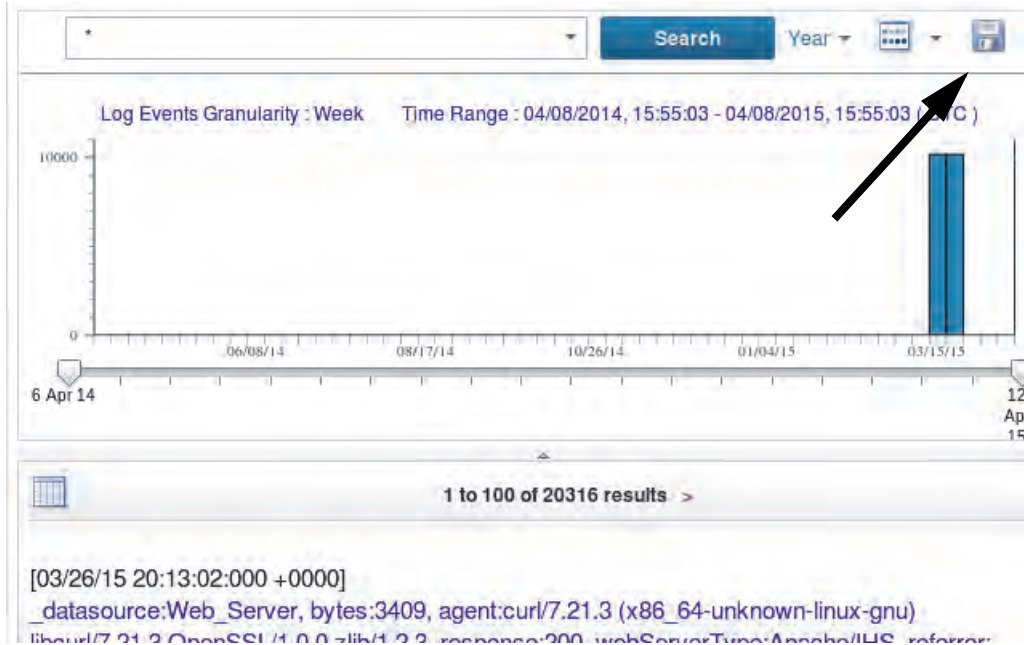
- a. Click the **Add Search** tab. Select **Last Year** as the time filter.



- b. Select only **Web\_Server** as the source. Click **Search**.



15. Log events load in to the user interface. Save your search:
- Click the **Save** button.



- Enter `Web_Diag` as the **Name** and click **OK**.

Save Quick Search

Enter a name for your search. To group similar searches, enter a tag name. Click **OK** to add your search, quick searches marked with the same tag are displayed under the same folder in the Quick Searches pane

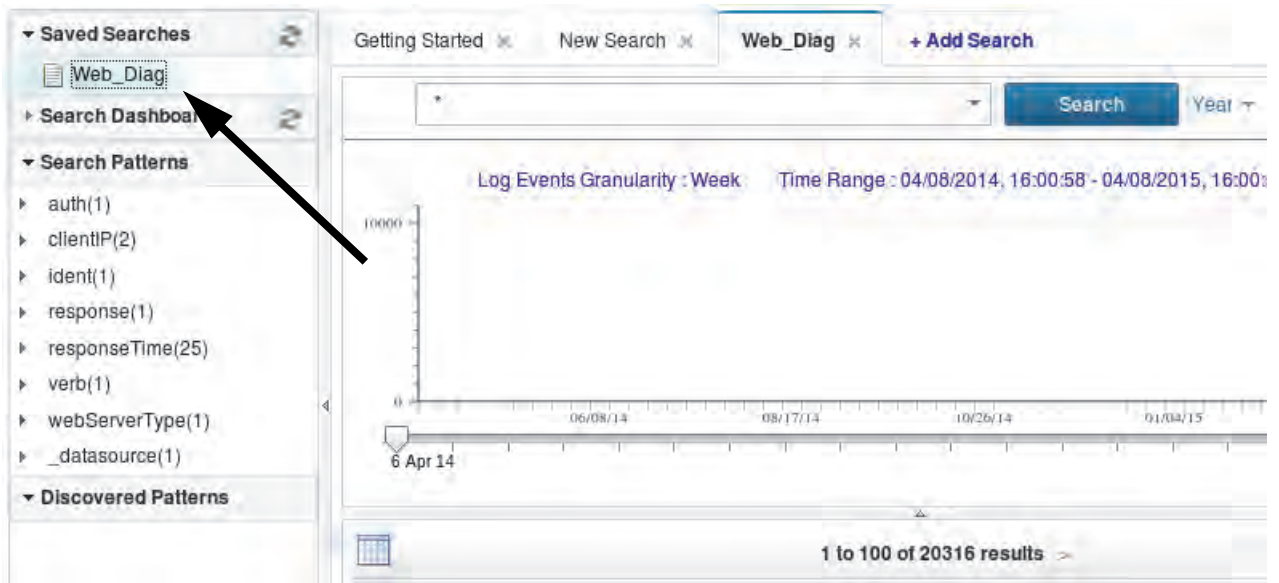
Name

Tag

Time Filter

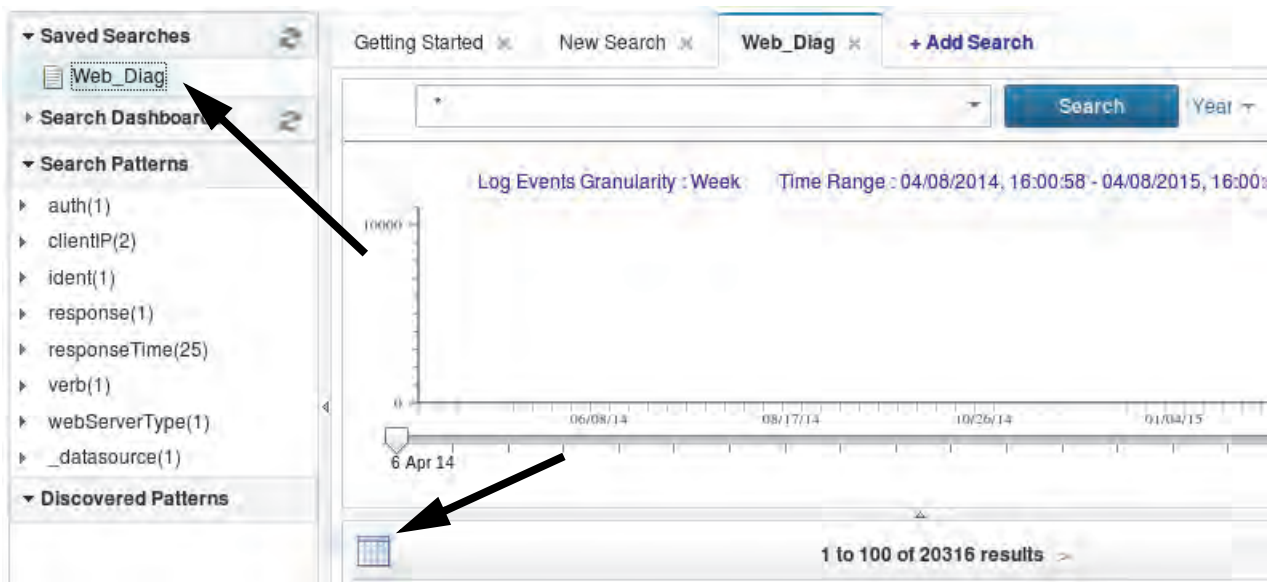
☐ Preserve the fixed time period: 4/8/2014, 3:55 PM - 4/8/2015, 3:55 PM

- c. Double-click the **Saved Searches > Web\_Diag** link to verify that the search was saved.

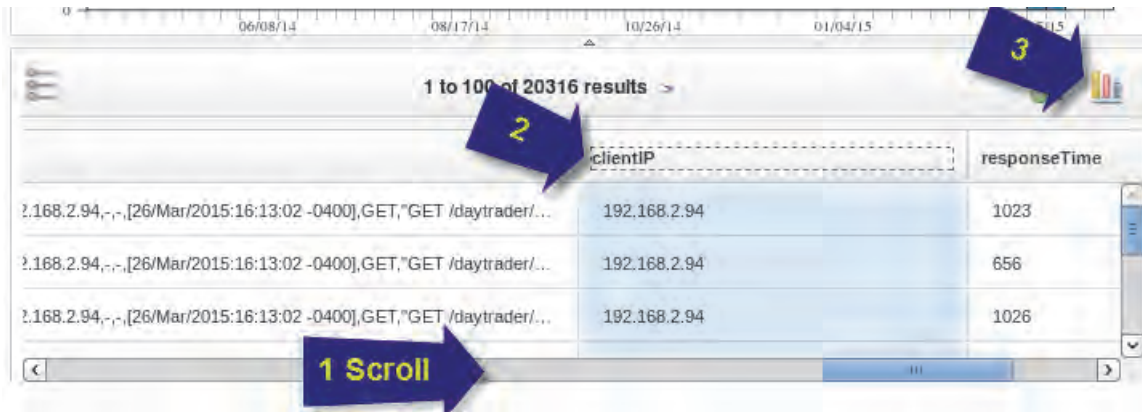


Dashboard pages contain charts that show log data. You can create charts and dashboards from any log file. In this example, you use a web access log file.

16. Create a dashboard page with two charts that show data from the web access log:
- Configure the first chart to show the number of messages by client IP. Run the **Web\_Diag** saved search. Click the **Grid View** button.



- b. Scroll to the right in the search results and click the **clientIP** column. Click the **Plot Column** button.



1 to 100 of 20316 results

	clientIP	responseTime
192.168.2.94,-,-,[26/Mar/2015:16:13:02 -0400],GET,"GET /daytrader/...	192.168.2.94	1023
192.168.2.94,-,-,[26/Mar/2015:16:13:02 -0400],GET,"GET /daytrader/...	192.168.2.94	656
192.168.2.94,-,-,[26/Mar/2015:16:13:02 -0400],GET,"GET /daytrader/...	192.168.2.94	1026

Annotations: Arrow 1 points to the scroll bar. Arrow 2 points to the 'clientIP' column header. Arrow 3 points to the 'Plot Column' button in the top right corner.


- c. Select **Generate Counts**.  
d. Click **Plot Chart (All Data)**.

### Plot Chart

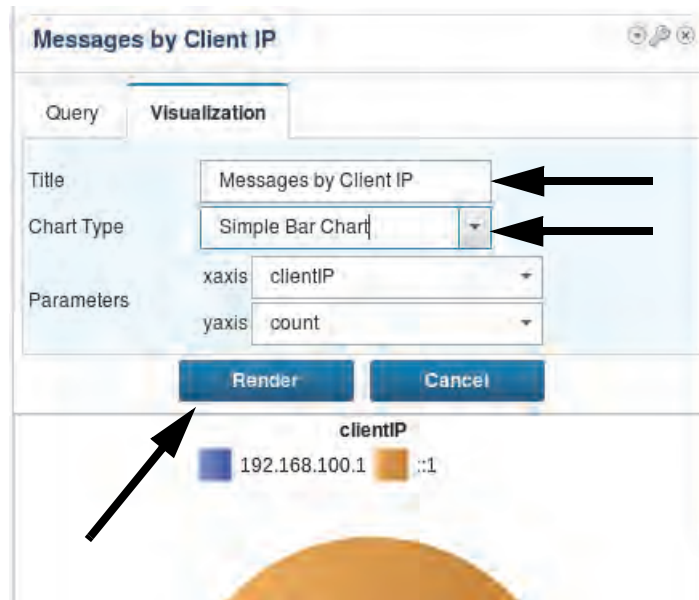
☒ Generate Count  
**Selected Columns**  
clientIP

Cancel Plot Chart (Current Page Data) Plot Chart (All Data)

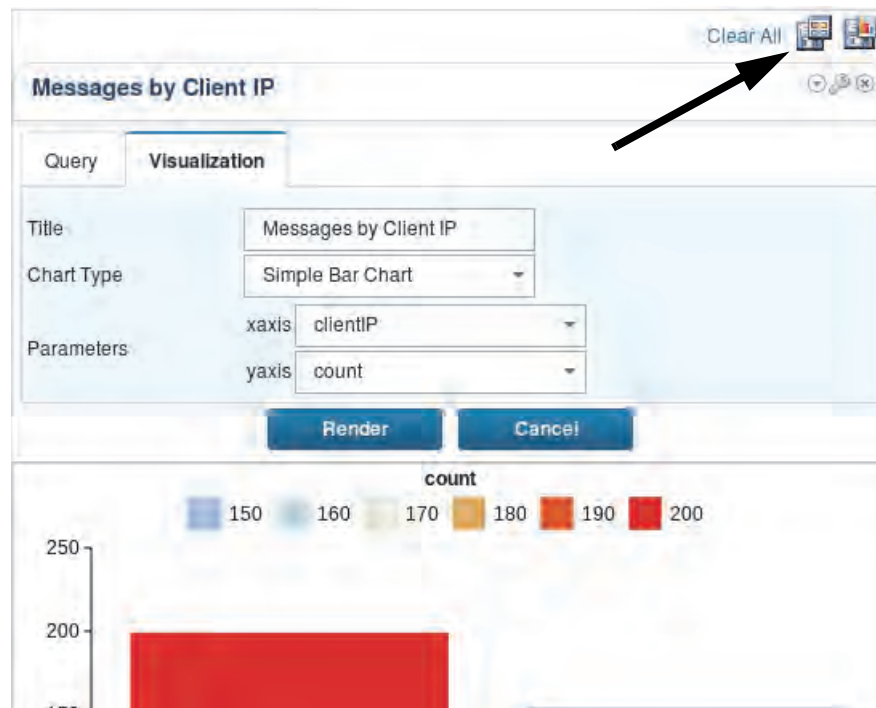
- e. Edit the chart by clicking the **Settings** button at the upper right of the chart.



- f. Change the title to `Messages by Client IP`. Change the chart type to **Simple Bar Chart**. Click **Render**.

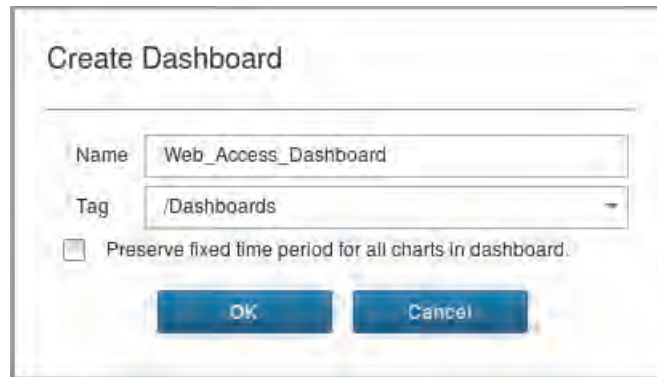


- g. Click the **Create New Dashboard** button at the top of the chart.



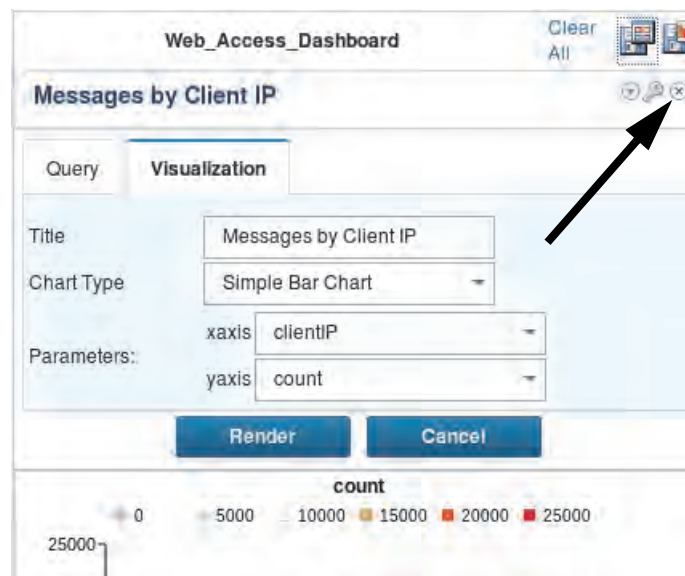
- h. Enter `Web_Access_Dashboard` for the dashboard name.

- i. Enter `/Dashboards` into the **Tag** field and click **OK**.



The 'Create Dashboard' dialog box is shown. It has a title bar 'Create Dashboard'. Below it, there are two input fields: 'Name' with the value 'Web\_Access\_Dashboard' and 'Tag' with the value '/Dashboards'. Below these fields is a checkbox labeled 'Preserve fixed time period for all charts in dashboard.' which is currently unchecked. At the bottom of the dialog are two buttons: 'OK' and 'Cancel'.

- j. Close the chart page. After you close the chart page, you return to the search results.



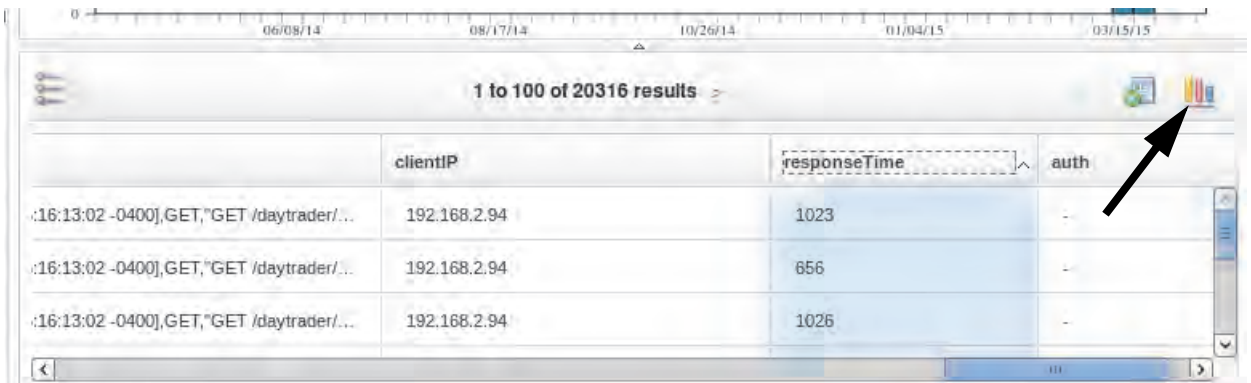
The 'Web\_Access\_Dashboard' configuration window is shown. It has a title bar 'Web\_Access\_Dashboard' and a 'Clear All' button. Below the title bar is a section titled 'Messages by Client IP'. There are two tabs: 'Query' and 'Visualization'. The 'Visualization' tab is selected. It contains a 'Title' field with the value 'Messages by Client IP', a 'Chart Type' dropdown menu with 'Simple Bar Chart' selected, and 'Parameters' for 'xaxis' (clientIP) and 'yaxis' (count). Below these fields are 'Render' and 'Cancel' buttons. An arrow points to the 'Close' button (represented by an 'X' icon) in the top right corner of the visualization configuration area.

- k. Create the second chart. Scroll in the search results and select these two column titles: **timestamp** and **responseTime**. Use the Ctrl key to select multiple columns. Click the **Plot Column** button.



**Note:** If you click the **timestamp** column first, it is used as the x-axis of the chart.





1 to 100 of 20316 results

	clientIP	responseTime	auth
:16:13:02 -0400],GET,"GET /daytrader/...	192.168.2.94	1023	-
:16:13:02 -0400],GET,"GET /daytrader/...	192.168.2.94	656	-
:16:13:02 -0400],GET,"GET /daytrader/...	192.168.2.94	1026	-

l. Click **Plot Chart (Current Page Data)**.

### Plot Chart

☐ Generate Count

**Selected Columns**

timestamp

responseTime

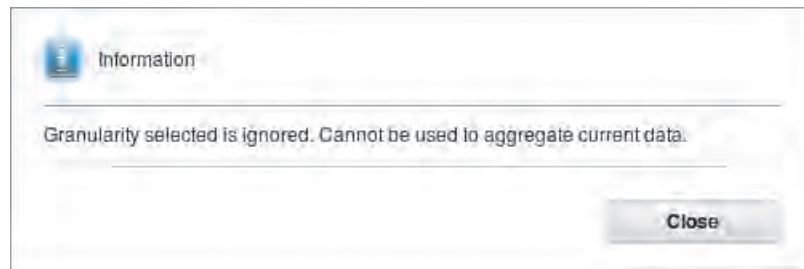
**Summary function**

☐ min

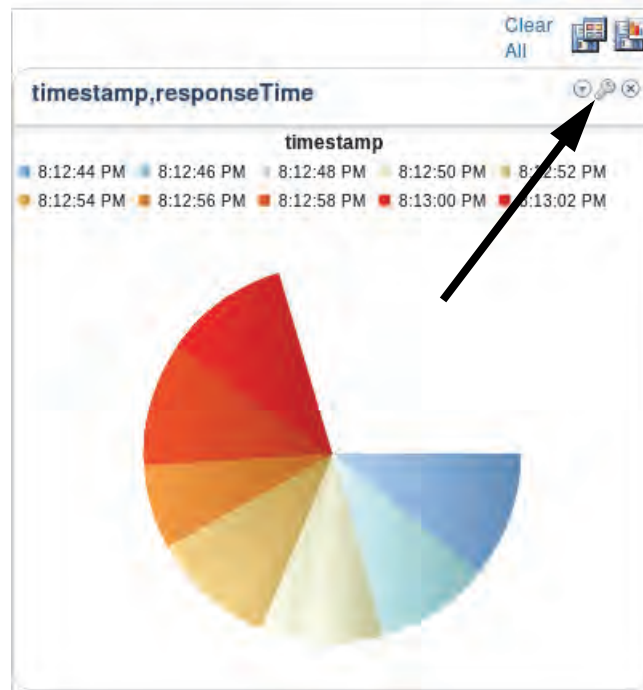
**Granularity**

☐ Hour

m. Click **Close** on the hint about granularity.



- n. Edit the chart by clicking the **Settings** button at the upper right of the chart.



- o. Change the title to Response Time. Change the chart type to **Point Chart**. Click **Render**.

Response Time

Query Visualization

Title Response Time

Chart Type Point Chart

Parameters xaxis timestamp yaxis responseTime

Render Cancel

timestamp

8:42 PM 8:45 PM 8:48 PM 8:51 PM 8:54 PM 8:57 PM 9:00 PM



- p. Click the **Add Chart** button at the top of the chart.

Clear All

**Response Time**

Query Visualization

Title Response Time

Chart Type Point Chart

Parameters xaxis timestamp yaxis responseTime

Render Cancel

responseTime

0 20000 40000 60000 80000 100000 120000 140000 160000

- q. Select **/Dashboards/Web\_Access\_Dashboard** and click **OK**.

Add chart to dashboard

Select Dashboard /Dashboards/Web\_Access\_Dashboard

☐ Preserve fixed time period for all charts in dashboard.

OK Cancel

- r. Close the chart page.

Clear All

**Response Time**

Query Visualization

Title Response Time

Chart Type Point Chart

Parameters xaxis timestamp yaxis responseTime

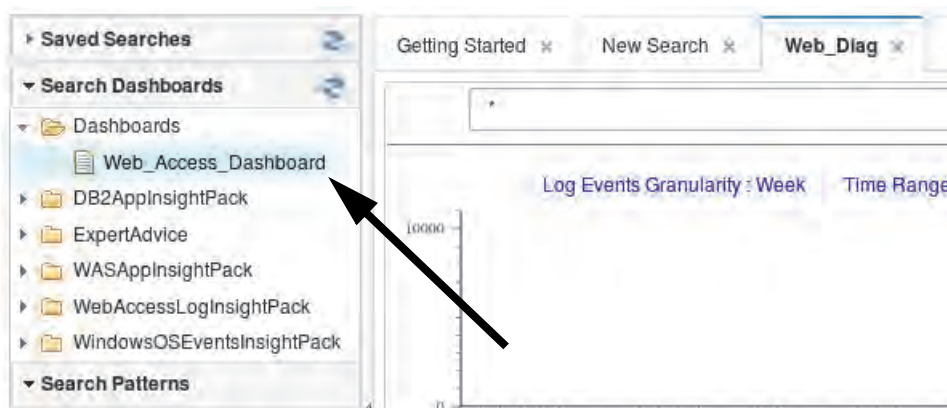
Render Cancel

responseTime

0 20000 40000 60000 80000 100000 120000 140000 160000

17. Test the dashboard:

- a. Expand **Search Dashboards > Dashboards** on the left side of the window. Double-click **Web\_Access\_Dashboards**.



- b. Verify that both charts are on the dashboard page. Leave the page open.

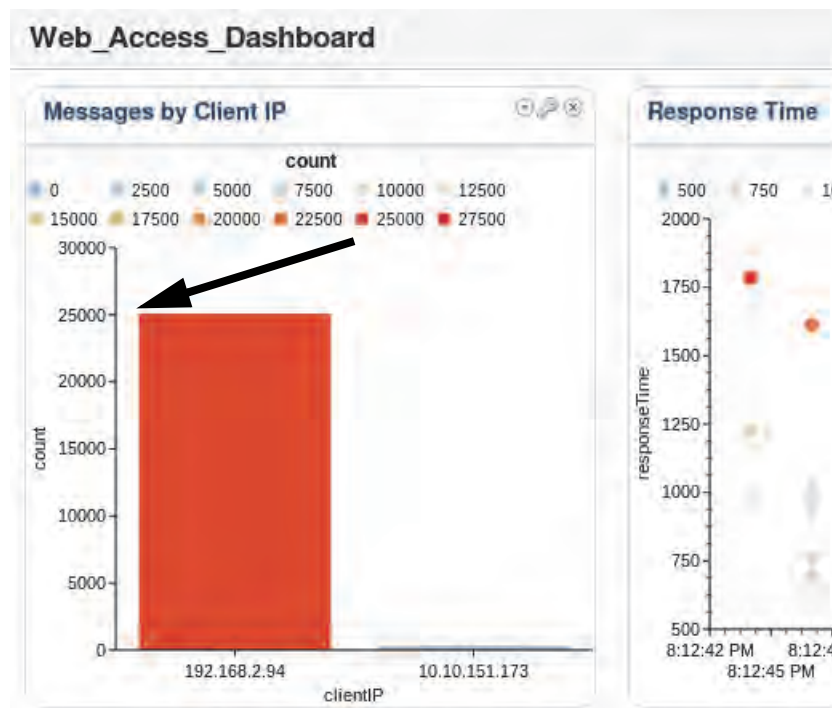


- c. Run the following command to generate more events in the web access log.  
`/software/log_samples/scripts/Web_Logs.sh`

- d. Return to the dashboard page. Click **Actions > Auto-Refresh > Every 1 minutes**.



- e. Wait 1 minute for the page to refresh. Notice that the y-axes of the charts change as more log messages are processed.

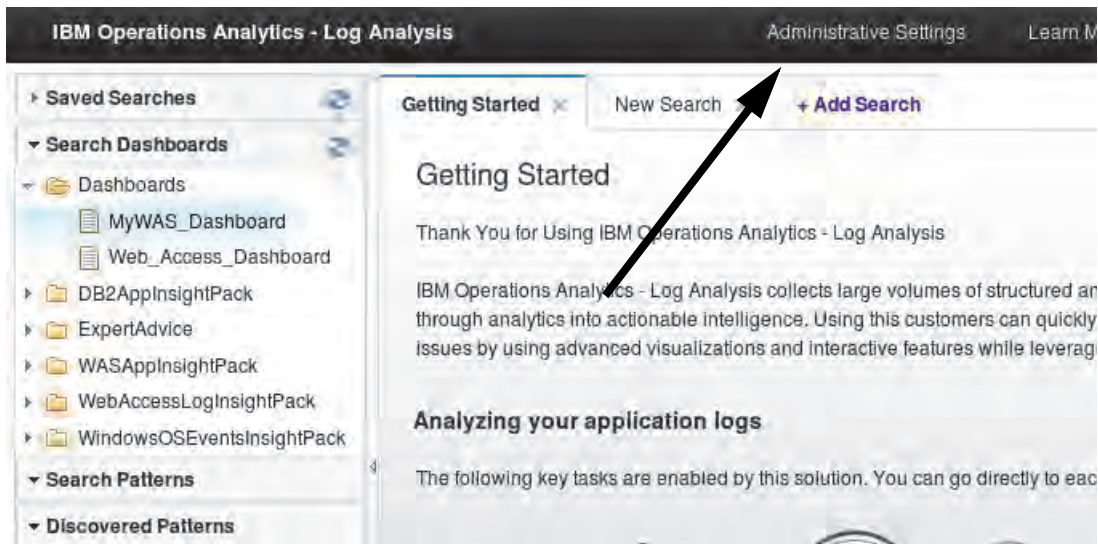


**Hint:** This lab uses different types of log files. Create dashboards for the other log types if you want more practice using the charting and dashboard features.

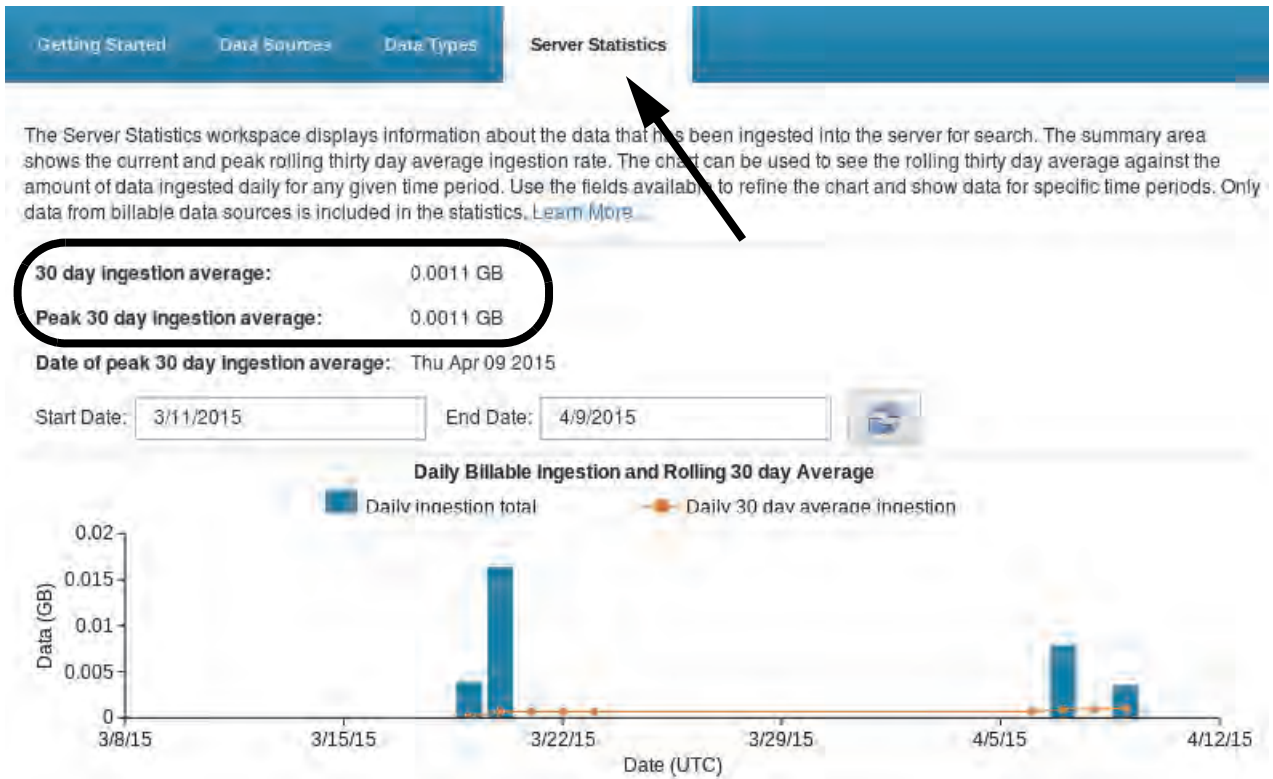
## Exercise 3 Deleting historical data

In this exercise, you view reports about the volume of log data that IBM Operations Analytics Log Analysis processes. You also delete historical data from the system.

1. Use the Server Statistics page to see how much data IBM Operations Analytics Log Analysis has processed.
  - a. Return to the user interface. Click **Administrative Settings**. The administration user interface opens in a new Firefox tab.



- b. Click the **Server Statistics** tab in the administration user interface. This page shows the volume of data that is processed with other product usage measurements. Notice the amount of data that is processed.



2. You can use the `export_statistics` utility to show the volume of data that is processed. Run the following commands to show the volume of data that was processed by data source. Notice the name of each data source and the name of the collection.

```
cd /opt/IBM/LogAnalysis/utilities
./export_statistics -u unityadmin -p unityadmin -t daily
```

Data Source	Collection	Date	Ingested Bytes
Billable	Log Path		Hostname
-----+-----+-----+-----			
-----+-----+-----+-----			
-----			
WAS_SystemErr	WAS_SystemErr	2015-04-07	533498
True	/software/log_samples/WAS_logs/SystemErr.log		
host2.tivoli.edu			
WAS_SystemOut	WAS_SystemOut	2015-04-07	1051931
True	/software/log_samples/WAS_logs/SystemOut.log		
host2.tivoli.edu			
Web_Server	Web_Server	2015-04-07	4749248
True	/software/log_samples/IHS_logs/IHS-access.log		
host2.tivoli.edu			



3. You use a command-line tool to delete data from the system. Use this tool to remove data from the WebSphere SystemErr data source that you previously created.

- a. You must configure the tool to delete only the data you define. Open the `delete.properties` file in a text editor to configure the tool. This example uses `vi`.

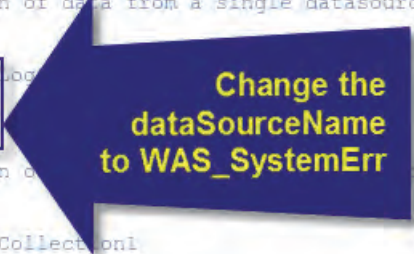
```
cd /opt/IBM/LogAnalysis/utilities/deleteUtility
vi delete.properties
```

- b. Make one change to this file. Change the value of the **dataSourceName** property to **WAS\_SystemErr**. Notice the use case number close to the top of the file. Make the `delete.properties` file look like the following example.

```
# properties and uses these for further delete operation.
[useCase]
useCaseNumber = useCase_1

# useCase_1 concerns the deletion of data from a single datasource. You must provide
# the same parameter.
# Example key - value pair
# dataSourceName = WASSystemOut-Collection1
[useCase_1]
dataSourceName = WAS_SystemErr

# useCase_2 concerns the deletion of data from a single datasource. You must provide
# the same parameter.
# Example key - value pair
# collectionName = WASSystemOut-Collection1
[useCase_2]
collectionName = SOLR_COLLECTION
```



- c. Save and exit the file after you finish editing it.
- d. Run the following command to delete data from the **WAS\_SystemErr** data source.

```
/usr/bin/python2.6 deleteUtility.py unityadmin
```

- e. Look at the log file that shows what data you deleted. The bottom of the log shows the records that were deleted. Use the following command to view the log.

```
tail -20 /opt/IBM/LogAnalysis/logs/DeleteApplication.log
```

```
...
04/09/15 13:43:41:339 UTC [main] INFO - DeleteManager : Datasource retrieved
:WAS_SystemErr
04/09/15 13:43:41:601 UTC [main] INFO - SolrDeleteManager : Solr collection
UnityCollection_07_04_2015_00_00_00_UTC, delete query:
+(_datasource:"WAS_SystemErr")
04/09/15 13:43:42:167 UTC [main] INFO - SolrDeleteManager : Solr collection
UnityCollection_09_04_2015_00_00_00_UTC, delete query:
+(_datasource:"WAS_SystemErr")
04/09/15 13:43:42:192 UTC [main] INFO - CommonUtil : 1 Trying to suspend
thread execution for 1000 miliseconds.
04/09/15 13:43:43:229 UTC [main] INFO - CommonUtil : 1 Successfully executed
POST request.
```



```
04/09/15 13:43:43:229 UTC [main] INFO - DeleteManager : Total number of
records deleted: All records deleted
04/09/15 13:43:43:229 UTC [main] INFO - DeleteManager : Total number of
records : All records deleted
```

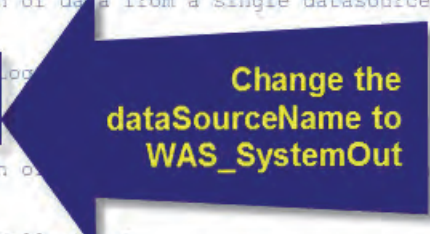
4. Use the tool again to remove data from the WebSphere SystemOut data source that you created.
  - a. Open the delete.properties file again a text editor. This example uses vi.  
vi delete.properties
  - b. Change the value of the **dataSourceName** property to **WAS\_SystemOut**. Make the delete.properties file look like the following example.

```
# properties and uses these for further delete operation.

[useCase]
useCaseNumber = useCase_1

# useCase_1 concerns the deletion of data from a single datasource. You must provide
# a parameter.
# Example key - value pair
# dataSourceName = WASSystemOut-Log
[useCase_1]
dataSourceName = WAS_SystemOut

# useCase_2 concerns the deletion of data from a collection. You must provide
# an onName parameter.
# Example key - value pair
# collectionName = WASSystemOut-Collection1
[useCase_2]
collectionName = SCALA_COLLECTION
```



Change the  
dataSourceName to  
WAS\_SystemOut

- c. Save and exit the file after you finish editing it.
- d. Run the following command to delete data from the WAS\_SystemOut data source.  
/usr/bin/python2.6 deleteUtility.py unityadmin

- e. Look at the log file that shows what data you deleted. The bottom of the log shows the number of records that are deleted. Use the following command to view the log.

```
tail -20 /opt/IBM/LogAnalysis/logs/DeleteApplication.log
```

```
...
```

```
04/09/15 13:52:07:377 UTC [main] INFO - DeleteManager : Datasource retrieved
:WAS_SystemOut
```

```
04/09/15 13:52:07:597 UTC [main] INFO - SolrDeleteManager : Solr collection
UnityCollection_07_04_2015_00_00_00_UTC, delete query:
```

```
+( _datasource: "WAS_SystemOut" )
```

```
04/09/15 13:52:08:046 UTC [main] INFO - SolrDeleteManager : Solr collection
UnityCollection_09_04_2015_00_00_00_UTC, delete query:
```

```
+( _datasource: "WAS_SystemOut" )
```

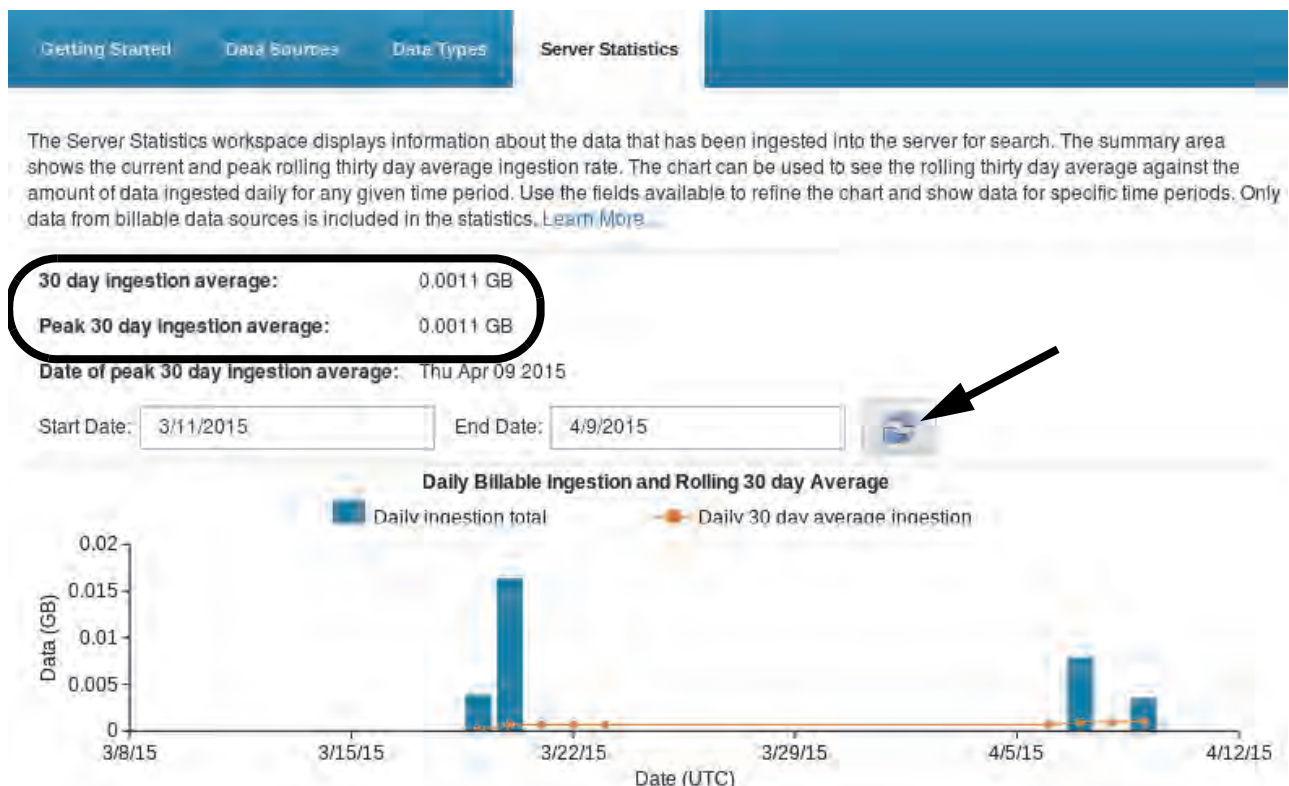
```
04/09/15 13:52:08:072 UTC [main] INFO - CommonUtil : 1 Trying to suspend
thread execution for 1000 miliseconds.
```

```
04/09/15 13:52:09:096 UTC [main] INFO - CommonUtil : 1 Successfully executed
POST request.
```

```
04/09/15 13:52:09:096 UTC [main] INFO - DeleteManager : Total number of
records deleted: All records deleted
```

```
04/09/15 13:52:09:097 UTC [main] INFO - DeleteManager : Total number of
records : All records deleted
```

5. Return to the **Server Statistics** page in the administration user interface. Click the **Refresh** button. Notice that the measurements of how much data is processed do not change, even after you delete historical data.





**Note:** The product usage measurements that you see with the export\_statistics tool are like the measurements on this page. They do not change even after you delete data.



## 2 Common configuration tasks exercises

### Exercise 1 Using the Generic Annotation Insight Pack

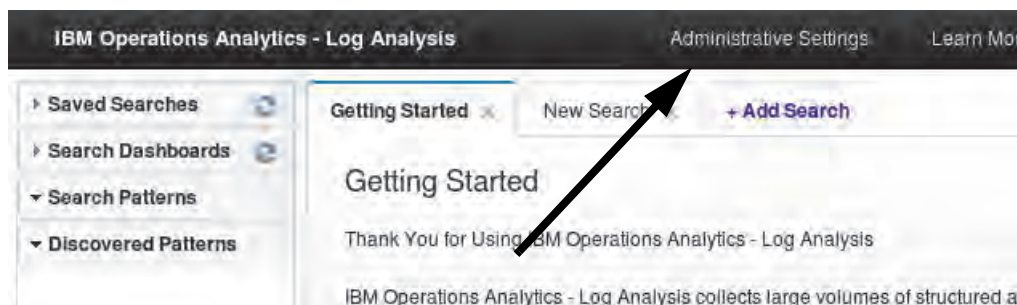
In this exercise, you use the Generic Annotation Insight Pack to analyze an OMNIBus ObjectServer log file.

1. Open the ObjectServer log file sample. Notice the format of the time stamp.
  - a. Run the following command to open the ObjectServer log file sample.

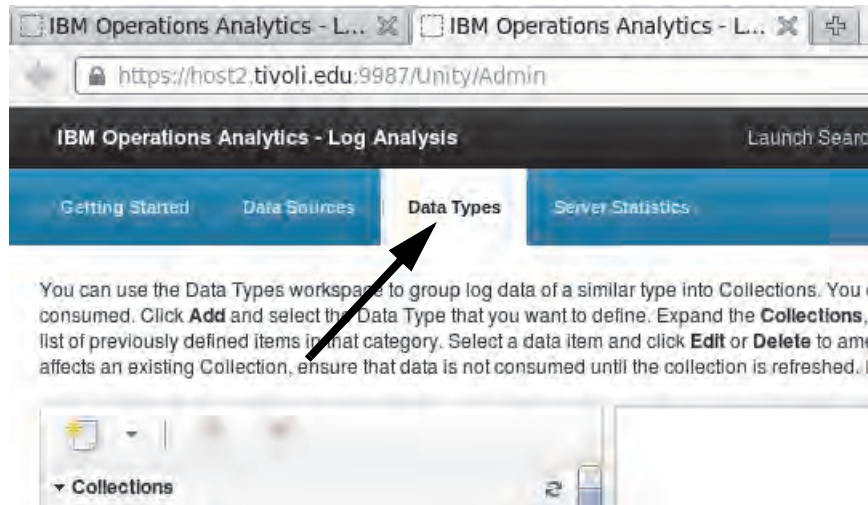
```
more /home/netcool/NYC_AGG_P.log
```
  - b. Look at the format of the time stamp in this log. The format is **yyyy-MM-dd'T'HH:mm:ss**.

```
2015-04-08T14:07:05: Warning: W-STO-103-001: Truncating column SiteState
2015-04-08T14:07:05: Warning: W-STO-103-001: Truncating column SiteState
2015-04-08T14:07:05: Warning: W-STO-103-001: Truncating column SiteState
```
2. Open the administration user interface.
  - a. Double-click the **Firefox** icon on the desktop.
  - b. Browse to the following address.

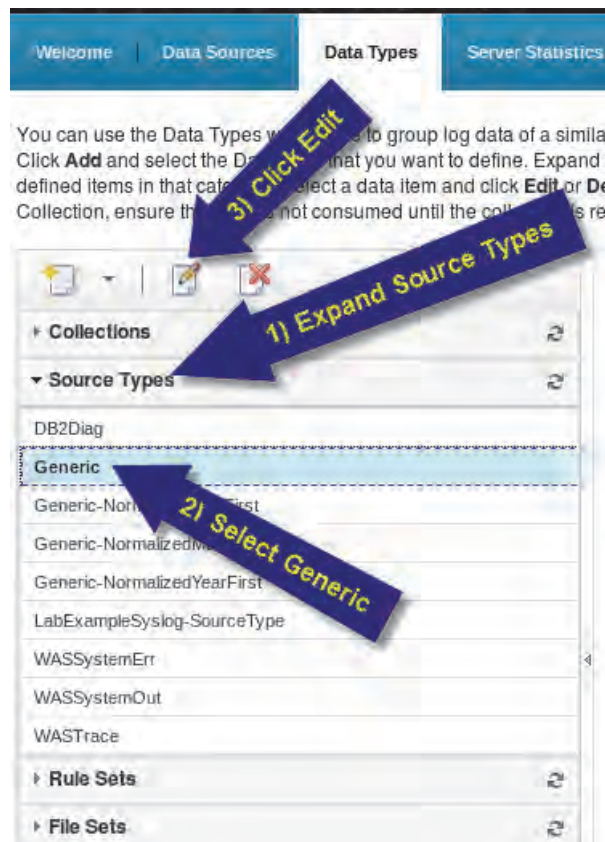
```
https://host2.tivoli.edu:9987/Unity
```
  - c. Log in to the user interface with the user name **unityadmin** and the password **unityadmin**.
  - d. Click **Administrative Settings**. The administration user interface opens in a new Firefox tab.



3. Open the **Generic** source type in the administration user interface. Copy the index configuration from the **Generic** source type.
  - a. Click the **Data Types** tab in the administration user interface. The administration user interface is in the second Firefox tab.



- b. Expand **Source Types**. Select the **Generic** source type and click **Edit**.



- c. Click **View Index Configuration**.

**Edit Generic** x

### Edit Source Type

A Source Type defines how a particular kind of data is split, annotated, and ind

\* Name:

\* Input type:

☒ Enable splitter

☐ Rule set

☐ File set

☒ Enable annotator

☐ Rule set

☐ File set

☐ Deliver data on annotator execution failure

**View Index Configuration**

\* Required

**Close**



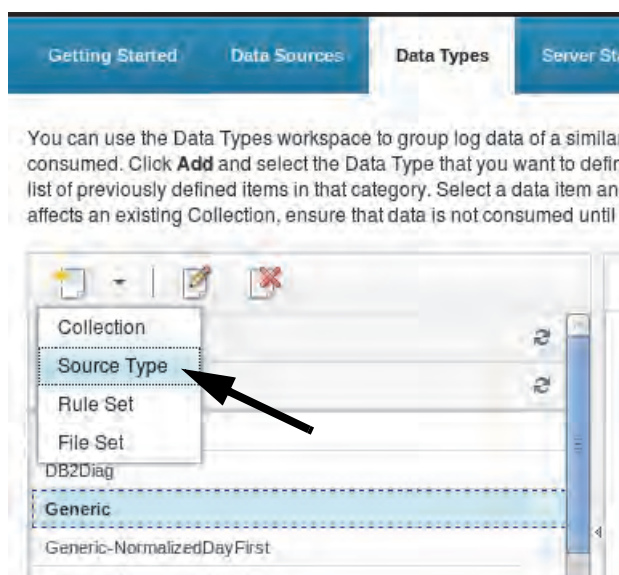
- d. Copy all of the text in the index configuration. Close it when you finish.



4. Create a source type name **Generic-ObjectServer-ST**. Use the values in the following table.

Field	Value
Name	Generic-ObjectServer-ST
Enable splitter	Select this option
Rule set	Generic-dateTime-Split
Enable annotator	Select this option
Rule set	Generic-Annotate
Deliver data on annotator execution failure	Leave this option clear
Edit Index Configuration	Enter the index configuration that you copied in a preceding step. Change the <b>dateFormats</b> field to " <b>yyyy-MM-dd'T'HH:mm:ss</b> "

- a. Click **Add > Source Type**.

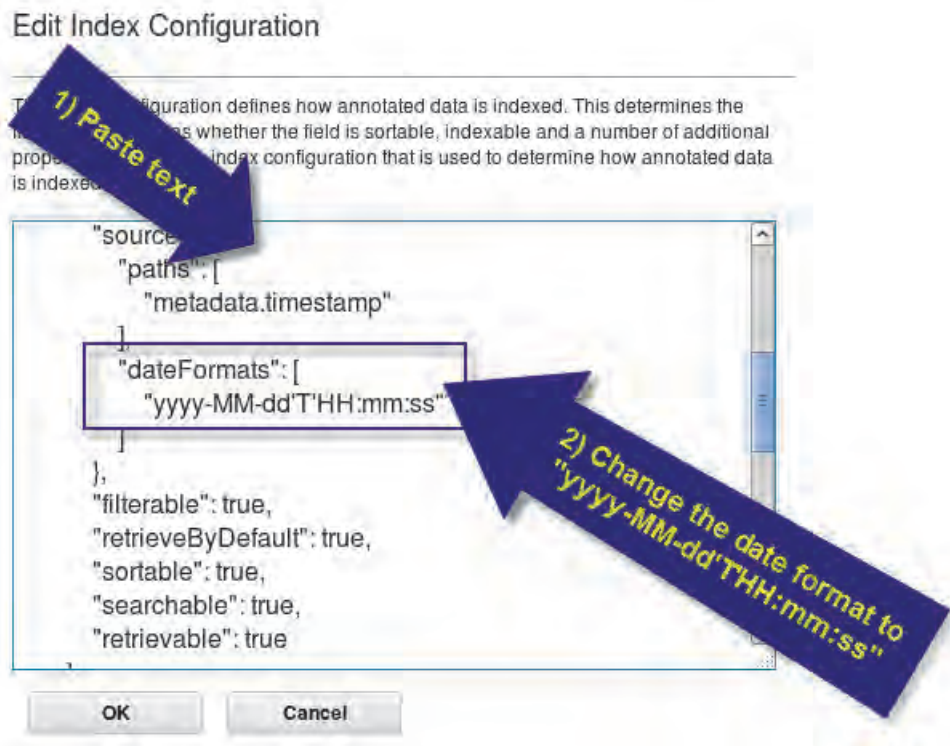


- b. Enter the values from the preceding table. Click **Edit Index Configuration**.

The screenshot shows the 'Add Source Type' dialog box. The title bar says 'Add Source Type'. The main title is 'Add Source Type'. Below the title, there is a description: 'A Source Type defines how a particular kind of data is split, annotated, and indexed for searching. [Learn More](#)'. The dialog contains several fields and checkboxes:

- Name:** A text box containing 'Generic-ObjectServer-ST'.
- Input type:** A dropdown menu.
- Enable splitter:** A checked checkbox.
- Rule set:** A radio button selected, with a dropdown menu showing 'Generic-dateTime-Split'.
- File set:** An unselected radio button.
- Enable annotator:** A checked checkbox.
- Rule set:** A radio button selected, with a dropdown menu showing 'Generic-Annotate'.
- File set:** An unselected radio button.
- Deliver data on annotator execution failure:** An unchecked checkbox.
- Edit Index Configuration:** A button highlighted with a black arrow.
- Required:** A label below the 'Edit Index Configuration' button.
- OK** and **Cancel** buttons at the bottom.

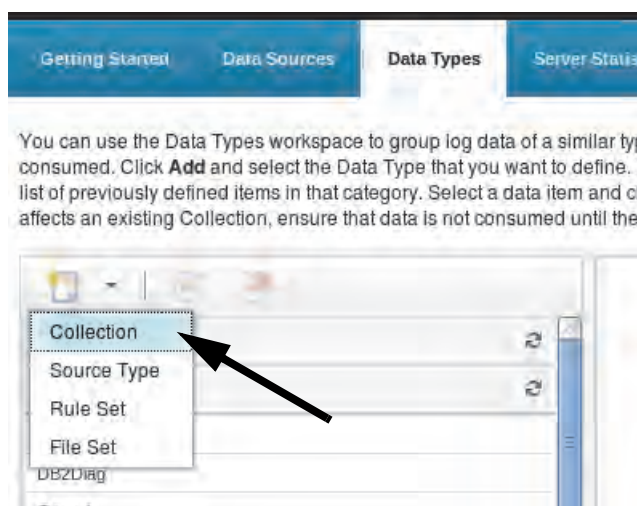
- c. Change the **dateFormats** field to match the format of the OMNIBus log file.
  - i. Remove all of the text from the index configuration.
  - ii. Paste the index configuration that you copied from the **Generic** source type.
  - iii. Change the **dateFormats** field as follows, including the quotation marks:  
"yyyy-MM-dd'T'HH:mm:ss"
  - iv. Click **OK**.



- d. Click **OK** to save the source type. Click **OK** again to confirm.
5. Create a collection named **Generic-ObjectServer-C**. Use the values in the following table.

Field	Value
Name	Generic-ObjectServer-C
Source Type	Generic-ObjectServer-ST

- a. Click **Add > Collection**.



- b. Enter the values from the preceding table. Click **OK**. Click **OK** to confirm.

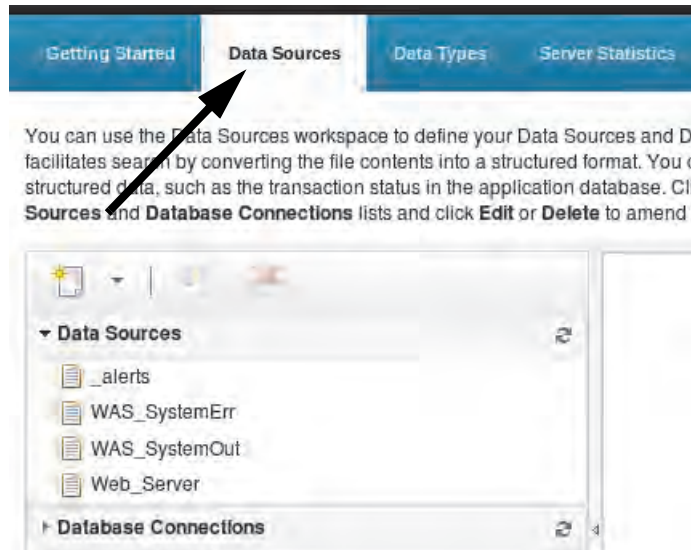


6. Create a data source named ObjectServer-Log. Use the values in the following table to complete the data source wizard.

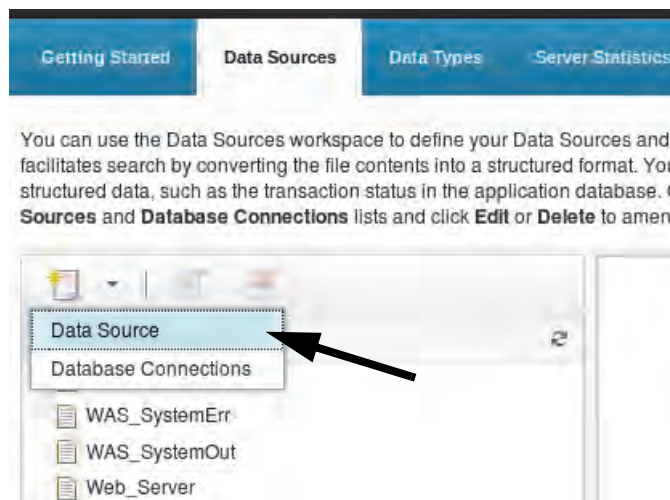
Field	Value
Location	Select Custom
Host name	host2
File Path	/opt/IBM/LogAnalysis/logsources/GAInsightPack/NYC_AGG_P.log
Type	Generic-ObjectServer-ST
Collection	Generic-ObjectServer-C
Name	ObjectServer-Log

Field	Value
Description	This source uses ObjectServer logs
Group	Leave this field blank

- a. Click the **Data Sources** tab in the administration user interface. The administration user interface is in the second Firefox tab.



- b. Click **Add > Data Source**.



- c. Select **Custom**. Enter **host2** as the **Host name** and click **Next**.

## Add Data Source

The screenshot shows the 'Add Data Source' wizard with three tabs: 'Select Location', 'Select Data', and 'Set Attributes'. The 'Select Location' tab is active. It contains a text box with instructions: 'If you want to ingest data into the Log Analytics server, use the wizard to configure a data source. Select Local or Remote file to monitor changes to a file. Select Custom when data is sent to the Log Analytics server from external sources such as a remote log file agent, Logstash, or the data collector client. [Learn More](#).' Below this are three radio buttons: 'Local file', 'Remote file', and 'Custom'. The 'Custom' radio button is selected. Below the radio buttons is a text box labeled 'Host name:' with the value 'host2'. At the bottom left, the word 'Required' is displayed. At the bottom right, there are four buttons: 'Back', 'Next', 'Finish', and 'Cancel'.

- d. Enter `/opt/IBM/LogAnalysis/logsources/GAInsightPack/NYC_AGG_P.log` as the file path.
- e. Select **Generic-ObjectServer-ST** as the **Type**.
- f. Select **Generic-ObjectServer-C** as the **Collection** and click **Next**.

## Add Data Source

The screenshot shows the 'Add Data Source' wizard with three tabs: 'Select Location', 'Select Data', and 'Set Attributes'. The 'Select Data' tab is active. It contains a text box with instructions: 'Enter the location and type of data for this data source. The file path is not validated when you select the custom option. [Learn More](#).' Below this are three text boxes: 'File path:' with the value '/opt/IBM/LogAnalysis/logsources/GAInsightPack/NYC\_AGG\_P.log', 'Type:' with the value 'Generic-ObjectServer-ST', and 'Collection:' with the value 'Generic-ObjectServer-C'. At the bottom left, the word 'Required' is displayed. At the bottom right, there are four buttons: 'Back', 'Next', 'Finish', and 'Cancel'.



- g. Enter **ObjectServer-Log** as the **Name** of the data source.
- h. Enter the **Description** from the preceding table and click **Finish**.

The screenshot shows the 'Add Data Source' dialog box with the 'Set Attributes' tab selected. The dialog has three tabs: 'Select Location', 'Select Data', and 'Set Attributes'. Below the tabs, there is a text area with the instruction: 'Enter a name for the new data source. Optionally, set a description and assign the source to a group of sources. [Learn More](#)'. There are three input fields: 'Name:' with the value 'ObjectServer-Log', 'Description:' with the value 'This source uses ObjectServer logs', and 'Group:' which is empty. Below these fields is a 'Required' label. At the bottom of the dialog are four buttons: 'Back', 'Next', 'Finish', and 'Cancel'.

- i. Click **OK** and **Close** in the confirmation windows.

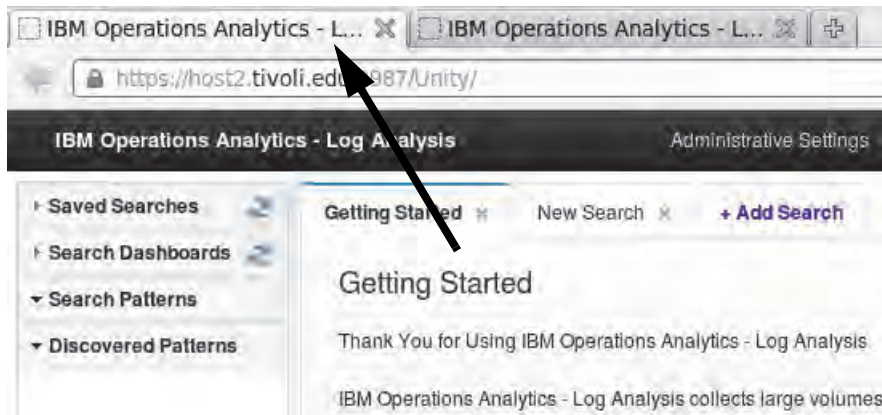
The screenshot shows a 'Confirmation' dialog box with a yellow warning icon. The text inside reads: 'Are you sure that you want to save the data source? After the data source is saved, the following attributes cannot be changed: Location, Name, and Type. Click OK to save. Click Cancel to return to the editor.' At the bottom right are two buttons: 'OK' and 'Cancel'.

The screenshot shows an 'Information' dialog box with a blue information icon. The text inside reads: 'The data source has been successfully configured. For streaming sources, only subsequent changes to the file will be ingested and indexed.' At the bottom right is a 'Close' button.

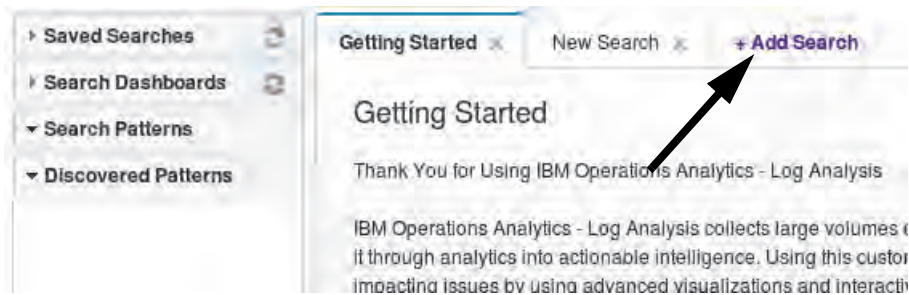
7. Run the following command to manually add some messages to the bottom of the NYC\_AGG\_P.log file. Enter the entire command on one line.

```
cat /home/netcool/NYC_AGG_P.log >>  
/opt/IBM/LogAnalysis/logsources/GAInsightPack/NYC_AGG_P.log
```

8. Return to the user interface by clicking the first Firefox tab.



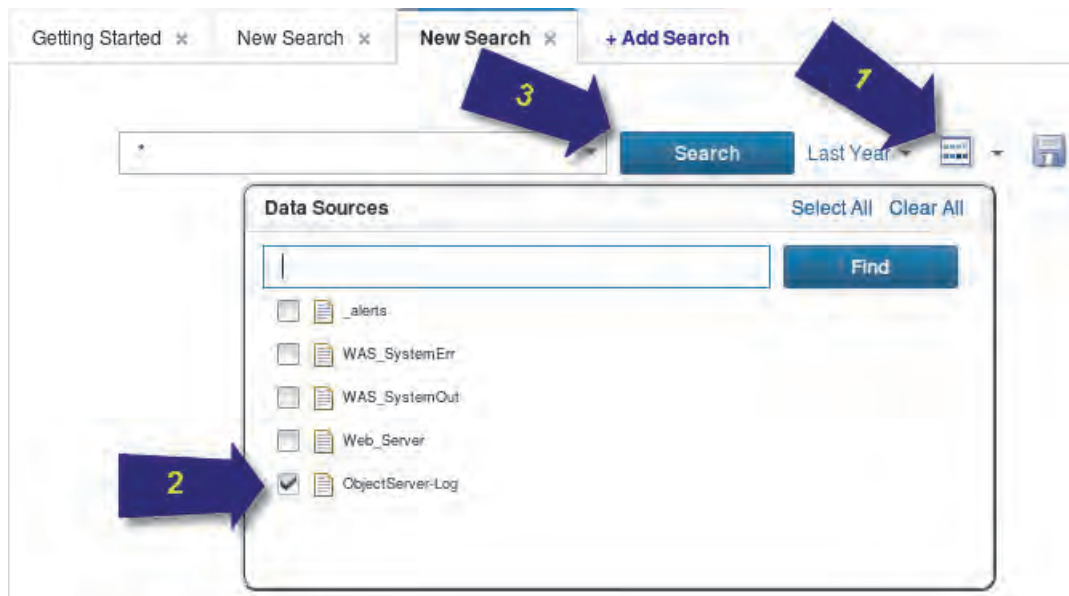
9. Search through the new log and create a saved search based on this log source.
  - a. Click **Add Search**.



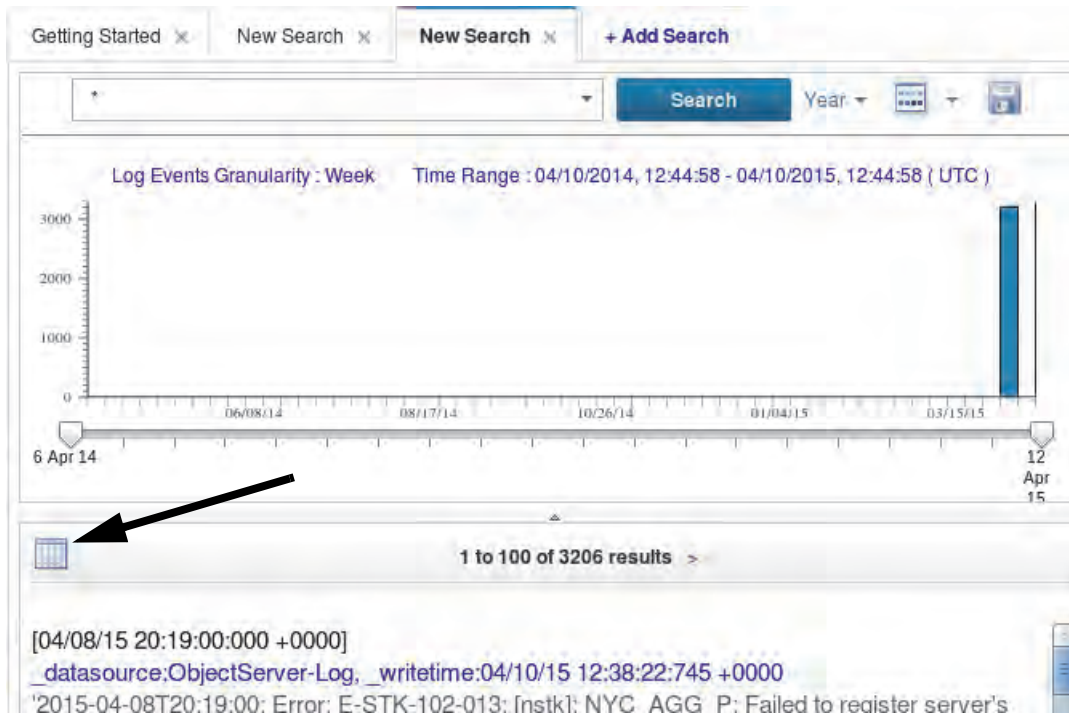
- b. Select **Last Year** as the time filter.



- c. Select only **ObjectServer-Log** as the log source.



- d. Click **Search**. Log events load in to the user interface. Switch to grid view.



- e. Scroll from left to right and notice the columns with log data in them.

timestamp	_writetime	_datasource	logRecord
04/08/15 20:19:00:000	04/10/15 12:38:22:745 +...	ObjectServer-Log	'2015-04-08T20:19:00:000; Error: E-STK-10...
04/08/15 20:18:52:000 +...	04/10/15 12:38:22:745 +...	ObjectServer-Log	'2015-04-08T20:18:52:000; Error: E-STK-10...
04/08/15 20:18:48:000 +...	04/10/15 12:38:22:745 +...	ObjectServer-Log	'2015-04-08T20:18:48:000; Error: E-STK-10...

- f. Look at the **Discovered Patterns** area on the left of the user interface. This area lists text strings that IBM Operations Analytics Log Analysis found recurring in the log file. You can use these patterns to filter log events.
- g. Click **Error** in the Discovered Patterns, and click **Search**. The log file is filtered to show messages that contain the word **Error**.

Search Criteria: `+_concept:"Error"` Search

Log Events Granularity : Week Time Range : 04/10/2014, 12:44:58 - 04/10/2015, 1

timestamp	_writetime	_datasource	logRecord
04/08/15 20:19:00:000 +...	04/10/15 12:38:22:745 +...	ObjectServer-Log	'2015-04-08T...


- h. Take some time to search through the OMNibus log with the IBM Operations Analytics Log Analysis user interface.

## Exercise 2 Using the DSV toolkit

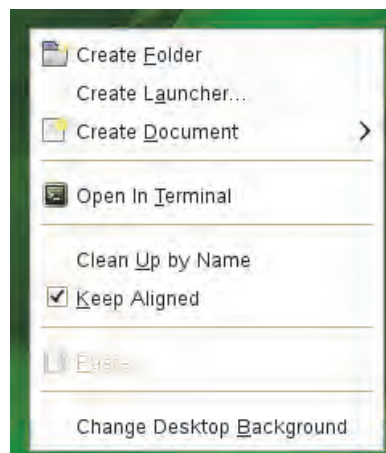
In this exercise, you use the DSV toolkit to create an Insight Pack. The log file that you build the pack for is generated by the UNIX Sar utility.

### Creating the Insight Pack

1. Log in to the host2 virtual machine (host2.tivoli.edu) with the user name **netcool** and the password **object00**, if you are not already logged in.



2. Right-click the desktop and select **Open In Terminal**.



3. Use the following command to verify that you are logged in to the correct host.

```
hostname  
host2.tivoli.edu
```



4. Open a sample of the log file and notice the following details about the format of the log messages. Run the following command to view a sample of the log file.
  - The log file uses a comma as the delimiting character.
  - There are nine fields, or columns, in this log.
  - The first field in this log file is a time stamp, and the format is dd-MMM-yyyy HH:mm:ss.
  - The fields in the log file contain mostly numeric values

```
more /home/netcool/sar_memory/fragments/mem.log
```

```
09-Apr-2015 02:48:05,962588,2159140,69.16,253748,813288,2096472,0,0.00
09-Apr-2015 02:48:17,961292,2160436,69.21,253748,813340,2096472,0,0.00
09-Apr-2015 02:48:30,962068,2159660,69.18,253748,813340,2096472,0,0.00
09-Apr-2015 02:48:42,962068,2159660,69.18,253748,813340,2096472,0,0.00
09-Apr-2015 02:48:54,961680,2160048,69.19,253748,813332,2096472,0,0.00
```

5. Run the following commands to create a new directory for your initial DSV properties file.

```
mkdir /opt/IBM/LogAnalysis/unity_content/DSVToolkit_v1.1.0.3/sarMem
cd /opt/IBM/LogAnalysis/unity_content/DSVToolkit_v1.1.0.3
```

6. Use the `primeProps.py` tool to generate a generic properties file for your log. Run the following command. Enter the entire command on one line.

```
/usr/bin/python2.6 primeProps.py
/opt/IBM/LogAnalysis/unity_content/DSVToolkit_v1.1.0.3/sarMem/
sarMemProperties.properties 9
```

7. Edit the `sarMemProperties.properties` file that you generated. This file is a list of 10 different fields: one field for the overall log record and nine fields for each of the fields in the log file.

- a. Open the `sarMemProperties.properties` file with a text editor. This example uses `vi`.

```
vi sarMem/sarMemProperties.properties
```

- b. Change the value of **scalaHome** to **/opt/IBM/LogAnalysis**.

- c. Change the **aqlModuleName** to **sarmemDSV9Column**.

- d. Change the name of field1 to **timestamp**. Change the **dataType** to **DATE**. Add a **dateFormat** line to the bottom of the field1 paragraph with the time stamp format you saw in the log file sample. Make field1 look like the following example.

```
[field1_indexConfig]
name: timestamp
dataType: DATE
retrievable: true
retrieveByDefault: true
sortable: true
filterable: true
searchable: true
dateFormat: dd-MMM-yyyy HH:mm:ss
```



- e. Change fields 2 - 9 so that their names are `kbMemoryFree`, `kbMemoryUsed`, `memoryUsed`, `kbBuffers`, `kbCached`, `kbSwapFree`, `kbSwapUsed`, `swapUsed`.
- f. Change fields 2,3,5,6,7 and 8 so that the `dataType` is `LONG`.
- g. Change fields 4 and 9 so that the `dataType` is `DOUBLE`.
- h. Edit fields 2 - 9 to match the following properties:  
`retrievable: true`  
`retrieveByDefault: true`  
`sortable: true`  
`filterable: false`  
`searchable: true`
- i. Verify that your `sarMemProperties.properties` file looks exactly like the following example. When you finish, save and exit the file.

```
[SCALA_server]
scalaHome: /opt/IBM/LogAnalysis
```

```
[DSV_file]
delimiter: ,
totalColumns: 9
aqlModuleName: sarmemDSV9Column
version: 1.0.0.0
```

```
[field0_indexConfig]
name: logRecord
dataType: TEXT
retrievable: true
retrieveByDefault: true
sortable: false
filterable: false
searchable: true
path_1: content.text
combine: FIRST
```

```
[field1_indexConfig]
name: timestamp
dataType: DATE
retrievable: true
retrieveByDefault: true
sortable: true
filterable: true
searchable: true
dateFormat: dd-MMM-yyyy HH:mm:ss
```

```
[field2_indexConfig]
```

```
name: kbMemoryFree
dataType: LONG
retrievable: true
retrieveByDefault: true
sortable: true
filterable: false
searchable: true
```

```
[field3_indexConfig]
name: kbMemoryUsed
dataType: LONG
retrievable: true
retrieveByDefault: true
sortable: true
filterable: false
searchable: true
```

```
[field4_indexConfig]
name: memoryUsed
dataType: DOUBLE
retrievable: true
retrieveByDefault: true
sortable: true
filterable: false
searchable: true
```

```
[field5_indexConfig]
name: kbBuffers
dataType: LONG
retrievable: true
retrieveByDefault: true
sortable: true
filterable: false
searchable: true
```

```
[field6_indexConfig]
name: kbCached
dataType: LONG
retrievable: true
retrieveByDefault: true
sortable: true
filterable: false
searchable: true
```

```
[field7_indexConfig]
name: kbSwapFree
dataType: LONG
retrievable: true
```

```
retrieveByDefault: true
sortable: true
filterable: false
searchable: true
```

```
[field8_indexConfig]
name: kbSwapUsed
dataType: LONG
retrievable: true
retrieveByDefault: true
sortable: true
filterable: false
searchable: true
```

```
[field9_indexConfig]
name: swapUsed
dataType: DOUBLE
retrievable: true
retrieveByDefault: true
sortable: true
filterable: false
searchable: true
```

8. Create the insight pack with the `dsvGen.py` tool. When you run this tool, it generates an Insight Pack based on the configuration you added in the properties file.

- a. Verify that you are in the `dsvGen.py` tool directory.

```
cd /opt/IBM/LogAnalysis/unity_content/DSVToolkit_v1.1.0.3/
```

- b. Run the following command to create the Insight Pack. Enter the command on a single line.

```
/usr/bin/python2.6 dsvGen.py sarMem/sarMemProperties.properties -o -d -u
unityadmin -p unityadmin
```

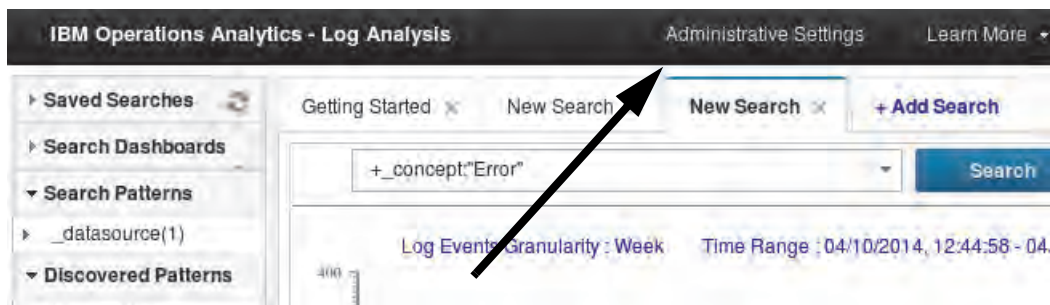
- c. Enter **y** when you are prompted to restart the log file agent.

```
Deploying/undeploying Log File Adapter configuration files requires the LFA
to to stopped before and restarted after. Do you want to continue (y/n)?
```

```
y
```

## Testing the Insight Pack

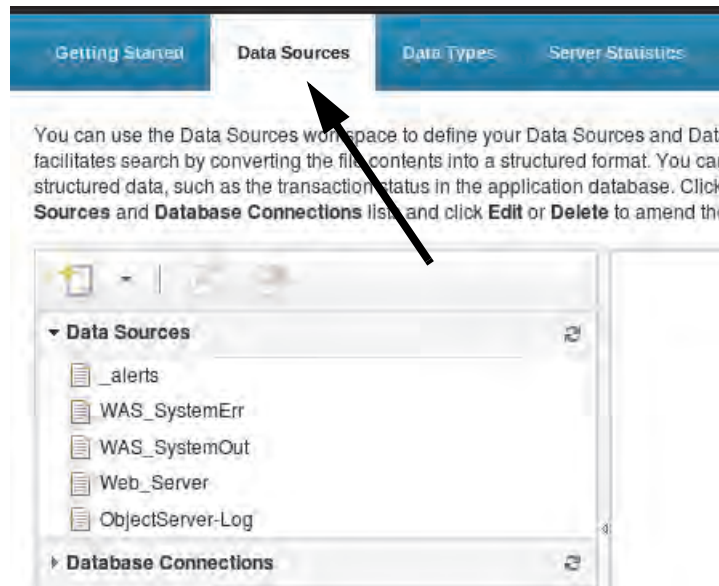
- Open the administration user interface. Click **Administrative Settings**. The administration user interface opens in a new Firefox tab.



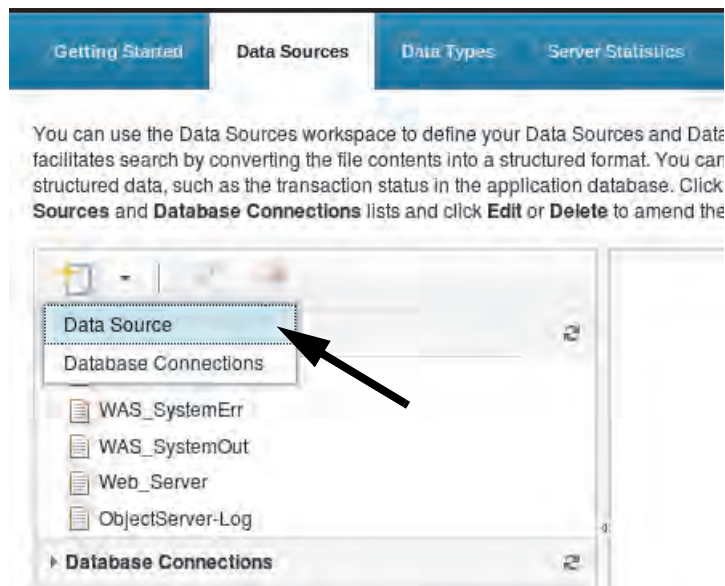
- Create a data source named **Sar\_Memory**. Use the values in the following table to complete the data source wizard.

Field	Value
Location	Select Custom
Host name	host2
File Path	/opt/IBM/LogAnalysis/logsources/ sarmemDSV9ColumnInsightPack/mem.log
Type	sarmemDSV9Column
Collection	sarmemDSV9Column-Collection
Name	Sar_Memory
Description	This source uses sysstat measurements
Group	Leave this field blank

- a. Click the **Data Sources** tab in the administration user interface. The administration user interface is in the second Firefox tab.



- b. Click **Add > Data Source**.



- c. Select **Custom**. Enter **host2** as the **Host name** and click **Next**.

Add Data Source

Select Location   Select Data   Set Attributes

If you want to ingest data into the Log Analytics server, use the wizard to configure a data source. Select Local or Remote file to monitor changes to a file. Select Custom when data is sent to the Log Analytics server from external sources such as a remote log file agent, Logstash, or the data collector client. [Learn More](#)

☐ Local file  
☐ Remote file  
☒ Custom

Host name:

Required

Back   Next   Finish   Cancel

- d. Enter `/opt/IBM/LogAnalysis/logsources/sarmemDSV9ColumnInsightPack/mem.log` as the file path.



**Important:** Ensure that you enter the path correctly.

- e. Select **sarmemDSV9Column** as the **Type**.



- f. Select **sarmemDSV9Column-Collection** as the **Collection** and click **Next**.

The screenshot shows the 'Add Data Source' dialog with the 'Select Data' tab selected. The 'File path' field contains 'sources/sarmemDSV9ColumnInsightPack/mem.log'. The 'Type' dropdown is set to 'sarmemDSV9Column'. The 'Collection' dropdown is set to 'sarmemDSV9Column-Collection'. The 'Required' checkbox is unchecked. The 'Next' button is highlighted.

Add Data Source

Select Location **Select Data** Set Attributes

Enter the location and type of data for this data source. The file path is not validated when you select the custom option. [Learn More](#)

File path:

Type:

Collection:

☐ Required

- g. Enter **Sar\_Memory** as the **Name** of the data source.
- h. Enter the **Description** from the preceding table and click **Finish**.

The screenshot shows the 'Add Data Source' dialog with the 'Set Attributes' tab selected. The 'Name' field contains 'Sar\_Memory'. The 'Description' field contains 'This source uses sysstat measurements'. The 'Group' dropdown is empty. The 'Next' button is highlighted.

Add Data Source

Select Location Select Data **Set Attributes**

Enter a name for the new data source. Optionally, set a description and assign the source to a group of sources. [Learn More](#)

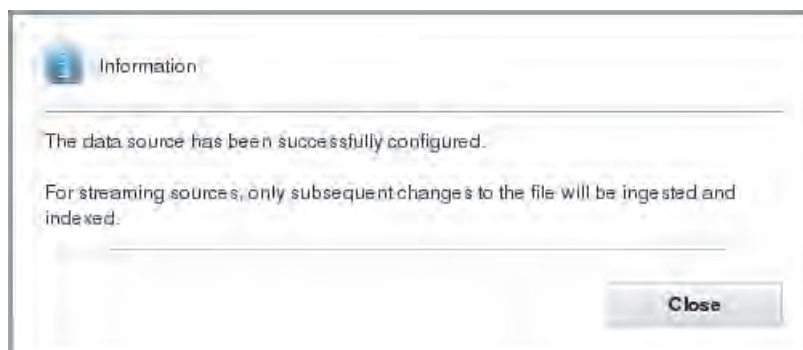

Name:

Description:

Group:

☐ Required

- i. Click **OK** and **Close** in the confirmation windows.

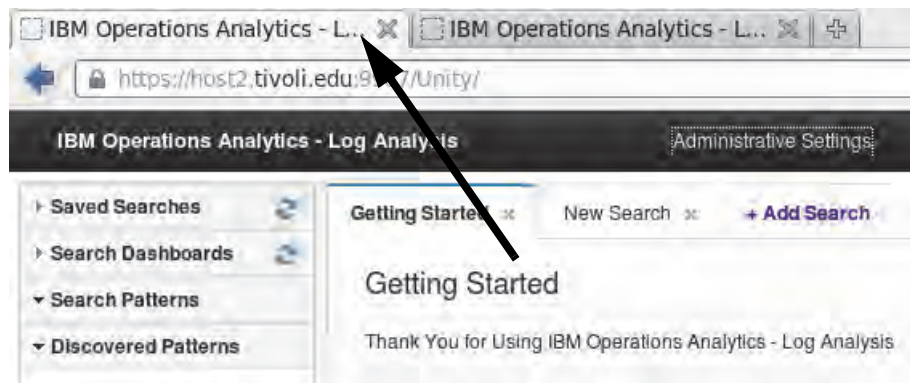


The sar memory log file for this lab is generated by a script.

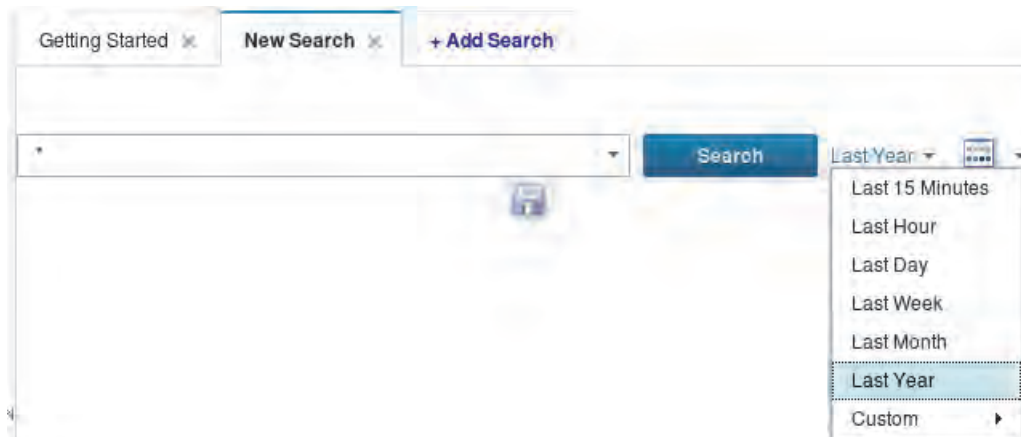
11. Return to the terminal window. Run the following command to generate events in the mem.log file.

```
/home/netcool/sar_memory/fragments/Generate_sarMemory.sh
```

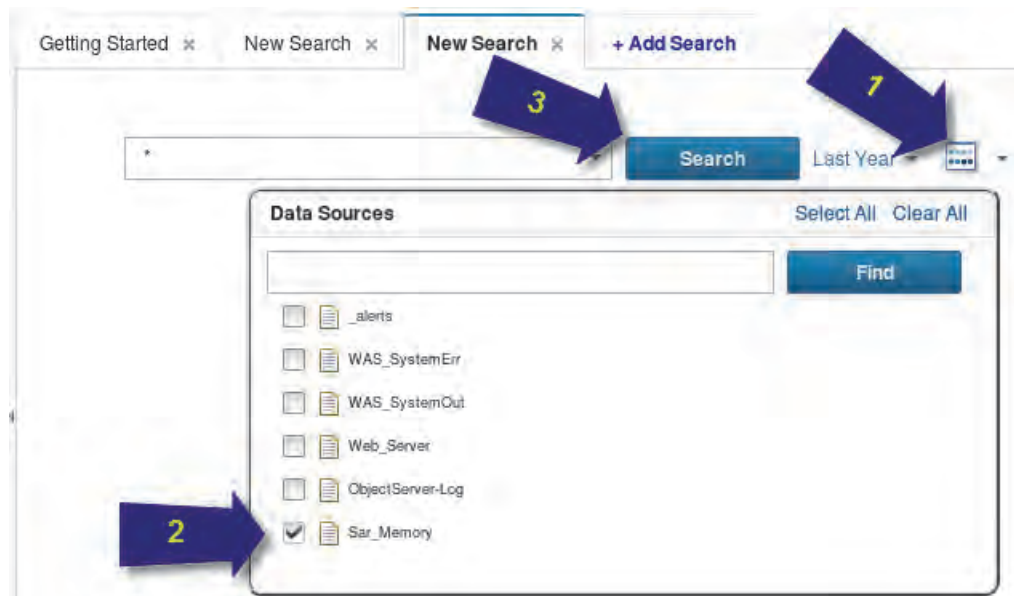
12. Return to the search interface by clicking the first Firefox tab.



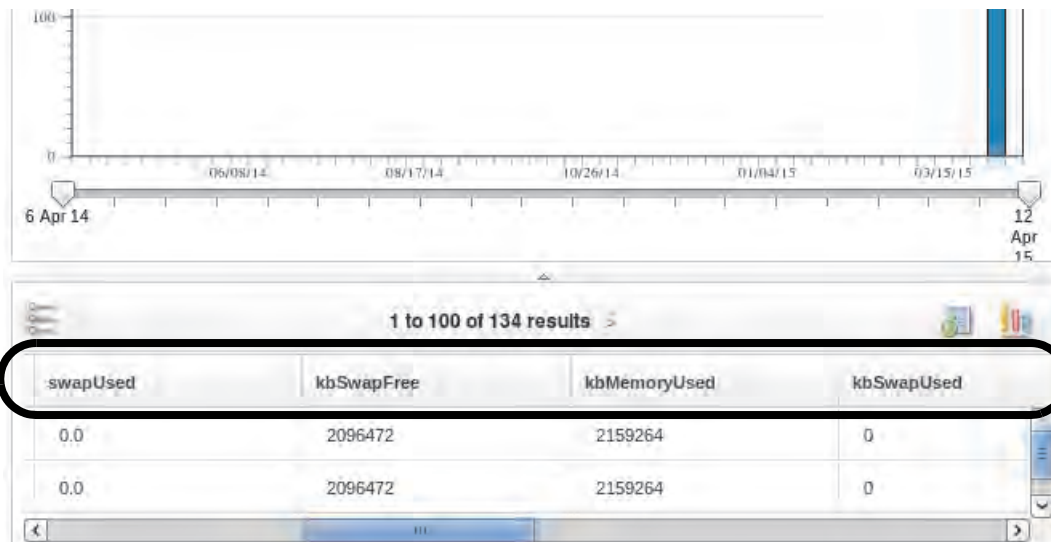
13. Search through the new log and verify that the log file was processed successfully.
- a. Click the **Add Search** tab. Select **Last Year** as the time filter.



- b. Select only **Sar\_Memory** as the data source. Click **Search**.



- c. Log events load in to the user interface. Click the **Grid View** button.



- d. Scroll left and right to view the columns. Verify that each of the fields you configured for the log file are now columns in the user interface. Users can now search through this log, create charts and dashboards, and so on.



## 3 Troubleshooting exercises

There are no student exercises for this unit.



## 4 Alerts exercises

In this exercise, you configure the product to generate alerts from conditions in log files.

### Exercise 1 Creating alert actions

In this exercise, you create two alert actions. One action adds an entry to a log file and the other action sends an email notification.

#### Index alert action

When a condition is detected in a log file, the product automatically indexes the occurrence of the alert. Users can query the indexed alerts in the search interface. This action is enabled by default.

1. Look at the index alert action with the `alerts.sh` utility.
  - a. Change to the alerts subdirectory:  

```
cd /opt/IBM/LogAnalysis/utilities/alerts
```
  - b. Run the following command to show all alert actions.  

```
./alerts.sh -getAlertAction
```

The index alert action is enabled by default. You use this action in alert conditions later in this exercise.

```
Name: index
Label: Index
Description: Alert action implementation that indexes triggered alerts
Template: index
Enabled: true
```

#### Log alert action

2. Create an alert action that adds a message to a log file.
  - a. Verify that you are in the alerts subdirectory.  

```
cd /opt/IBM/LogAnalysis/utilities/alerts
```



- b. Make a copy of the `logAlertAction.json` file. Name the copy `LAB_logAlertAction.json`.
- ```
cp logAlertAction.json LAB_logAlertAction.json
```

- c. Open the `LAB_logAlertAction.json` file in a text editor. Edit the following three fields in the file. Make the `LAB_logAlertAction.json` file match the following example.

```
vi LAB_logAlertAction.json
```

```
{
  "name": "LAB-log-action",
  "description": "This action adds a message to a log file",
  "alertActionTemplateName": "log",
  "parameterValues": {
    "filePath": "/tmp/LAB_alert.log"
  }
}
```

- d. Save and close the file when you are done.

3. Use the `alerts.sh` utility to create the new alert action.

- a. Run the following command to create the alert action.

```
./alerts.sh -createAlertAction LAB_logAlertAction.json
```

- b. Run the following command to verify that the alert action was created.

```
./alerts.sh -getAlertAction
```

```
...
```

```
Name: LAB-log-action
```

```
Description: This action adds a message to a log file
```

```
Template: log
```

```
Parameter Values:
```

```
    filePath: /tmp/LAB_alert.log
```

```
Enabled: true
```

## Email alert action

4. Create an alert action that sends an email to the **netcool** user.

- a. Make a copy of the `emailAlertAction.json` file. Name the copy `LAB_emailAlertAction.json`.

```
cp emailAlertAction.json LAB_emailAlertAction.json
```

- b. Open the `LAB_emailAlertAction.json` file in a text editor. Edit the following six fields in the file. Make the `LAB_emailAlertAction.json` file match the following example.

```
vi LAB_emailAlertAction.json
```

```
{
  "name": "LAB-email-action",
  "description": "This action sends an e-mail to netcool",
```

```

    "alertActionTemplateName": "email",
    "parameterValues": {
      "smtpMailServer": "localhost",
      "secure": false,
      "from": "from@ibm.com",
      "to": ["netcool"],
      "cc": [],
      "bcc": [],
      "subjectPrefix": "An alert from a log record requires your attention",
      "header": "Dear User,",
      "footer": "*** This is a system generated e-mail, please do not reply to
this e-mail ***\n",
      "attachLogRecordAnnotations": true,
    }
  }
}

```

c. Save and close the file when you are done.

5. Use the `alerts.sh` utility to create the new alert action.

a. Run the following command to create the alert action.

```
./alerts.sh -createAlertAction LAB_emailAlertAction.json
```

b. Run the following command to verify that the alert action was created.

```
./alerts.sh -getAlertAction
```

```
...
```

```
Name: LAB-email-action
```

```
Description: This action sends an e-mail to netcool
```

```
Template: email
```

```
Parameter Values:
```

```
  to: [netcool]
```

```
  footer: *** This is a system generated e-mail, please do not reply to
this e-mail ***
```

```
  attachLogRecordAnnotations: true
```

```
  secure: false
```

```
  bcc: []
```

```
  subjectPrefix: An alert from a log record requires your attention
```

```
  from: from@ibm.com
```

```
  header: Dear User,
```

```
  cc: []
```

```
  smtpMailServer: localhost
```

```
Enabled: true
```

## Exercise 2 Creating base conditions

In this exercise, you create base conditions. These conditions use queries to detect text patterns in data sources.

### WebSphere SystemOut log base condition

1. Create a base condition that queries the WAS\_SystemOut data source and uses all of the actions you configured in the preceding exercise. Query for the text `java.sql.SQLException` in the WebSphere SystemOut log.

- a. Verify that you are in the alerts subdirectory.

```
cd /opt/IBM/LogAnalysis/utilities/alerts
```

- b. Make a copy of the `queryBaseCondition.json` file. Name the copy `WASOUT_queryBaseCondition.json`.

```
cp queryBaseCondition.json WASOUT_queryBaseCondition.json
```

- c. Open the `WASOUT_queryBaseCondition.json` file in a text editor. Edit the following five fields in the file. Make the `WASOUT_queryBaseCondition.json` file match the following example.

```
vi WASOUT_queryBaseCondition.json
```

```
{
  "name": "WASOUT-base-condition",
  "description": "This condition detects java sql exceptions",
  "baseConditionTemplateName": "query",
  "datasourceName": "WAS_SystemOut",
  "parameterValues": { "query" :
    "+javaException:java.sql.SQLException"},
  "actions": ["index","LAB-log-action","LAB-email-action"]
}
```

- d. Save and close the file when you are done.

2. Use the `alerts.sh` utility to create the new condition.

- a. Run the following command to create the condition.

```
./alerts.sh -createBaseCondition WASOUT_queryBaseCondition.json
```

- b. Run the following command to verify that the condition was created.

```
./alerts.sh -getBaseCondition
...
Name: WASOUT-base-condition
Description: This condition detects java sql exceptions
Template: query
Datasource: WAS_SystemOut
Parameter Values:
  query: +javaException:java.sql.SQLException
Actions: [LAB-email-action, LAB-log-action, index]
Enabled: true
```

## WebSphere SystemErr log base condition

3. Create a base condition that queries the WAS\_SystemErr data source and uses all of the actions that you configured earlier. Query for the text `java.sql.SQLException` in the WebSphere SystemErr log. This condition is identical to the `WASOUT_queryBaseCondition.json` you just created, except that it uses a different data source (`WAS_SystemErr`).

- a. Verify that you are in the alerts subdirectory.

```
cd /opt/IBM/LogAnalysis/utilities/alerts
```

- b. Make a copy of the `WASOUT_queryBaseCondition.json` file. Name the copy `WASERR_queryBaseCondition.json`.

```
cp WASOUT_queryBaseCondition.json WASERR_queryBaseCondition.json
```

- c. Open the `WASERR_queryBaseCondition.json` file in a text editor. Edit the following two fields in the file. Make the `WASERR_queryBaseCondition.json` file match the following example.

```
vi WASERR_queryBaseCondition.json
```

```
{
  "name": "WASERR-base-condition",
  "description": "This condition detects java sql exceptions",
  "baseConditionTemplateName": "query",
  "datasourceName": "WAS_SystemErr",
  "parameterValues": { "query" : "+javaException:java.sql.SQLException" },
  "actions": ["index", "LAB-log-action", "LAB-email-action"]
}
```

- d. Save and close the file when you are done.

4. Use the `alerts.sh` utility to create the new condition.

- a. Run the following command to create the condition.

```
./alerts.sh -createBaseCondition WASERR_queryBaseCondition.json
```

- b. Run the following command to verify that the condition was created.

```
./alerts.sh -getBaseCondition
...
Name: WASERR-base-condition
Description: This condition detects java sql exceptions
Template: query
Datasource: WAS_SystemErr
Parameter Values:
  query: +javaException:java.sql.SQLException
Actions: [LAB-email-action, LAB-log-action, index]
Enabled: true
```

## Web access log base condition

5. Create a base condition that queries the Web\_Server data source and uses all of the actions you configured earlier. Query for server response times that are over 30,000 microseconds.

- a. Verify that you are in the alerts subdirectory.

```
cd /opt/IBM/LogAnalysis/utilities/alerts
```

- b. Make a copy of the queryBaseCondition.json file. Name the copy WEB\_queryBaseCondition.json.

```
cp queryBaseCondition.json WEB_queryBaseCondition.json
```

- c. Open the WEB\_queryBaseCondition.json file in a text editor. Edit the following five fields in the file. Make the WEB\_queryBaseCondition.json file match the following example.

```
vi WEB_queryBaseCondition.json

{
  "name": "WEB-base-condition",
  "description": "This condition detects high response times",
  "baseConditionTemplateName": "query",
  "datasourceName": "Web_Server",
  "parameterValues": { "query" : "+responseTime:[30000 TO *]"},
  "actions": ["index", "LAB-log-action", "LAB-email-action"]
}
```

- d. Save and close the file when you are done.

6. Use the alerts.sh utility to create the new condition.

- a. Run the following command to create the condition.

```
./alerts.sh -createBaseCondition WEB_queryBaseCondition.json
```

- b. Run the following command to verify that the condition was created.

```
./alerts.sh -getBaseCondition
...
Name: WEB-base-condition
Description: This condition detects high response times
Template: query
Datasource: Web_Server
Parameter Values:
  query: +responseTime:[30000 TO *]
Actions: [LAB-email-action, LAB-log-action, index]
Enabled: true
```

## Exercise 3 Testing base conditions

In this exercise, you test the alert actions and conditions you created in a preceding exercise.

1. Run the following commands to add more messages to the WAS and web server log files.

```
/software/log_samples/scripts/WAS_Logs.sh
/software/log_samples/scripts/Web_Logs.sh
```

2. Open the /tmp/LAB\_alert.log file. Notice the alerts in this log file that were triggered by the WAS and the Web\_Server conditions.

```
more /tmp/LAB_alert.log
```

```
{ "conditionName": "WASOUT-base-condition", "conditionType": "base", "datasources": [
  "WAS_SystemOut"], "triggeringInput": { "datasourceHostname": [ "host2.tivoli.edu" ], "
  threadID": [ "00000000" ], "packageName": [ "com.ibm.tivoli.ncw.datasource" ], "excepti
  onClassName": [ "ConnectionPool" ], "severity": [ "E" ], "exceptionMethodName": [ "create
  Connection" ], "timestamp": [ 1426946599768 ], "_writetime": [ 1429035550400 ]
  ...
  { "conditionName": "WASERR-base-condition", "conditionType": "base", "datasources": [
    "WAS_SystemErr"], "triggeringInput": { "datasourceHostname": [ "host2.tivoli.edu" ], "
    threadID": [ "0000000e" ], "packageName": [ "com.ibm.tivoli.ncw.ncosvmm.sql" ], "except
    ionClassName": [ "FailBackDataSourceSelector" ], "severity": [ "R" ], "exceptionMethodN
    ame": [ "getDataSource" ], "timestamp": [ 1427037406733 ], "_writetime": [ 1429035560583 ]
    ...
    { "conditionName": "WEB-base-condition", "conditionType": "base", "datasources": [ "We
    b_Server" ], "triggeringInput": { "bytes": [ "3409" ], "agent": [ "curl\7.21.3
    (x86_64-unknown-linux-gnu) libcurl\7.21.3 OpenSSL\1.0.0
    zlib\1.2.3" ], "timestamp": [ 1426794280000 ], "response": [ "200" ], "webServerType": [ "
    Apache\IHS" ], "referrer": [ "-" ], "ident": [ "-" ], "verb": [ "GET" ], "request": [ "GET
    \daytrader\app?action=sell&holdingID=5015 HTTP\1.1" ]
```



3. Check the mail for the **netcool** user. Notice the email notifications that were triggered by the WAS and the Web\_Server conditions.

- a. Run the following command to check the mail for the **netcool** user.

```
mail
```

```
>U 1 from@ibm.com          Tue Apr 14 18:19 251/12689 "An alert from a log
record requires your attention [Alert condition WASOUT-base-condition
trigger]"
  U 2 from@ibm.com          Tue Apr 14 18:19 253/12699 "An alert from a log
record requires your attention [Alert condition WASOUT-base-condition
trigger]"
  U 3 from@ibm.com          Tue Apr 14 18:19 234/12740 "An alert from a log
record requires your attention [Alert condition WASERR-base-condition
trigger]"
  U 4 from@ibm.com          Tue Apr 14 18:19 234/12741 "An alert from a log
record requires your attention [Alert condition WASERR-base-condition
trigger]"
  U 5 from@ibm.com          Tue Apr 14 18:19 234/12739 "An alert from a log
record requires your attention [Alert condition WASERR-base-condition
trigger]"
  U 6 from@ibm.com          Tue Apr 14 18:20 76/2163  "An alert from a log
record requires your attention [Alert condition WEB-base-condition
triggered]"
  U 7 from@ibm.com          Tue Apr 14 18:20 76/2192  "An alert from a log
record requires your attention [Alert condition WEB-base-condition
triggered]"
  U 8 from@ibm.com          Tue Apr 14 18:20 76/2599  "An alert from a log
record requires your attention [Alert condition WEB-base-condition
triggered]"
...
```

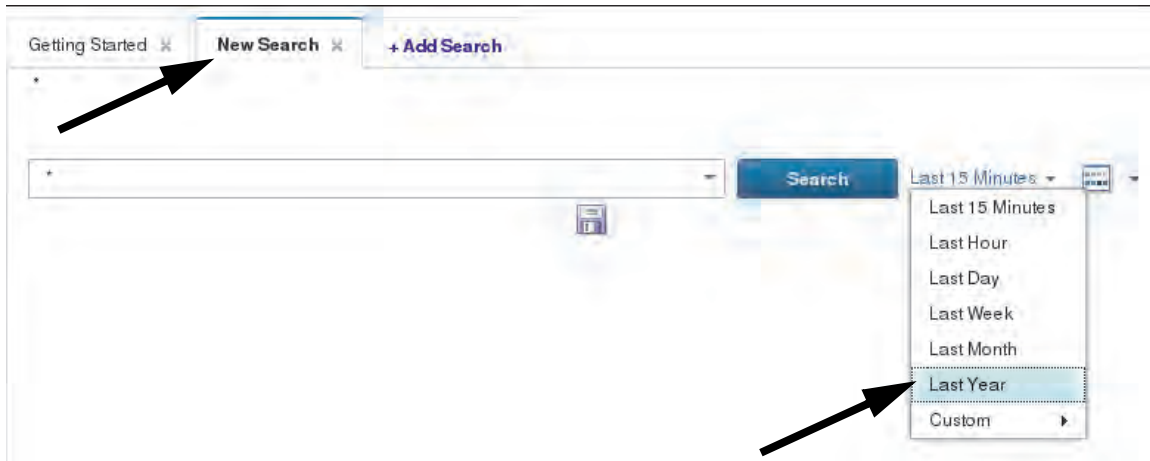
- b. Run a command like the following example to delete all mail messages.

```
& del 1-12
```

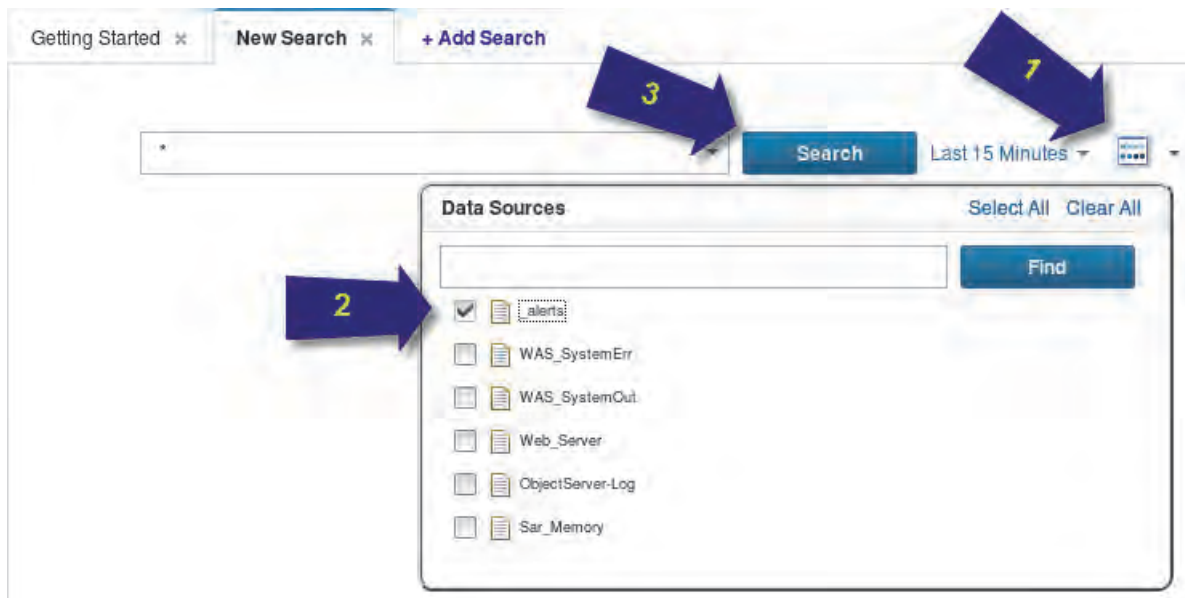
- c. Quit the mail program.

```
& quit
```

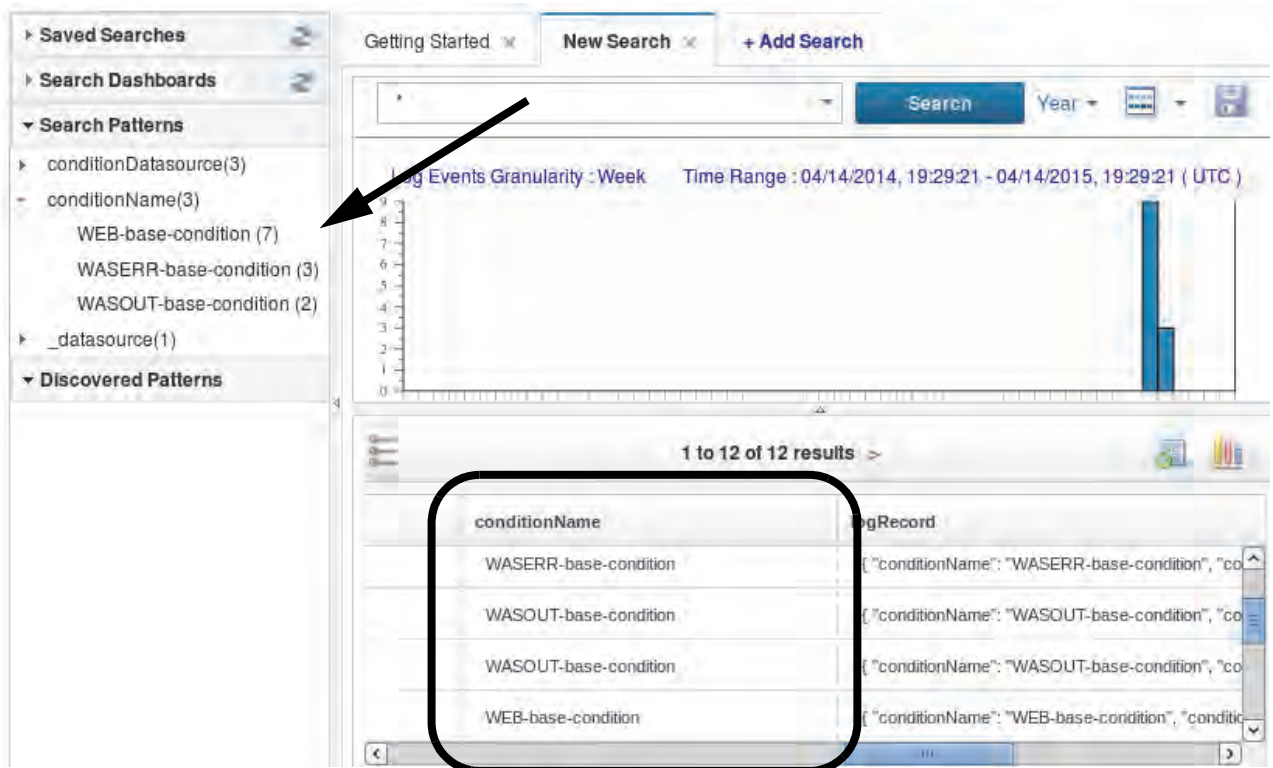
4. Open a Firefox browser and search for the log alerts.
  - a. Open a Firefox browser and browse to the following address. Log in with the user name **unityadmin** and the password **unityadmin**.  
<https://host2.tivoli.edu:9987/Unity>
  - b. Click the **New Search** tab. Select **Last Year** as the time filter.



- c. Select only **\_alerts** as the log source. Click **Search**.



Indexed alerts load in to the user interface. Notice the alerts that were triggered by the WAS and the Web\_Server conditions.



## Exercise 4 Creating composite conditions

In this exercise, you create composite conditions. These conditions accept input from base conditions.

### WebSphere multi-condition-window composite condition

1. Create a composite condition that triggers an alert when the text `java.sql.SQLException` occurs in the `WAS_SystemOut` **and** the `WAS_SystemErr` data sources within a 90-second window. Use the two WAS base conditions that you created in a preceding exercise. Use all of the actions that you configured earlier.
  - a. Verify that you are in the alerts subdirectory.

```
cd /opt/IBM/LogAnalysis/utilities/alerts
```
  - b. Make a copy of the `multiConditionWindow.json` file. Name the copy `WAS_multiConditionWindow.json`.

```
cp multiConditionWindow.json WAS_multiConditionWindow.json
```

- c. Open the `WAS_multiConditionWindow.json` file in a text editor. Edit the following five fields in the file. Make the `WAS_multiConditionWindow.json` file match the following example.

```
vi WAS_multiConditionWindow.json
```

```
{
  "name": "WAS-multi-condition",
  "description": "This condition detects java sql exceptions in two WAS logs",
  "compositeConditionTemplateName": "multi-condition-window",
  "inputConditions": ["WASOUT-base-condition", "WASERR-base-condition"],
  "parameterValues": { "windowDuration" : "90s"},
  "actions": ["index", "LAB-log-action", "LAB-email-action"]
}
```

- d. Save and close the file when you are done.

2. Use the `alerts.sh` utility to create the new condition.

- a. Run the following command to create the condition.

```
./alerts.sh -createCompositeCondition WAS_multiConditionWindow.json
```

- b. Run the following command to verify that the condition was created.

```
./alerts.sh -getCompositeCondition
```

```
...
```

```
Name: WAS-multi-condition
```

```
Description: This condition detects java sql exceptions in two WAS logs
```

```
Template: multi-condition-window
```

```
Parameter Values:
```

```
  windowDuration: 90s
```

```
Actions: [LAB-email-action, LAB-log-action, index]
```

```
Input conditions: [WASERR-base-condition, WASOUT-base-condition]
```

```
Enabled: true
```

## Web access single-condition-count composite condition

1. Create a composite condition that triggers an alert when the value of `responseTime` is over 30,000 five times within a 20-minute window. Use the `WEB-base-condition` that you created in a preceding exercise. Use all of the actions that you configured earlier.

- a. Verify that you are in the `alerts` subdirectory.

```
cd /opt/IBM/LogAnalysis/utilities/alerts
```

- b. Make a copy of the `singleConditionCount.json` file. Name the copy `WEB_singleConditionCount.json`.

```
cp singleConditionCount.json WEB_singleConditionCount.json
```

- c. Open the `WEB_singleConditionCount.json` file in a text editor. Edit the following five fields in the file. Make the `WEB_singleConditionCount.json` file match the following example.

```
vi WEB_singleConditionCount.json
```

```
{
  "name": "WEB-window-count-condition",
  "description": "This condition detects when 5 response times over 30K
  occur within 20 minutes",
  "compositeConditionTemplateName": "single-condition-count",
  "inputConditions": ["WEB-base-condition"],
  "parameterValues": { "windowDuration" : "20m", "threshold": 5},
  "actions": ["index", "LAB-log-action", "LAB-email-action"]
}
```

- d. Save and close the file when you are done.
2. Use the `alerts.sh` utility to create the new condition.
    - a. Run the following command to create the condition.

```
./alerts.sh -createCompositeCondition WEB_singleConditionCount.json
```

- b. Run the following command to verify that the condition was created.

```
./alerts.sh -getCompositeCondition
```

```
...
```

```
Name: WEB-window-count-condition
```

```
Description: This condition detects when 5 response times over 30K occur
within 20 minutes
```

```
Template: single-condition-count
```

```
Parameter Values:
```

```
    windowDuration: 20m
```

```
    threshold: 5
```

```
Actions: [LAB-email-action, LAB-log-action, index]
```

```
Input conditions: [WEB-base-condition]
```

```
Enabled: true
```

## Exercise 5 Testing composite conditions

In this exercise, you test the composite conditions that you created in the preceding exercise.

1. Run the following commands to add more messages to the WAS and web server log files.

```
/software/log_samples/scripts/WAS_Logs.sh
```

```
/software/log_samples/scripts/Web_Logs.sh
```

2. Open the `/tmp/LAB_alert.log` file. Notice the alerts in this log file that were triggered by the composite conditions.

```
more /tmp/LAB_alert.log
```

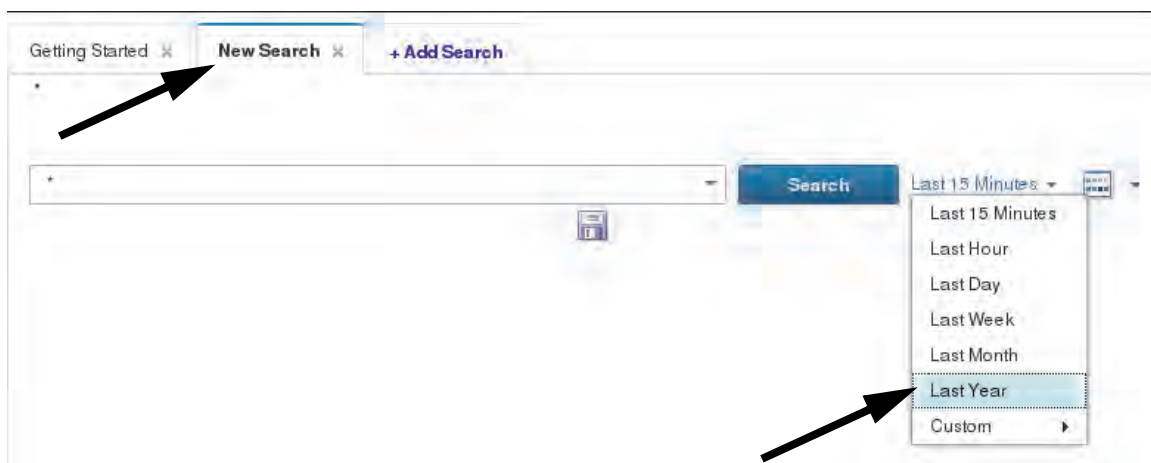
```
...
{"conditionName":"WAS-multi-condition","conditionType":"composite","datasources":
:["WAS_SystemErr","WAS_SystemOut"],"timestamp":1427037406612,"date":"2015-03-2
2T15:16:46.612Z","alertDetails":{"intervalSeconds":0,"maxDate":"2015-03-22T15:1
6:46.841Z","minDate":"2015-03-22T15:16:46.612Z","minTimestamp":1427037406612,"m
axTimestamp":1427037406841}}
...
{"conditionName":"WEB-window-count-condition","conditionType":"composite","data
sources":["Web_Server"],"timestamp":1426795356000,"date":"2015-03-19T20:02:36.0
00Z","alertDetails":{"minTimestamp":1426794823000,"minDate":"2015-03-19T19:53:4
3.000Z","maxTimestamp":1426795875000,"maxDate":"2015-03-19T20:11:15.000Z","inte
rvalSeconds":1052,"count":5}}
...
```

3. Open a Firefox browser if one is not already open. Search for the alerts that were triggered by the composite conditions.

- a. Open a Firefox browser and browse to the following address. Log in with the user name **unityadmin** and the password **unityadmin**.

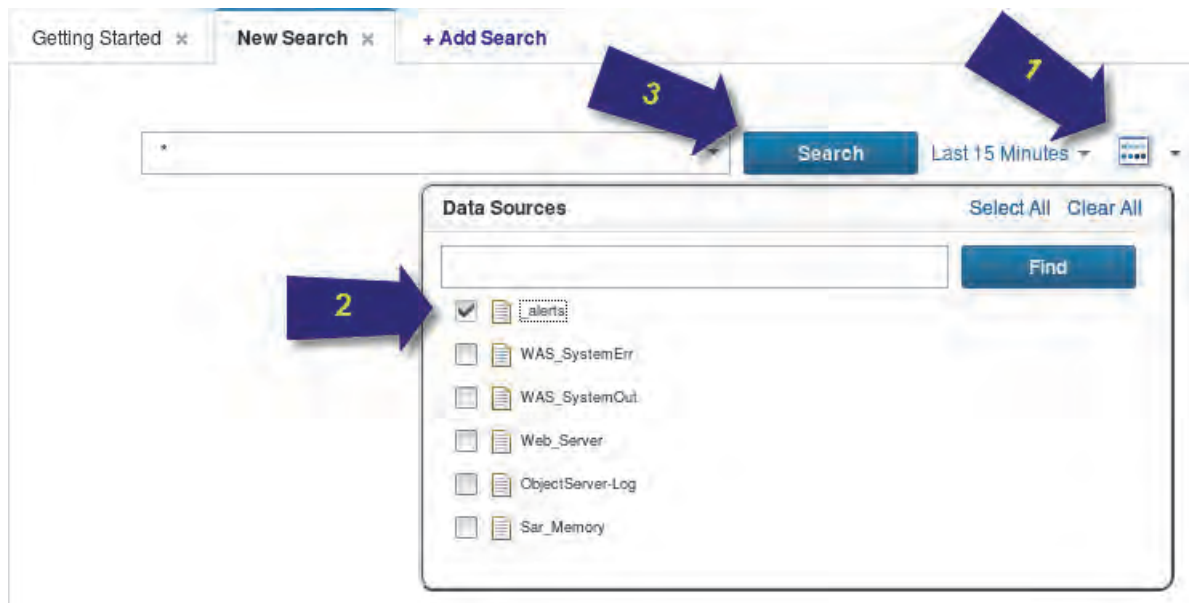
<https://host2.tivoli.edu:9987/Unity>

- b. Click the **New Search** tab. Select **Last Year** as the time filter.

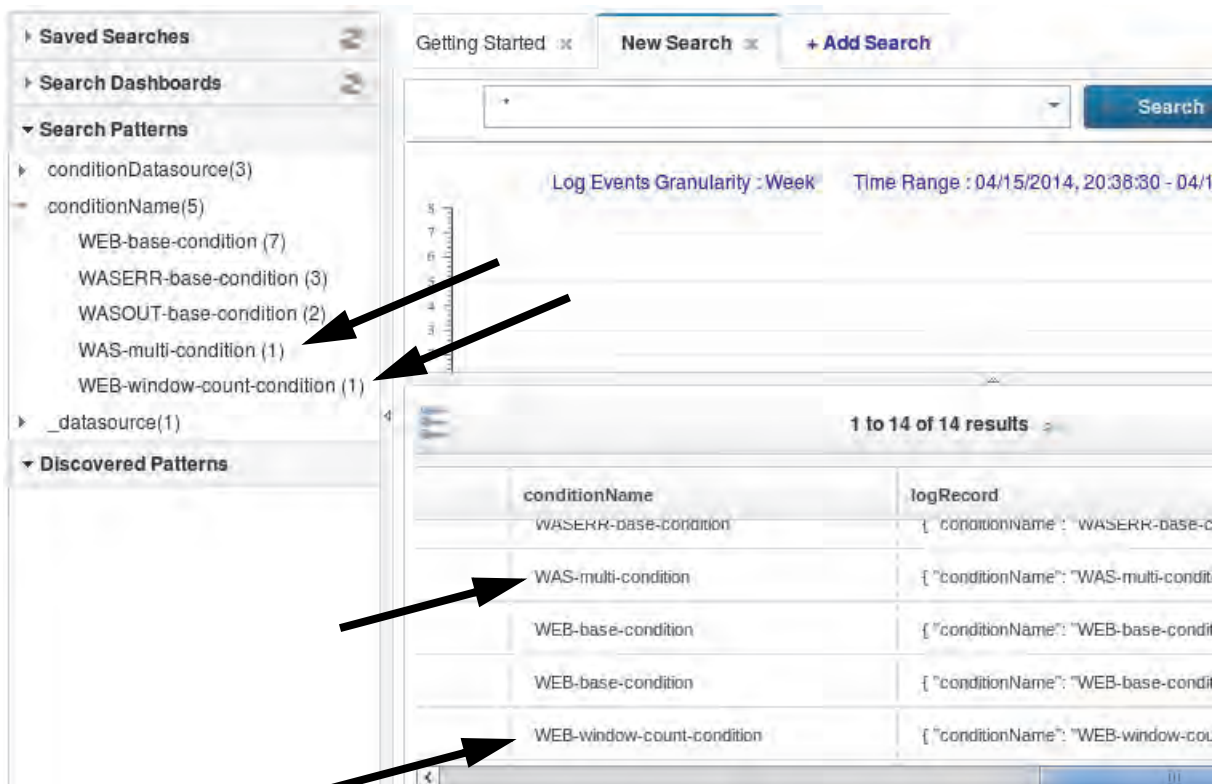




c. Select only **\_alerts** as the log source. Click **Search**.



Alerts that have been indexed load in to the user interface. Notice the alerts that were triggered by the composite conditions.





## 5 Hadoop Distributed File System (HDFS) integration exercises

In these exercises, you configure IBM Operations Analytics Log Analysis to use Hadoop Distributed File System (HDFS) for long-term data storage.

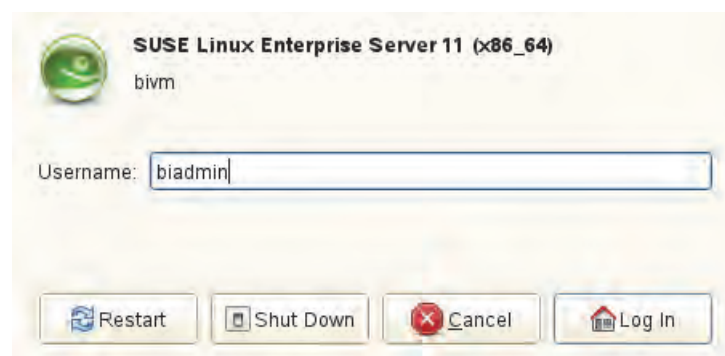
### Exercise 1 Configuring passwordless SSH

The integration of Log Analysis and HDFS requires passwordless SSH authentication between hosts for the user that owns Log analysis. You must set up passwordless SSH in a full mesh configuration for all hosts in the data node cluster, including authentication for the hosts connecting to themselves.

You use two hosts in your lab environment. The host named `host2.tivoli.edu` runs the Log Analysis software. The host named `bivm.ibm.com` runs IBM InfoSphere BigInsights version 3.0.0.1, which includes HDFS and map-reduce.



1. Start the virtual machine named `bivm.ibm.com`. Log in to the `bivm.ibm.com` host with the user name **biadmin** and the password **object00**.



2. Create a user on bivm.ibm.com that matches the **netcool** user on host2.tivoli.edu.

- a. Open a terminal window on the bivm.ibm.com host.
- b. Run the following command to verify that you are working on the correct host.

```
hostname  
bivm
```

- c. Switch to the **root** user. The password is **object00**.

```
su - root  
Password: object00
```

- d. Run the following commands to create the **netcool** user and the **ncoadmin** group.

```
groupadd -g 502 ncoadmin  
useradd -g 502 -u 501 -m -d /home/netcool -k /etc/skel -s /bin/bash netcool
```

- e. Set **object00** as the password for the **netcool** user. Ignore the warning about the bad password.

```
passwd netcool
```

```
Changing password for netcool.  
New Password: object00  
Bad password: it is based on a dictionary word  
Reenter New Password: object00  
Password changed.
```

3. Configure passwordless SSH authentication from host2.tivoli.edu to bivm.ibm.com.

- a. Switch to the host2.tivoli.edu host and open a terminal window as the **netcool** user.
- b. Run the following command to verify that you are working on the correct host.

```
hostname  
host2.tivoli.edu
```

- c. Change to the **/home/netcool** directory.

```
cd /home/netcool
```

- d. Run the following command to generate a pair of public keys. Press Enter to accept all of the default values.

```
ssh-keygen -t rsa
```

```
Generating public/private rsa key pair.  
Enter file in which to save the key (/home/netcool/.ssh/id_rsa):  
Enter passphrase (empty for no passphrase):  
Enter same passphrase again:  
Your identification has been saved in /home/netcool/.ssh/id_rsa.  
Your public key has been saved in /home/netcool/.ssh/id_rsa.pub.  
The key fingerprint is:
```

## Exercise 1 Configuring passwordless SSH

```
aa:c1:3d:69:ff:bb:05:ac:f2:8a:1c:1b:36:68:80:eb netcool@host2.tivoli.edu
```

The key's randomart image is:

```
+--[ RSA 2048]-----+
|
|
|
| .      .
| o      S o
| o o . o . .
| . o B B . .
| .. o X = .
| E = ..o.+o
+-----+
```

- e. Run the following command to create an `.ssh` subdirectory for `netcool` on the `bivm.ibm.com` host. Enter **yes** when you are prompted to continue. Enter **object00** as the password.

```
ssh netcool@bivm.ibm.com mkdir -p .ssh
```

```
The authenticity of host 'bivm.ibm.com (192.168.100.166)' can't be
established.
```

```
RSA key fingerprint is eb:35:c3:3d:bb:27:4e:0d:fe:24:fe:19:e2:ed:36:80.
```

```
Are you sure you want to continue connecting (yes/no)? yes
```

```
Warning: Permanently added 'bivm.ibm.com,192.168.100.166' (RSA) to the list
of known hosts.
```

```
Password: object00
```

- f. Copy the new public key (`id_rsa.pub`) to a file named `authorized_keys` on `bivm.ibm.com`. Enter **object00** as the password.

```
cat .ssh/id_rsa.pub | ssh netcool@bivm.ibm.com 'cat >> .ssh/authorized_keys'
```

```
Password: object00
```

- g. Run the following command to set permissions for the `.ssh` subdirectory and `authorized_keys` file on `bivm.ibm.com`.

```
ssh netcool@bivm.ibm.com "chmod 700 .ssh; chmod 640 .ssh/authorized_keys"
```

- h. Test the configuration by connecting to `bivm.ibm.com` as the **netcool** user. Run the following command and verify that you are not prompted for a password.

```
[netcool@host2 ~]$ ssh bivm.ibm.com
netcool@bivm:~>
```

- i. Type `exit` to close the SSH session to `bivm.ibm.com`.

```
exit
```

```
logout
```

```
Connection to bivm.ibm.com closed.
```

4. Configure passwordless SSH authentication from bivm.ibm.com to itself (bivm.ibm.com).

- a. Switch to the bivm.ibm.com host and open a terminal window as the **netcool** user.

```
su - netcool
Password: object00
```

- b. Run the following command to verify that you are working on the correct host.

```
hostname
bivm
```

- c. Change to the /home/netcool directory.

```
cd /home/netcool
```

- d. Run the following command to generate a pair of public keys. Press Enter to accept all of the default values.

```
ssh-keygen -t rsa
```

Generating public/private rsa key pair.

Enter file in which to save the key (/home/netcool/.ssh/id\_rsa):

Enter passphrase (empty for no passphrase):

Enter same passphrase again:

Your identification has been saved in /home/netcool/.ssh/id\_rsa.

Your public key has been saved in /home/netcool/.ssh/id\_rsa.pub.

The key fingerprint is:

86:36:e7:5b:8f:0a:cf:f0:0b:b8:b6:01:bc:14:37:67 netcool@bivm

The key's randomart image is:

```
+--[ RSA 2048 ]-----+
|
|
| . o E
| . o + .
| +   + S
| . o o =
| . o + . .
|   .o B o o
|   .o. Bo. .
+-----+
```

- e. Add the new public key (`id_rsa.pub`) to the file named `authorized_keys` on `bivm.ibm.com`. Enter **yes** when you are prompted to continue. Enter **object00** as the password.

```
cat .ssh/id_rsa.pub | ssh netcool@bivm.ibm.com 'cat >> .ssh/authorized_keys'
```

The authenticity of host 'bivm.ibm.com (192.168.100.166)' can't be established.

RSA key fingerprint is eb:35:c3:3d:bb:27:4e:0d:fe:24:fe:19:e2:ed:36:80.

Are you sure you want to continue connecting (yes/no)? **yes**

Warning: Permanently added 'bivm.ibm.com,192.168.100.166' (RSA) to the list of known hosts.

Password: **object00**

- f. Test the configuration by connecting to `bivm.ibm.com` as the **netcool** user. Run the following command and verify that you are not prompted for a password.

```
netcool@bivm:~> ssh bivm.ibm.com
```

Last login: Fri Apr 24 14:30:53 2015 from host2.tivoli.edu

```
netcool@bivm:~>
```

- g. Type `exit` to close the SSH session to `bivm.ibm.com`.

```
exit
```

```
logout
```

```
Connection to bivm.ibm.com closed.
```

5. Configure passwordless SSH authentication from `bivm.ibm.com` to `host2.tivoli.edu`.

- a. Open a terminal window as the **netcool** user, if you have not already done so.

- b. Run the following command to verify that you are working on the correct host.

```
hostname
```

```
bivm
```

- c. Change to the `/home/netcool` directory.

```
cd /home/netcool
```

- d. Add the public key on `bivm.ibm.com` (`id_rsa.pub`) to the file named `authorized_keys` on `host2.tivoli.edu`. Enter **yes** when you are prompted to continue. Enter **object00** as the password.

```
cat .ssh/id_rsa.pub | ssh netcool@host2.tivoli.edu 'cat >>
.ssh/authorized_keys'
```

The authenticity of host 'host2.tivoli.edu (192.168.100.161)' can't be established.

RSA key fingerprint is 1c:2c:83:be:ca:fd:a4:86:14:29:16:2f:76:65:af:55.

Are you sure you want to continue connecting (yes/no)? **yes**

Warning: Permanently added 'host2.tivoli.edu,192.168.100.161' (RSA) to the list of known hosts.

netcool@host2.tivoli.edu's password: **object00**



- e. Test the configuration by connecting to host2.tivoli.edu as the **netcool** user. Run the following command and verify that you are not prompted for a password.

```
netcool@bivm:~> ssh host2.tivoli.edu
```

```
Last login: Thu Mar 19 13:49:39 2015 from 192.168.100.1  
[netcool@host2 ~]$
```

- f. Type **exit** to close the SSH session to host2.tivoli.edu.

```
exit
```

```
logout
```

```
Connection to host2.tivoli.edu closed.
```

## Exercise 2 Configuring BigInsights and Hadoop

1. Copy the .jar files that Hadoop requires from host2.tivoli.edu to bivm.ibm.com.
  - a. Switch to the bivm.ibm.com host and open a terminal window as the **netcool** user, if you have not already done so.

```
su - netcool  
Password: object00
```

- b. Run the following command to verify that you are working on the correct host.

```
hostname  
bivm
```

- c. Change to the /home/netcool directory.

```
cd /home/netcool
```

- d. Run the following commands to create directories to save the .jar files.

```
mkdir LA_SERVICE_HOME  
mkdir SEARCH_JARS_TEMP
```

- e. Change to the following directory:

```
cd /home/netcool/LA_SERVICE_HOME/
```

- f. Use FTP to connect to host2.tivoli.edu. Log in with the user name **netcool** and the password **object00**.

```
ftp host2.tivoli.edu
```

```
Connected to host2.tivoli.edu.  
220 (vsFTPd 2.2.2)  
Name (host2.tivoli.edu:biadmin): netcool  
331 Please specify the password.  
Password: object00  
230 Login successful.  
Remote system type is UNIX.  
Using binary mode to transfer files.
```

- g. Run the following command to change to the directory where the .jar files are saved.

```
cd /opt/IBM/LogAnalysis/utilities/hadoop/
```

- h. Run the following command to copy the file named service.zip.

```
get service.zip
```

- i. Run the following command to change your local directory.

```
lcd /home/netcool/SEARCH_JARS_TEMP
```

- j. Run the following command to copy the file named search.zip.

```
get search.zip
```

- k. Close the FTP session.

```
quit
```

2. Decompress the service.zip file and start the Log Analysis - BigInsights ingestion service.

- a. Change to the directory where you saved the service.zip file.

```
cd /home/netcool/LA_SERVICE_HOME
```

- b. Decompress the service.zip file.

```
unzip service.zip
```

- c. Start the Hadoop service.

```
bin/server.sh clusterStart
```



```
Starting server on [bivm.ibm.com] ... starting rpc server at port 9003.  
Server started. Appending logs to file  
/home/netcool/LA_SERVICE_HOME/logs/rpcServer.out
```

- d. Check the status of the service.


```
bin/server.sh clusterStatus
```

```
Getting status for [bivm.ibm.com] ... Server running with process id=9479.
```

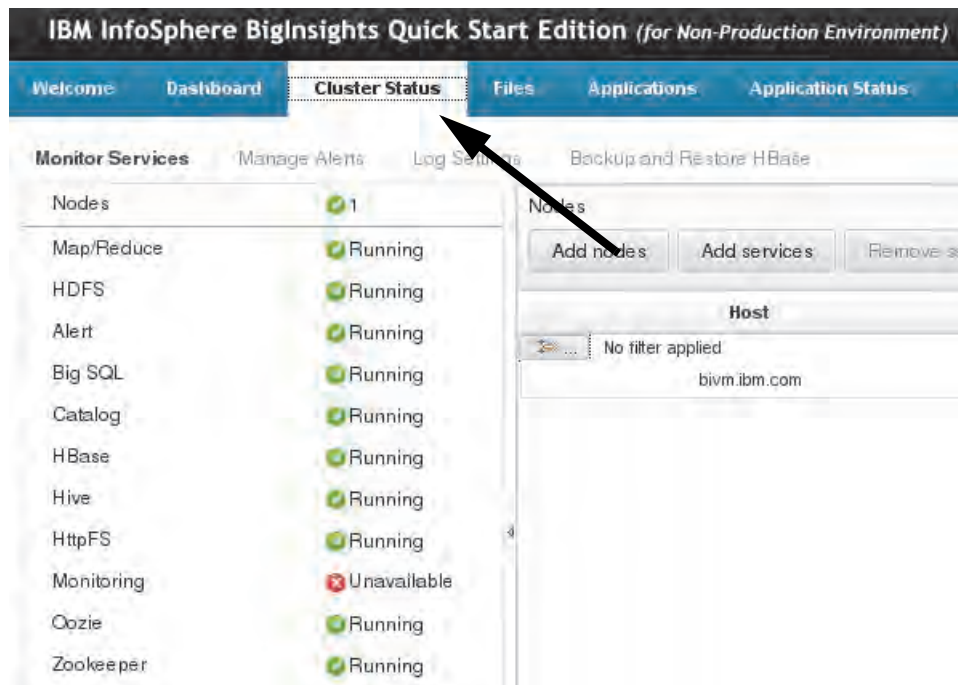
3. Start IBM InfoSphere BigInsights. Verify that BigInsights and all of its applications are available.
  - a. Double-click the Start BigInsights icon on the biadmin desktop. Wait for BigInsights and all of its applications to start.



- b. Double-click the Web console shortcut on the biadmin desktop. Log in with the user name **biadmin** and the password **object00**.

A screenshot of the IBM InfoSphere BigInsights Quickstart Edition login screen. The title is 'IBM® InfoSphere® BigInsights™ Quickstart Edition'. Below the title, it says 'Please enter your information'. There are two input fields: 'User name:' with the text 'biadmin' entered, and 'Password:' with a masked password represented by ten dots. At the bottom, there are 'Login' and 'Cancel' buttons. At the very bottom, there is a small copyright notice: 'Licensed Materials - Property of IBM Corp. © Copyright 2010, 2014. IBM Corporation. IBM, InfoSphere and BigInsights are trademarks of IBM Corporation, registered in many jurisdictions worldwide.'

- c. Click the **Cluster Status** tab. Verify that all nodes are running.



**Note:** You can ignore the unavailable Monitoring node.

4. Create directories for Log Analysis data on the HDFS data node as the **biadmin** user.
- Open a new terminal window and switch to the biadmin user. The password is **object00**.  

```
su - biadmin  
Password: object00
```
  - Run the following command to verify that you are working on the correct host.  

```
hostname  
bivm
```
  - Run the following commands to create the required directories.  

```
hadoop fs -mkdir /la-hadoop-tier  
hadoop fs -mkdir /la-hadoop-tier/data  
hadoop fs -mkdir /la-hadoop-tier/jars  
hadoop fs -mkdir /la-hadoop-tier/output
```
  - Run the following command to change the owner of these directories to **netcool**.  

```
hadoop fs -chown -R netcool:ncoadmin /la-hadoop-tier
```

- e. Run the following commands to verify that the directories were created and that **netcool** owns them.

```
hadoop fs -ls /
```

```
drwxrwxr-x   - hdfs    biadmin          0 2014-09-18 23:55 /biginsights
drwxr-xr-x   - hdfs    biadmin          0 2014-09-18 23:32 /hadoop
drwxr-xr-x   - hbase    biadmin          0 2015-04-24 15:37 /hbase
drwxr-xr-x   - netcool ncoadmin         0 2015-04-24 16:20 /la-hadoop-tier
drwxrwxrwt   - hdfs    biadmin          0 2014-09-19 00:34 /tmp
drwxrwxrwx   - hdfs    biadmin          0 2014-09-25 20:11 /user
```

```
hadoop fs -ls /la-hadoop-tier
```

```
drwxr-xr-x   - netcool ncoadmin          0 2015-04-24 16:20
/la-hadoop-tier/data
drwxr-xr-x   - netcool ncoadmin          0 2015-04-24 16:20
/la-hadoop-tier/jars
drwxr-xr-x   - netcool ncoadmin          0 2015-04-24 16:20
/la-hadoop-tier/output
```

5. As the **netcool** user, copy the service .jar files to the new `/la-hadoop-tier/jars` directory.

- a. Open a terminal window and switch to the **netcool** user if you do not already have one open.

```
su - netcool
Password: object00
```

- b. Run the following command to verify that you are working on the correct host.

```
hostname
bivm
```

- c. Run the following command to copy the service .jar files to the correct directory.

```
hadoop fs -copyFromLocal /home/netcool/LA_SERVICE_HOME/lib/*.jar
/la-hadoop-tier/jars/
```

6. As the **netcool** user, copy the search .jar files to the new `/la-hadoop-tier/jars` directory.

- a. Run the following commands to decompress the search .jar files.

```
cd /home/netcool/SEARCH_JARS_TEMP
```

```
unzip search.zip
```

- b. Run the following command to copy the search .jar files to the correct directory.

```
hadoop fs -copyFromLocal /home/netcool/SEARCH_JARS_TEMP/*.jar
/la-hadoop-tier/jars/
```



**Note:** You can ignore the File exists messages.

7. Verify that all of the .jar files were copied to the /la-hadoop-tier/jars directory.

```
hadoop fs -ls /la-hadoop-tier/jars
```

## Exercise 3 Configuring Log Analysis

1. Edit the unitysetup.properties file on the Log Analysis server to enable the Hadoop integration.

- a. Switch to the host2.tivoli.edu host and open a terminal window as the **netcool** user.
- b. Run the following command to verify that you are working on the correct host.

```
hostname
host2.tivoli.edu
```

- c. Open the unitysetup.properties file with a text editor.

```
vi
/opt/IBM/LogAnalysis/wlp/usr/servers/Unity/apps/Unity.war/WEB-INF/unitysetup
.properties
```

- d. Find the **INDEX\_IMPLEMENTATION** property. Change the value to **SOLR,HADOOP**.

```
...
#Number of charts allowed in dynamic dashboards
DASHBOARD_CHARTS_LIMIT=8

#Underlying search index implementation
INDEX_IMPLEMENTATION=SOLR,HADOOP

#Setting range on minute to calculate Granularity (in minutes)
GRANULARITY_MINUTE_LOWER_LIMIT=15
GRANULARITY_MINUTE_UPPER_LIMIT=60

#Wait time for get search result call (in ms)
...
```



- e. Find the HADOOP\_TIER\_HDFS\_BASE\_DIR and HADOOP\_TIER\_JOB\_TRACKER\_URI properties. Change the value of these properties to match the following example.

```
...
#Maximum no. of charts allowed to auto refresh across dashboards
MAX_AUTO_REFRESH_CHARTS=20

#Hadoop-tier properties
#Enable HADOOP_TIER by adding HADOOP to INDEX_IMPLEMENTATION property ->
INDEX_IMPLEMENTATION=SOLR,HADOOP
HADOOP_TIER_HDFS_BASE_DIR=hdfs://bivm.ibm.com:9000/la-hadoop-tier
HADOOP_TIER_HDFS_ADMIN_USER=hdfs
HADOOP_TIER_JOB_TRACKER_URI=bivm.ibm.com:9001
HADOOP_TIER_SERVER_PORT=9003
...
```

- f. Save and close the file when you are finished.

2. Copy the core-site.xml and hdfs-site.xml configuration files from the bivm.ibm.com host to the host2.tivoli.edu host.

- a. Change to the directory where Log Analysis expects to see the core-site.xml and hdfs-site.xml files.

```
cd
/opt/IBM/LogAnalysis/wlp/usr/servers/Unity/apps/Unity.war/WEB-INF/classes/
```

- b. Run the following commands to copy the files from bivm.ibm.com and save them locally.

```
scp netcool@bivm.ibm.com:/opt/ibm/biginsights/hadoop-conf/core-site.xml .
scp netcool@bivm.ibm.com:/opt/ibm/biginsights/hadoop-conf/hdfs-site.xml .
```

- c. Verify that the files were correctly copied.

```
ls -l
-rw-r--r-- 1 netcool ncoadmin 3416 Apr 24 20:47 core-site.xml
-rw-r--r-- 1 netcool ncoadmin 3906 Apr 24 20:47 hdfs-site.xml
-rw-r--r-- 1 netcool ncoadmin 3520 Mar 19 18:35 log4j.properties
```

3. Stop Log Analysis.

```
/opt/IBM/LogAnalysis/utilities/unity.sh -stop
```

4. Start Log Analysis.

```
/opt/IBM/LogAnalysis/utilities/unity.sh -start
```

## Exercise 4 Verifying the integration

After you enable the integration, Log Analysis stores log data to both Solr and HDFS. In this exercise, you verify that Log Analysis is writing data to the HDFS system.



**Important:** Run all of the steps in this exercise on the host2.tivoli.edu virtual machine.


1. Run the following commands to create more log messages for Log Analysis to process.

```
/software/log_samples/scripts/WAS_Logs.sh  
/software/log_samples/scripts/Web_Logs.sh
```

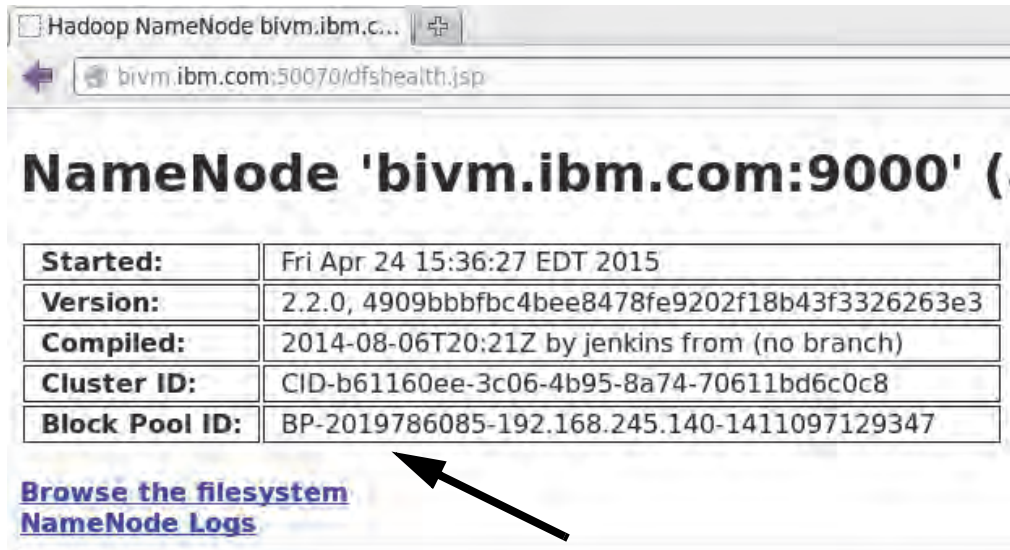
2. Open the `/opt/IBM/LogAnalysis/logs/Hadooptier.log` file. Look for messages like the following example. These messages verify that Log Analysis is communicating with the new BigInsights ingestion service.

```
04/27/15 13:36:55:437 UTC [Thread-46] INFO  
com.ibm.tivoli.unity.hadoop.ingestion.client.ServiceNode - Retrieved batches  
status for service node [192.168.100.166:9003]  
04/27/15 13:37:00:438 UTC [Thread-46] INFO  
com.ibm.tivoli.unity.hadoop.ingestion.client.ServiceNode - Retrieving batches  
status from service node [192.168.100.166:9003]  
04/27/15 13:37:00:439 UTC [Thread-46] INFO  
com.ibm.tivoli.unity.hadoop.ingestion.client.ServiceNode - Received updates for  
[3] batches from service node [192.168.100.166:9003]. [0] batches finished with  
errors.
```

3. Use a browser to view the HDFS file system and verify that log data is present.
  - a. Double-click the **Firefox** icon on the desktop.



- b. Browse to the address <http://bivm.ibm.com:50070>. Click **Browse the filesystem**.



|                       |                                                |
|-----------------------|------------------------------------------------|
| <b>Started:</b>       | Fri Apr 24 15:36:27 EDT 2015                   |
| <b>Version:</b>       | 2.2.0, 4909bbbfb4bee8478fe9202f18b43f3326263e3 |
| <b>Compiled:</b>      | 2014-08-06T20:21Z by jenkins from (no branch)  |
| <b>Cluster ID:</b>    | CID-b61160ee-3c06-4b95-8a74-70611bd6c0c8       |
| <b>Block Pool ID:</b> | BP-2019786085-192.168.245.140-1411097129347    |

[Browse the filesystem](#)  
[NameNode Logs](#)

- c. Click **la-hadoop-tier**.

| Name                           | Type | Size | Replication | Block Size | Modification Time | Permission | Owner   | Group    |
|--------------------------------|------|------|-------------|------------|-------------------|------------|---------|----------|
| <a href="#">biginsights</a>    | dir  |      |             |            | 2014-09-18 23:55  | rw-rw-r--  | hdfs    | biadmin  |
| <a href="#">hadoop</a>         | dir  |      |             |            | 2014-09-18 23:32  | rw-rw-r--  | hdfs    | biadmin  |
| <a href="#">hbase</a>          | dir  |      |             |            | 2015-04-24 15:37  | rw-rw-r--  | hbase   | biadmin  |
| <a href="#">la-hadoop-tier</a> | dir  |      |             |            | 2015-04-24 16:20  | rw-rw-r--  | netcool | ncoadmin |
| <a href="#">tmp</a>            | dir  |      |             |            | 2014-09-19 00:34  | rw-rw-rw-  | hdfs    | biadmin  |
| <a href="#">user</a>           | dir  |      |             |            | 2014-09-25 20:11  | rw-rw-rw-  | hdfs    | biadmin  |

- d. Click **data**.

[Go to parent directory](#)

| Name                   | Type | Size | Replication | Block Size | Modification Time | Permission | Owner   | Group    |
|------------------------|------|------|-------------|------------|-------------------|------------|---------|----------|
| <a href="#">data</a>   | dir  |      |             |            | 2015-04-27 09:35  | rw-rw-r--  | netcool | ncoadmin |
| <a href="#">jars</a>   | dir  |      |             |            | 2015-04-24 16:29  | rw-rw-r--  | netcool | ncoadmin |
| <a href="#">output</a> | dir  |      |             |            | 2015-04-24 16:20  | rw-rw-r--  | netcool | ncoadmin |

- e. Click **UnityCollection\_<date>\_UTC**.

[Go to parent directory](#)

| Name                                                    | Type | Size | Replication | Block Size | Mod  |
|---------------------------------------------------------|------|------|-------------|------------|------|
| <a href="#">.tmp</a>                                    | dir  |      |             |            | 2015 |
| <a href="#">UnityCollection_27_04_2015_00_00_00_UTC</a> | dir  |      |             |            | 2015 |

- f. Notice the subdirectories that are named like the Log Analysis data sources. Click one of the subdirectories.

[Go to parent directory](#)

| Name                          | Type | Size | Replication | Block Size | Modification |
|-------------------------------|------|------|-------------|------------|--------------|
| <a href="#">WAS_SystemErr</a> | dir  |      |             |            | 2015-04-27 0 |
| <a href="#">WAS_SystemOut</a> | dir  |      |             |            | 2015-04-27 0 |
| <a href="#">Web_Server</a>    | dir  |      |             |            | 2015-04-27 0 |

- g. Click one of the date directories.

[Go to parent directory](#)

| Name                          | Type | Size | Replication | Block Size | Modification Time | Permission | Owner   | Group   |
|-------------------------------|------|------|-------------|------------|-------------------|------------|---------|---------|
| <a href="#">2015-03-03_00</a> | dir  |      |             |            | 2015-04-27 09:36  | rwxr-xr-x  | netcool | ncadmin |
| <a href="#">2015-03-04_00</a> | dir  |      |             |            | 2015-04-27 09:36  | rwxr-xr-x  | netcool | ncadmin |
| <a href="#">2015-03-05_00</a> | dir  |      |             |            | 2015-04-27 09:36  | rwxr-xr-x  | netcool | ncadmin |
| <a href="#">2015-03-06_00</a> | dir  |      |             |            | 2015-04-27 09:36  | rwxr-xr-x  | netcool | ncadmin |
| <a href="#">2015-03-07_00</a> | dir  |      |             |            | 2015-04-27 09:36  | rwxr-xr-x  | netcool | ncadmin |

Notice the .avro file or files. These files contain log data in compressed binary format.

Goto :

[Go to parent directory](#)

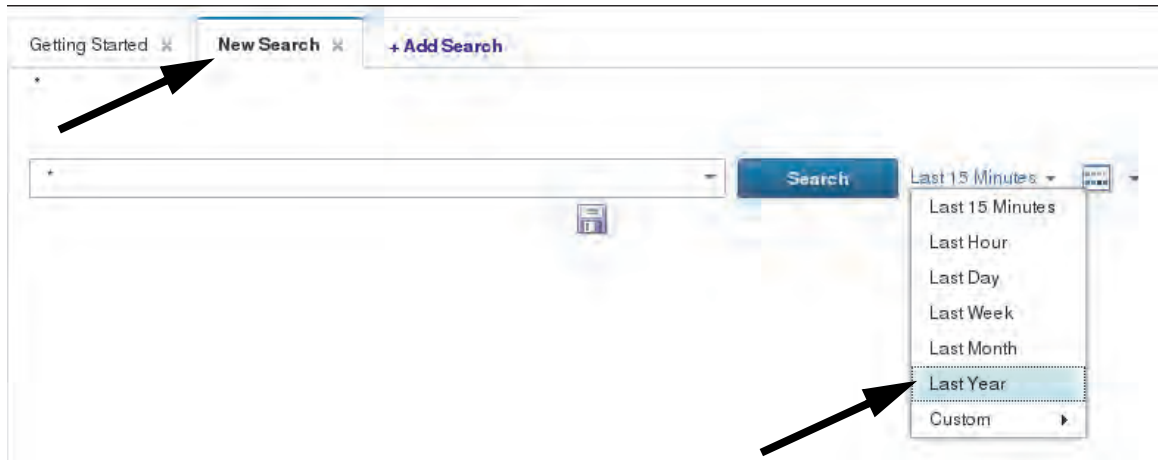
| Name                  | Type | Size    | Replication | Block Size | Modification Time | Per |
|-----------------------|------|---------|-------------|------------|-------------------|-----|
| <a href="#">Lavro</a> | file | 4.00 KB | 1           | 128 MB     | 2015-04-27 09:36  | rw- |

[Go back to DFS home](#)

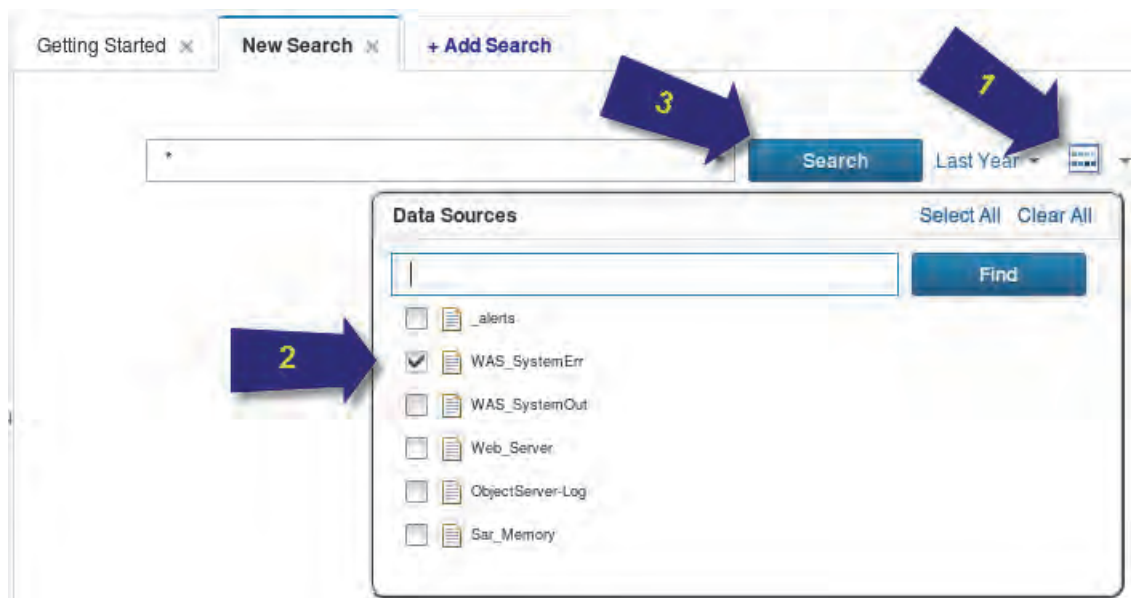
4. Open a Firefox browser and search for the log messages. Search through one of the data sources you saw in the HDFS directory structure, such as WAS\_SystemErr.
- a. Open a Firefox browser and browse to the following address. Log in with the user name **unityadmin** and the password **unityadmin**.
- <https://host2.tivoli.edu:9987/Unity>



- b. Click the **New Search** tab. Select **Last Year** as the time filter.



- c. Select only **WAS\_SystemErr** as the log source. Click **Search**.



- d. Log messages load in to the user interface. Leave the search results open.
5. Open the `/opt/IBM/LogAnalysis/logs/UnityApplication.log` file. Look for messages like the following example. These messages verify that Log Analysis is searching the HDFS node for data. Notice that Log analysis is also searching for data in the Solr file system.

```
04/27/15 14:35:55:733 UTC [Default Executor-thread-15] INFO -  
JAXRSUnitySearchServlet : New search query is being POSTed. Mode=async  
04/27/15 14:35:55:749 UTC [Default Executor-thread-15] INFO - USRHandler :  
Collection: WAS_SystemErr, bean:  
id=31 name=WAS_SystemErr sourceType=13 indexingConfig=null  
annotator=0  
04/27/15 14:35:55:755 UTC [Default Executor-thread-15] INFO -  
InsightPackResourceBundleManager : __ENTRY translatedIndexConfigFields()
```

```

04/27/15 14:35:55:756 UTC [Default Executor-thread-15] INFO -
InsightPackArtifactMapHandler : Not in the cache! Going to DB for SOURCETYPE id
13
04/27/15 14:35:55:910 UTC [Default Executor-thread-15] INFO -
InsightPackResourceBundleManager : __EXIT translatedIndexConfigFields()
04/27/15 14:35:55:920 UTC [Default Executor-thread-15] INFO - UnitySearchQuery
: Disabled highlighting for * query
04/27/15 14:35:55:971 UTC [Default Executor-thread-15] INFO -
UnitySearchRuntime : Checking for hadoop data in query, time-stamp range =
[1398609355000, 1430145355401]
04/27/15 14:35:55:972 UTC [Default Executor-thread-15] INFO -
UnitySearchRuntime : Min-colr-tier-timestamp: 1426723200000,
queryReferencesHadoopData: true
04/27/15 14:35:55:972 UTC [Default Executor-thread-15] INFO -
UnitySearchRuntime : Search Query ID: Query mode: asynchronous
04/27/15 14:35:55:976 UTC [Default Executor-thread-15] INFO -
UnitySearchRuntime : Submitted Solr search query: 1
04/27/15 14:35:55:976 UTC [Thread-87] INFO - SearchQueryRunner : Starting query
runner thread for Solr query 1
04/27/15 14:35:55:977 UTC [Thread-87] INFO - SolrSearchQuery : Solr search
query: ***
04/27/15 14:35:55:981 UTC [Default Executor-thread-15] INFO -
UnitySearchRuntime : Submitted Hadoop search query: 1
04/27/15 14:35:55:982 UTC [Thread-88] INFO - SearchQueryRunner : Starting query
runner thread for hadoop query 1
04/27/15 14:35:55:982 UTC [Thread-88] INFO - HadoopTierSearchQuery : Retrieving
hadoop-tier partitions: maxWriteTime=1426723200000
04/27/15 14:35:55:989 UTC [Thread-87] INFO - SolrSearchQuery : Solr filter
query: +(_datasource:"WAS_SystemErr") +timestamp:[2014-04-27T14:35:55.000Z TO
2015-04-27T14:35:55.401Z}
04/27/15 14:35:56:130 UTC [Thread-88] INFO - HadoopTierSearchQuery : Retrieved
partitions, num = 0
04/27/15 14:35:56:130 UTC [Thread-88] INFO - HadoopTierSearchQuery : No
hadoop-tier partitions, state = COMPLETE
04/27/15 14:35:56:130 UTC [Thread-88] INFO - SearchQueryRunner : Completed
executing hadoop-tier queries: 1
04/27/15 14:35:56:130 UTC [Thread-88] INFO - SearchQueryRunner : Completed
query execution for hadoop-tier query ID 1, status = COMPLETE
04/27/15 14:35:56:196 UTC [Default Executor-thread-16] INFO - LogSourcesHandler
: 6
04/27/15 14:35:56:976 UTC [Thread-87] INFO - SolrSearchQuery : Solr Query
1[1/2], search Time: 984
04/27/15 14:35:56:976 UTC [Thread-87] INFO - SolrSearchQuery : Solr Query 1,
total Results: 557
04/27/15 14:35:56:977 UTC [Thread-87] INFO - SolrSearchQuery : Solr Query 1,
num Results: 557
04/27/15 14:35:57:019 UTC [Thread-87] INFO - SolrSearchQuery : Solr search
query: ***

```

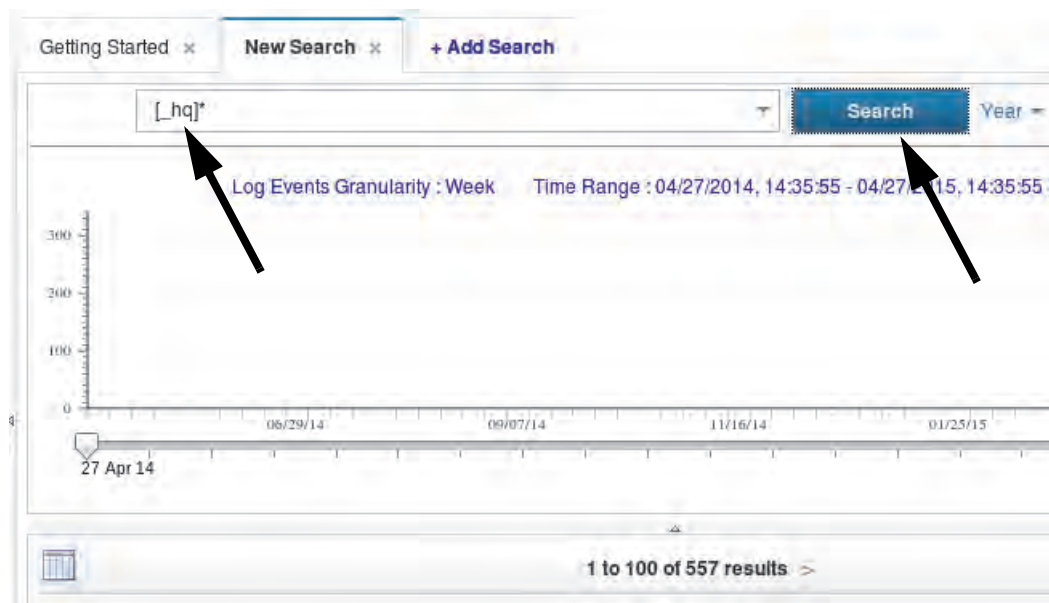


```
04/27/15 14:35:57:019 UTC [Thread-87] INFO - SolrSearchQuery : Solr filter
query: +(_datasource:"WAS_SystemErr") +timestamp:[2014-04-27T14:35:55.000Z TO
2015-04-27T14:35:55.401Z}
...
04/27/15 14:35:57:532 UTC [Default Executor-thread-3] INFO - UnitySearchResult
: COMBINED-STATE: COMPLETE
```



**Important:** When a user searches for log messages, Log Analysis queries the Solr and Hadoop file systems. The data in both file systems is combined and displayed in the user search results. If there is identical data in both file systems, Log Analysis only shows data from Solr to the user. In this example, only data from Solr is returned to the search results. This is why you see messages like No hadoop-tier partitions in the UnityApplication log file.

6. Force the Log Analysis user interface to search for data only in the HDFS node.
  - a. Return to the Firefox browser where the Log Analysis user interface is open.
  - b. Verify that you are searching the **WAS\_SystemErr** data source. Type `[_hq]*` in the search filter field. Click **Search**. Log messages load in to the search results.



7. Open the `/opt/IBM/LogAnalysis/logs/UnityApplication.log` file again. Look for messages like the following example. These messages verify that Log Analysis is retrieving data from the HDFS node.

```
04/27/15 15:11:26:718 UTC [Thread-90] INFO - HadoopTierSearchQuery : Hadoop
Query 2[1/1], search Time: 35871
04/27/15 15:11:26:718 UTC [Thread-90] INFO - HadoopTierSearchQuery : Hadoop
Query 2, total Results: 557
04/27/15 15:11:26:718 UTC [Thread-90] INFO - HadoopTierSearchQuery : Hadoop
Query 2, num Results: 90
```

```
04/27/15 15:11:26:718 UTC [Thread-90] INFO - SearchQueryRunner : Completed
executing hadoop-tier queries: 2
04/27/15 15:11:26:718 UTC [Thread-90] INFO - SearchQueryRunner : Completed
query execution for hadoop-tier query ID 2, status = COMPLETE
```

## Exercise 5 Disabling the HDFS integration

These steps show you how to disable the HDFS integration.



**Important:** Run all of the steps in this exercise on the **host2.tivoli.edu** virtual machine.

1. Edit the `unitysetup.properties` file on the Log Analysis server to disable the Hadoop integration.

- a. Open a terminal window as the **netcool** user, if you do not already have one open.

- b. Open the `unitysetup.properties` file with a text editor.

```
vi
/opt/IBM/LogAnalysis/wlp/usr/servers/Unity/apps/Unity.war/WEB-INF/unitysetup
.properties
```

- c. Find the **INDEX\_IMPLEMENTATION** property. Change the value to **SOLR**.

```
...
#Number of charts allowed in dynamic dashboards
DASHBOARD_CHARTS_LIMIT=8

#Underlying search index implementation
INDEX_IMPLEMENTATION=SOLR

#Setting range on minute to calculate Granularity (in minutes)
GRANULARITY_MINUTE_LOWER_LIMIT=15
GRANULARITY_MINUTE_UPPER_LIMIT=60

#Wait time for get search result call (in ms)
...
```

- d. Save and close the file when you are finished.

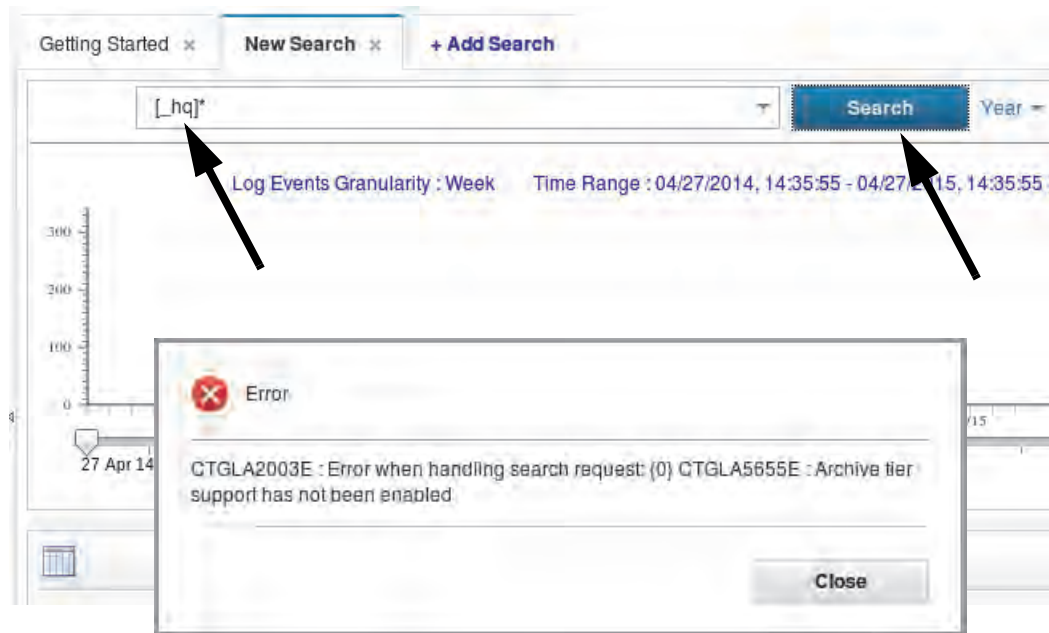
2. Stop Log Analysis.

```
/opt/IBM/LogAnalysis/utilities/unity.sh -stop
```

3. Start Log Analysis.

```
/opt/IBM/LogAnalysis/utilities/unity.sh -start
```

4. Force the Log Analysis user interface to search for data only in the HDFS node again to verify that the HDFS integration has been disabled.
  - a. Return to the Firefox browser where the Log Analysis user interface is open.
  - b. Verify that you are searching the **WAS\_SystemErr** data source. Type `[_hq]*` in the search filter field. Click **Search**. An error message about archive tier support confirms that you have disabled the HFDS integration.



**Note:** You can find similar error messages in the `UnityApplication.log` file.

5. Shut down the virtual machine named **bivm.ibm.com**.



**Important:** You do not use the **bivm.ibm.com** virtual machine for any other exercises during this course.



## 6 Performance tuning exercises

### Exercise 1 Tuning the EIF Receiver

In this exercise, you change the configuration of the EIF (Tivoli® Event Integration Facility) Receiver. You use IBM Operations Analytics Log Analysis log files to confirm the configuration change. This exercise shows you how to meter incoming log messages from streaming log sources.

1. Monitor the IBM Operations Analytics Log Analysis log files to look for messages that are related to data ingestion.
  - a. Open a new terminal and tail the `GenericReceiver.log` file.

```
tail -f /opt/IBM/LogAnalysis/logs/GenericReceiver.log
```
  - b. Open a new terminal and tail the `UnityEifReceiver.log` file.

```
tail -f /opt/IBM/LogAnalysis/logs/UnityEifReceiver.log
```
  - c. Run the following command to create more log messages in the WAS\_SystemOut data source. Watch the `UnityEifReceiver.log` and `GenericReceiver.log` log messages. Notice the delay between new messages in `UnityEifReceiver.log` and new messages in `GenericReceiver.log`.

```
/software/log_samples/scripts/WAS_Logs.sh
```



**Note:** `UnityEifReceiver.log` shows messages that are sent from the Log File Agent. `GenericReceiver.log` shows messages from the Log Analysis common data ingestion interface.

2. Wait about 90 seconds, or until new log messages stop arriving. Press Ctrl+C to stop the tail in both logs. Look for messages in `UnityEifReceiver.log` that show the EIF Receiver posting the JSON payload and the size of the posted data. Notice the number of successes and failures in the post event. Notice the indexed source volume. Messages like this one have a corresponding message in `GenericReceiver.log`.

```
04/16/15 16:15:30:235 UTC [pool-6-thread-1] INFO - LogEventPoster : -----
Posting Event to UNITY DATA COLLECTOR -
https://host2.tivoli.edu:9987/Unity/DataCollector
04/16/15 16:15:30:235 UTC [pool-6-thread-1] INFO - LogEventPoster : Post
Queue-Remove Operation:0
```

```
04/16/15 16:15:30:240 UTC [pool-6-thread-1] INFO - LogEventPoster : Posting  
Post Data Json of size:275607  
04/16/15 16:15:30:736 UTC [pool-6-thread-1] INFO - LogEventPoster : Total  
Response time taken(sec):0  
04/16/15 16:15:30:736 UTC [pool-6-thread-1] INFO - LogEventPoster : ++++++++  
RESPONSE MESSAGE ++++++++  
04/16/15 16:15:30:736 UTC [pool-6-thread-1] INFO - LogEventPoster : OK  
04/16/15 16:15:30:736 UTC [pool-6-thread-1] INFO - LogEventPoster :  
{ "BATCH_STATUS": { "failures": [], "stream": "_unity_default_stream", "writeTime": "20  
15-04-16T16:04:04.994+0000", "indexedSourceVolume": 266774, "indexNumSuccessful": 5  
57, "indexNumFailures": 0, "indexBatchIds": "1429200244994_54", "batchSize": 557, "ba  
tchId": "1429200930584_58", "numSuccessful": 557, "numFailures": 0 }, "RESPONSE_MESSAG  
E": "INPUT_BATCH_PROCESSED", "RESPONSE_CODE": 200 }
```

3. Look for messages in GenericReceiver.log that show the generic receiver processing messages that were sent from the EIF receiver.

```
04/16/15 16:15:35:275 UTC [Default Executor-thread-8] INFO -  
UnityFlowController : Batch Status for -> WAS_SystemErr , Size: 557 , Num  
successful: 557 , Num failures: 0 , Indexed Source volume: 0  
04/16/15 16:15:35:275 UTC [Default Executor-thread-8] INFO -  
DataCollectorRestServlet : Batch of Size 557 processed and encountered 0  
failures  
04/16/15 16:15:35:056 UTC [Thread-60] INFO - IndexStatusChecker : Updating  
statistics for data source [WAS_SystemErr], stream [_unity_default_stream],  
ingested bytes [266774], write date [Wed Thu 16 16:15:35 UTC 2015].
```

4. In a separate terminal, stop IBM Operations Analytics Log Analysis.
- /opt/IBM/LogAnalysis/utilities/unity.sh -stop
5. Edit the EIF receiver configuration. Change the following properties. The properties are at the bottom of the file.

- logsource.buffer.wait.timeout=2
- logsource.max.buffer.size=1024

- a. Open the EIF receiver configuration file.

```
vi /opt/IBM/LogAnalysis/UnityEIFReceiver/config/unity.conf
```

- b. Change the following two properties to match the following example.

```
#Timeout in Seconds  
logsource.buffer.wait.timeout=2  
#Buffer Size in Bytes  
logsource.max.buffer.size=1024
```

- c. Save and close the file when you are done.



**Important:** These settings are not appropriate for a production system. The settings in this example are set intentionally low.

6. Back up and remove the `UnityEifReceiver.log` and `GenericReceiver.log` files. Removing these log files make it easier to find new log messages.

```
cd /opt/IBM/LogAnalysis/logs/  
mv GenericReceiver.log GenericReceiver.log.old  
mv UnityEifReceiver.log UnityEifReceiver.log.old
```

7. Start IBM Operations Analytics Log Analysis.

```
/opt/IBM/LogAnalysis/utilities/unity.sh -start
```

8. Wait for all of the log analysis components to start. Resume the tail of the `UnityEifReceiver.log` and `GenericReceiver.log` files.
  - a. Find the two terminals that you opened for viewing these log files in a preceding step.
  - b. Resume the tail of the log files. Run these commands in two different terminal windows.

```
tail -f /opt/IBM/LogAnalysis/logs/GenericReceiver.log  
tail -f /opt/IBM/LogAnalysis/logs/UnityEifReceiver.log
```



**Hint:** When you view these log files, it might be easier to stop and start the tail on the `UnityEifReceiver.log` and `GenericReceiver.log` files to search for specific messages. You can start the tail with the two preceding commands. You can stop the tail with `Ctrl+C`.

9. Look at the first messages in `UnityEifReceiver.log` after it starts. Notice the updated properties in the log file. The full path to this log file is

```
/opt/IBM/LogAnalysis/logs/UnityEifReceiver.log.  
Unity Data Collector  
KEYSTORE=/opt/IBM/LogAnalysis/wlp/usr/servers/Unity/keystore/unity.ks  
Unity Data Collector LOGSOURCE TIMEOUT=2  
Unity Data Collector LOGSOURCE BUFFER SIZE=1024  
Unity Data Collector MAX EV SERVICE JSON QUEUE SIZE=80000  
Unity Data Collector MAX EV POSTER QUEUE SIZE=500
```

10. Run the following command to create more log messages in the `WAS_SystemOut` data source.

```
/software/log_samples/scripts/WAS_Logs.sh
```

11. Wait approximately 60 seconds. Watch the `UnityEifReceiver.log` and `GenericReceiver.log` log messages. Notice that the delay between new messages in `UnityEifReceiver.log` and new messages in `GenericReceiver.log` is less than before the configuration change.

12. Look for messages in `UnityEifReceiver.log` that show the EIF Receiver posting the JSON payload.

```
04/16/15 17:02:37:820 UTC [pool-6-thread-1] INFO - LogEventPoster : Posting  
Post Data Json of size:1151  
..  
04/16/15 17:02:37:853 UTC [pool-6-thread-1] INFO - LogEventPoster : Posting  
Post Data Json of size:1165
```



```
...  
04/16/15 17:02:37:925 UTC [pool-6-thread-1] INFO - LogEventPoster : Posting  
Post Data Json of size:897
```



**Important:** The size of the posted data is much smaller than before you made the configuration change. The EIF receiver is posting data to the generic receiver more frequently now, but the size of each post is smaller.

13. Look for messages in GenericReceiver.log about the size of each batch of messages. This is the number of log messages in each batch. The full path to this log file is  
/opt/IBM/LogAnalysis/logs/GenericReceiver.log.

```
04/16/15 17:02:37:818 UTC [Default Executor-thread-8] INFO -  
DataCollectorRestServlet : Batch of Size 5 processed and encountered 0 failures  
...  
04/16/15 17:02:37:851 UTC [Default Executor-thread-5] INFO -  
DataCollectorRestServlet : Batch of Size 6 processed and encountered 0 failures  
...  
04/16/15 17:02:37:924 UTC [Default Executor-thread-3] INFO -  
DataCollectorRestServlet : Batch of Size 4 processed and encountered 0 failures
```



**Important:** Notice that the number of messages in each batch are much lower after you made the configuration change. The numbers changed because the message buffer of the EIR receiver is smaller. The EIF receiver now sends a batch of messages to the generic receiver when the EIF receiver buffer reaches 1024 bytes.

14. If you are tailing the log files, press Ctrl+C to stop the tail in both logs.

## Exercise 2 Solr administration

In this exercise, you view details about the Solr file system.

1. Look at the file system where Solr is indexing log data.
  - a. Change to the following directory:  

```
cd /opt/IBM/LogAnalysis/solr-4.7.1/scala_instance1/solr
```
  - b. List the contents of the directory. There is one subdirectory for each shard and each day of log data.

```
ls -l  
total 84
```

```
drwxr-xr-x 2 netcool ncoadmin 4096 Feb 19 2014 bin
-rw-r--r-- 1 netcool ncoadmin 2473 Feb 19 2014 README.txt
-rw-r--r-- 1 netcool ncoadmin 1715 Feb 19 2014 solr.xml
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 15 19:06
UnityCollection_15_04_2015_00_00_00_UTC_shard1_replica1
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 15 19:06
UnityCollection_15_04_2015_00_00_00_UTC_shard2_replica1
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 16 16:03
UnityCollection_16_04_2015_00_00_00_UTC_shard1_replica1
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 16 16:03
UnityCollection_16_04_2015_00_00_00_UTC_shard2_replica1
```



**Note:** In this example, two shards are configured. There is log data from two different days: April 15 and April 16. The number of days in your environment might be different from this example.

2. Look at the indexed log files in the Solr file system.

- a. Change to one of the **UnityCollection\_<date>\_UTC\_shard2\_replica1** directories. Your directory names have a more recent date.

```
cd UnityCollection_15_04_2015_00_00_00_UTC_shard2_replica1
```

- b. List the contents of the **data/index** subdirectory.

```
ls -l data/index/
...
-rw-r--r-- 1 netcool ncoadmin 172827 Apr 15 19:46 _r.fdt
-rw-r--r-- 1 netcool ncoadmin 232 Apr 15 19:46 _r.fdx
-rw-r--r-- 1 netcool ncoadmin 4332 Apr 15 19:46 _r.fnm
-rw-r--r-- 1 netcool ncoadmin 92942 Apr 15 19:46 _r_Lucene41_0.doc
-rw-r--r-- 1 netcool ncoadmin 93939 Apr 15 19:46 _r_Lucene41_0.pos
-rw-r--r-- 1 netcool ncoadmin 138017 Apr 15 19:46 _r_Lucene41_0.tim
-rw-r--r-- 1 netcool ncoadmin 2488 Apr 15 19:46 _r_Lucene41_0.tip
-rw-r--r-- 1 netcool ncoadmin 29210 Apr 15 19:46 _r_Lucene45_0.dvd
-rw-r--r-- 1 netcool ncoadmin 1238 Apr 15 19:46 _r_Lucene45_0.dvm
-rw-r--r-- 1 netcool ncoadmin 10106 Apr 15 19:46 _r.nvd
-rw-r--r-- 1 netcool ncoadmin 79 Apr 15 19:46 _r.nvm
-rw-r--r-- 1 netcool ncoadmin 381 Apr 15 19:46 _r.si
...
```



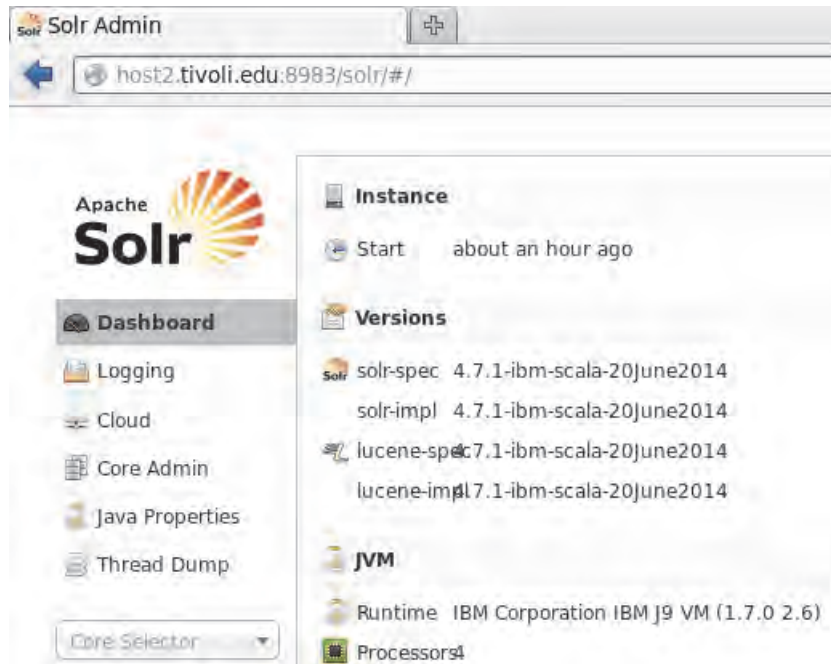
**Note:** The files in the **data/index** subdirectory are the raw and annotated log messages that IBM Operations Analytics Log Analysis has processed. These files are in compressed binary format.

3. Use the Solr administration interface to look at the Solr file system and indexed data.

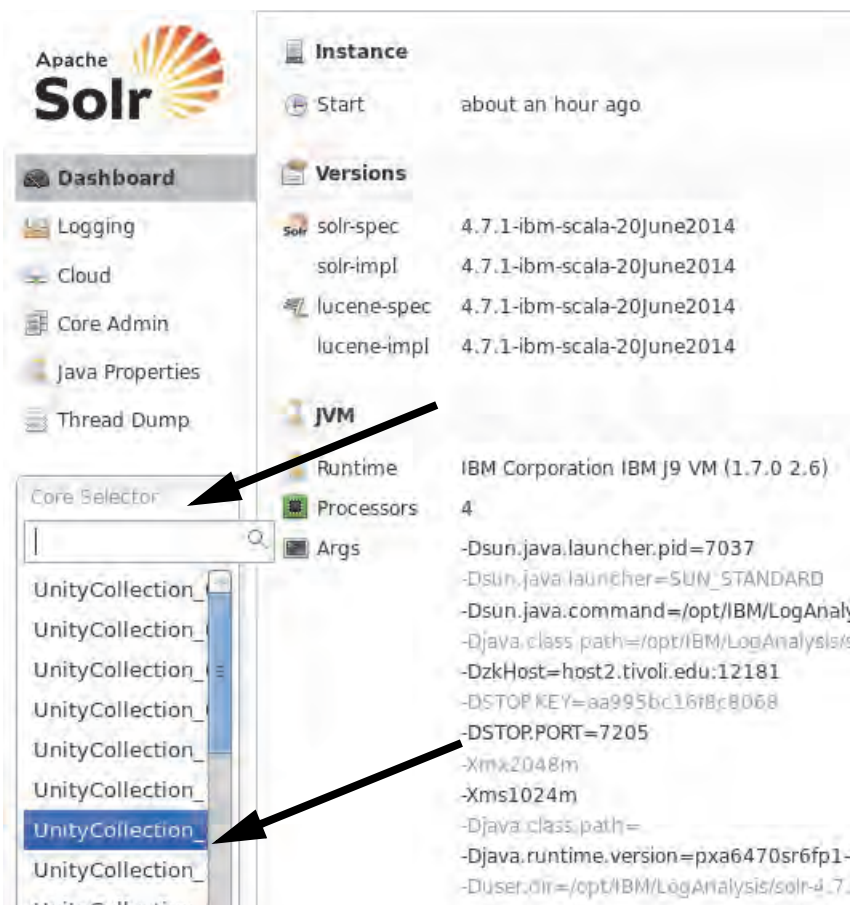
- a. Open a Firefox browser. Enter the following URL:

```
http://host2.tivoli.edu:8983/solr
```

- b. Press F11 or click **View > Full Screen** in Firefox to view the entire Solr administration interface.



- c. Click **Core Selector**, and select one of the shards in your system.



- d. Look at the directory structure in the **Instance** summary. This structure matches the file system that you saw in a preceding step.

The screenshot displays the Solr administration interface. On the left, the **Statistics** section shows various metrics: Last (about 23 hours ago), Modified, Num Docs (26650), Max Doc (26667), Heap Memory (61080), Usage, Deleted Docs (17), Version (62), Segment (8), Count, Optimized (with a red icon), and Current (with a green icon). Below these is an **optimize now** button. On the right, the **Instance** section shows the directory structure for the Solr instance. It lists CWD, Instance, Data, Index, and Impl. The Data and Index paths are identical: `/opt/IBM/LogAnalysis/solr-4.7.1/scala_instance1/solr/UnityCollection_15_04_2015_00_00_00 UTC_4.7.1/scala_instance1/`. The Index path is `/opt/IBM/LogAnalysis/solr-4.7.1/scala_instance1/solr/UnityCollection_15_04_2015_00_00_00 UTC_4.7.1/scala_instance1/data/index`. The Impl is `org.apache.solr.core.NRTCachingDirectoryFactory`. An arrow points from the **optimize now** button in the Statistics section to the **Data** field in the Instance section.

- e. Click **Query** in the navigation menu on the left.
- f. Enter **100** in the rows field.

- g. Click **Execute Query**. You might have to scroll down to see the Execute Query button.

The screenshot shows the Apache Solr Admin UI. On the left is a sidebar with navigation links: Dashboard, Logging, Cloud, Core Admin, Java Properties, Thread Dump, and a dropdown menu for 'UnityCollection...'. Below these are links for Overview, Analysis, Dataimport, Documents, Files, Ping, Plugins / Status, **Query** (highlighted with arrow 1), Replication, and Schema Browser. The main content area on the right contains query execution fields: 'q' (query text), 'fq' (facet query), 'sort', 'start, rows' (with '0' and '100' in input boxes), 'fl' (fields list), and 'df' (distribution fields). Below these is a 'Raw Query Parameters' section showing 'key1=val1&key2=val2'. Further down are 'wt' (set to 'json') and several checkboxes: 'indent' (checked), 'debugQuery', 'dismax', 'edismax', 'facet', 'spatial', and 'spellcheck'. At the bottom is a blue 'Execute Query' button, which is pointed to by arrow 3. A large blue arrow labeled 'Scroll' points downwards, indicating the need to scroll to reach the button. Arrow 2 points to the 'rows' input field.



- h. Scroll down through the search results.

```
{
  "id": "1429124789503_1429124789506_15_75",
  "message_tf": [
    "The com.ibm.tivoli.reporting.ejb.TivoliReportingEJBRemote interface, which is specified fo",
  ],
  "datasourceHostname_tf": [
    "host2.tivoli.ibm.com"
  ],
  "timestamp": "2015-03-11T13:32:32.342",
  "threadID_tf": [
    "0000000b"
  ],
  "_writetime": "2015-04-15T13:06:29.506Z",
  "_datasource": "WAS_SystemOut",
  "shortname_s": [
    "EJBInjectionB"
  ],
  "severity_ts": "W",
  "msgclassifier_tf": [
    "CWNN0033W"
  ],
  "logRecord": [
    "[03/11/15 13:32:32.340 UTC] 0000000b EJBInjectionB W CWNN0033W: The com.ibm.tivoli.repo",
  ],
  "_version_": 1498545955493904400
},
```



**Note:** This query is showing the last 100 annotated messages that were saved in this shard.

- i. Look at a graphical representation of the shards. Click **Cloud > Graph**. Your system is saving data across two shards.



- j. Close Firefox.

4. Edit the Solr configuration to save data over more shards.

- a. Open the **unitysetup.properties** file with a text editor.

```
vi
/opt/IBM/LogAnalysis/wlp/usr/servers/Unity/apps/Unity.war/WEB-INF/
unitysetup.properties
```

- b. Find the following two properties.

- ◆ INDEX\_NUM\_SHARDS
- ◆ COLLECTION\_ASYNC\_WINDOW

```
#velocity specific properties
# Boolean flag specifying synchronous (true) or asynchronous (false) indexing
asyncIndexing=true
```

```
INDEX_NUM_SHARDS=2
```

```
# Async time window for creating new time based collection. This value should
NOT be modified by the user
# It can be only an Integer value in hours or days ( < 365 ). If in hours it
has to be a factor of 24 greater than or equal to 6.
```

```
# It is specified as nh (n hours) or nd (n days); units being h (hour), or d
(day)
```

```
COLLECTION_ASYNC_WINDOW = 1d
```

```
#####
```



**Note:** The number of shards is set to two. The collection window is one day. This is why you have two subdirectories (shards) for each day of data.

- c. Change the **INDEX\_NUM\_SHARDS** property to 8.

```
# Boolean flag specifying synchronous (true) or asynchronous (false) indexing
asyncIndexing=true
```

```
INDEX_NUM_SHARDS=8
```

```
# Async time window for creating new time based collection. This value should
NOT be modified by the user
```

- d. Save and close the file.

- e. Stop IBM Operations Analytics Log Analysis.

```
/opt/IBM/LogAnalysis/utilities/unity.sh -stop
```

- f. Start IBM Operations Analytics Log Analysis.

```
/opt/IBM/LogAnalysis/utilities/unity.sh -start
```





**Important:** The new shard configuration takes effect at the start of the next collection window. Remember that the collection window is one day (`COLLECTION_ASYNC_WINDOW = 1d`). This setting means that you see the new shard directories at the start of the next day.

## Optional steps: Verifying the Solr configuration change



**Hint:** If you do not want to wait a full day, run the following steps to change the system time to the next day.

5. Change to the root user and change the system clock to the next day.
  - a. Open a new terminal window. Run the following command to switch to the **root** user. The password is **object00**.

```
su -  
Password: object00
```
  - b. Run the following command to view the current system date.

```
date  
Thu Apr 16 19:45:07 UTC 2015
```
  - c. Run a command like the following example to set the system date to the next day. You must change the date string in the command. Enter the date of the next day.

```
date +%Y%m%d -s "20150417"
```
  - d. Run the `date` command again to confirm that the system clock was changed.

```
date  
Fri Apr 17 00:00:08 UTC 2015
```
  - e. Close the terminal window. Complete the rest of the steps in this exercise as the **netcool** user.



**Note:** The next steps are optional. These steps show you how to verify the change to the shard configuration. Follow these steps only if you have waited a day after you changed the **INDEX\_NUM\_SHARDS** property to **8**. You can also use these steps if you changed the system clock to the next day.

6. Run the following command as the **netcool** user to create more log messages in the WAS\_SystemOut data source.

```
/software/log_samples/scripts/WAS_Logs.sh
```

7. Return to the Solr administration interface to look at the Solr file system and indexed data.
  - a. Open a Firefox browser. Enter the following URL:  
`http://host2.tivoli.edu:8983/solr`
  - b. Click **Cloud > Graph**.

The screenshot shows the Apache Solr administration interface. On the left is a sidebar with navigation links: Dashboard, Logging, Cloud (selected), Tree, Graph (highlighted), Graph (disabled), Dump, Core Admin, Java Properties, and Thread Dump. The main content area shows a tree structure of collections and shards. A red box highlights the 'UnityCollection\_17' collection, which contains 8 shards (shard1 through shard8). A blue arrow points to the 'Graph' link in the left sidebar with the text 'Click Graph'. Another blue arrow points to the highlighted shards with the text 'You have 8 shards after the change'.

Notice that after the change, and after data was ingested in a new collection window (the day after the change), there are eight shards instead of two.

8. Look at the Solr directory structure that was created to support the new shards.
  - a. Change to the following directory:  
`cd /opt/IBM/LogAnalysis/solr-4.7.1/scala_instance1/solr`

- b. List the contents of the directory. There is one subdirectory for each shard and each day of log data.

```
ls -l
...
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 17 00:01
UnityCollection_17_04_2015_00_00_00_UTC_shard1_replica1
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 17 00:01
UnityCollection_17_04_2015_00_00_00_UTC_shard2_replica1
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 17 00:01
UnityCollection_17_04_2015_00_00_00_UTC_shard3_replica1
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 17 00:01
UnityCollection_17_04_2015_00_00_00_UTC_shard4_replica1
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 17 00:01
UnityCollection_17_04_2015_00_00_00_UTC_shard5_replica1
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 17 00:01
UnityCollection_17_04_2015_00_00_00_UTC_shard6_replica1
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 17 00:01
UnityCollection_17_04_2015_00_00_00_UTC_shard7_replica1
drwxr-xr-x 3 netcool ncoadmin 4096 Apr 17 00:01
UnityCollection_17_04_2015_00_00_00_UTC_shard8_replica1
...
```



**Note:** Notice that you have eight subdirectories (one for each shard) for the current day collection window.



## **7 Backing up and restoring IBM Operations Analytics Log Analysis exercises**

There are no student exercises for this unit.





